Distributed Provenance Compression

Lay Kuan Loh

lloh@ece.cmu.edu

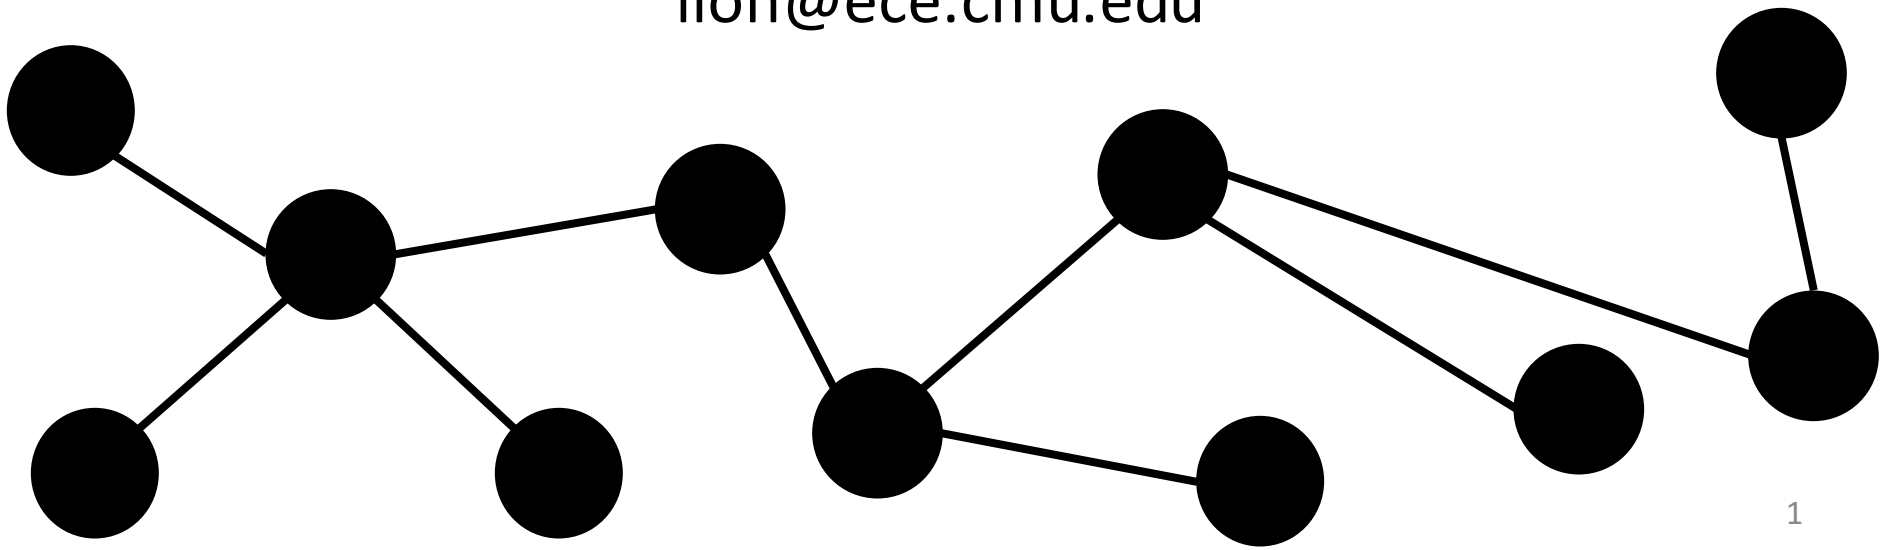

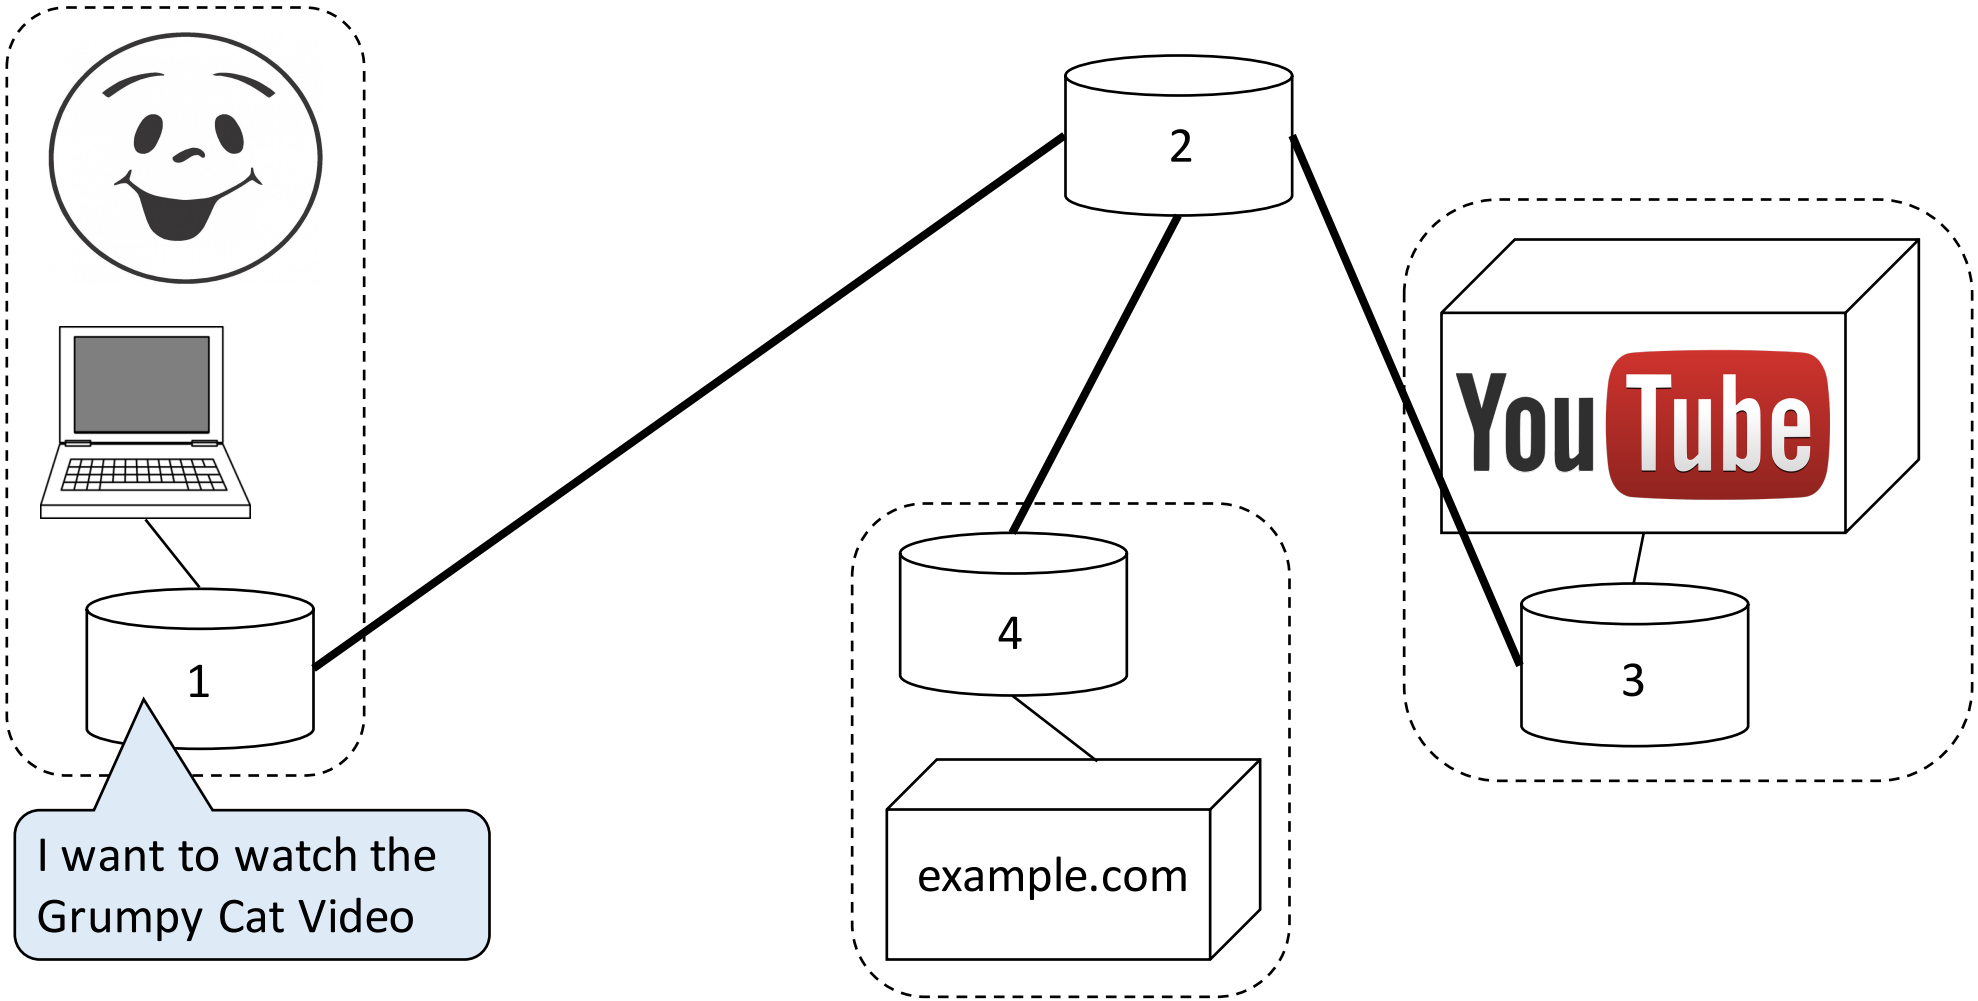
2
YouTube
4
1
3
example.com
I want to watch the
Grumpy Cat Video

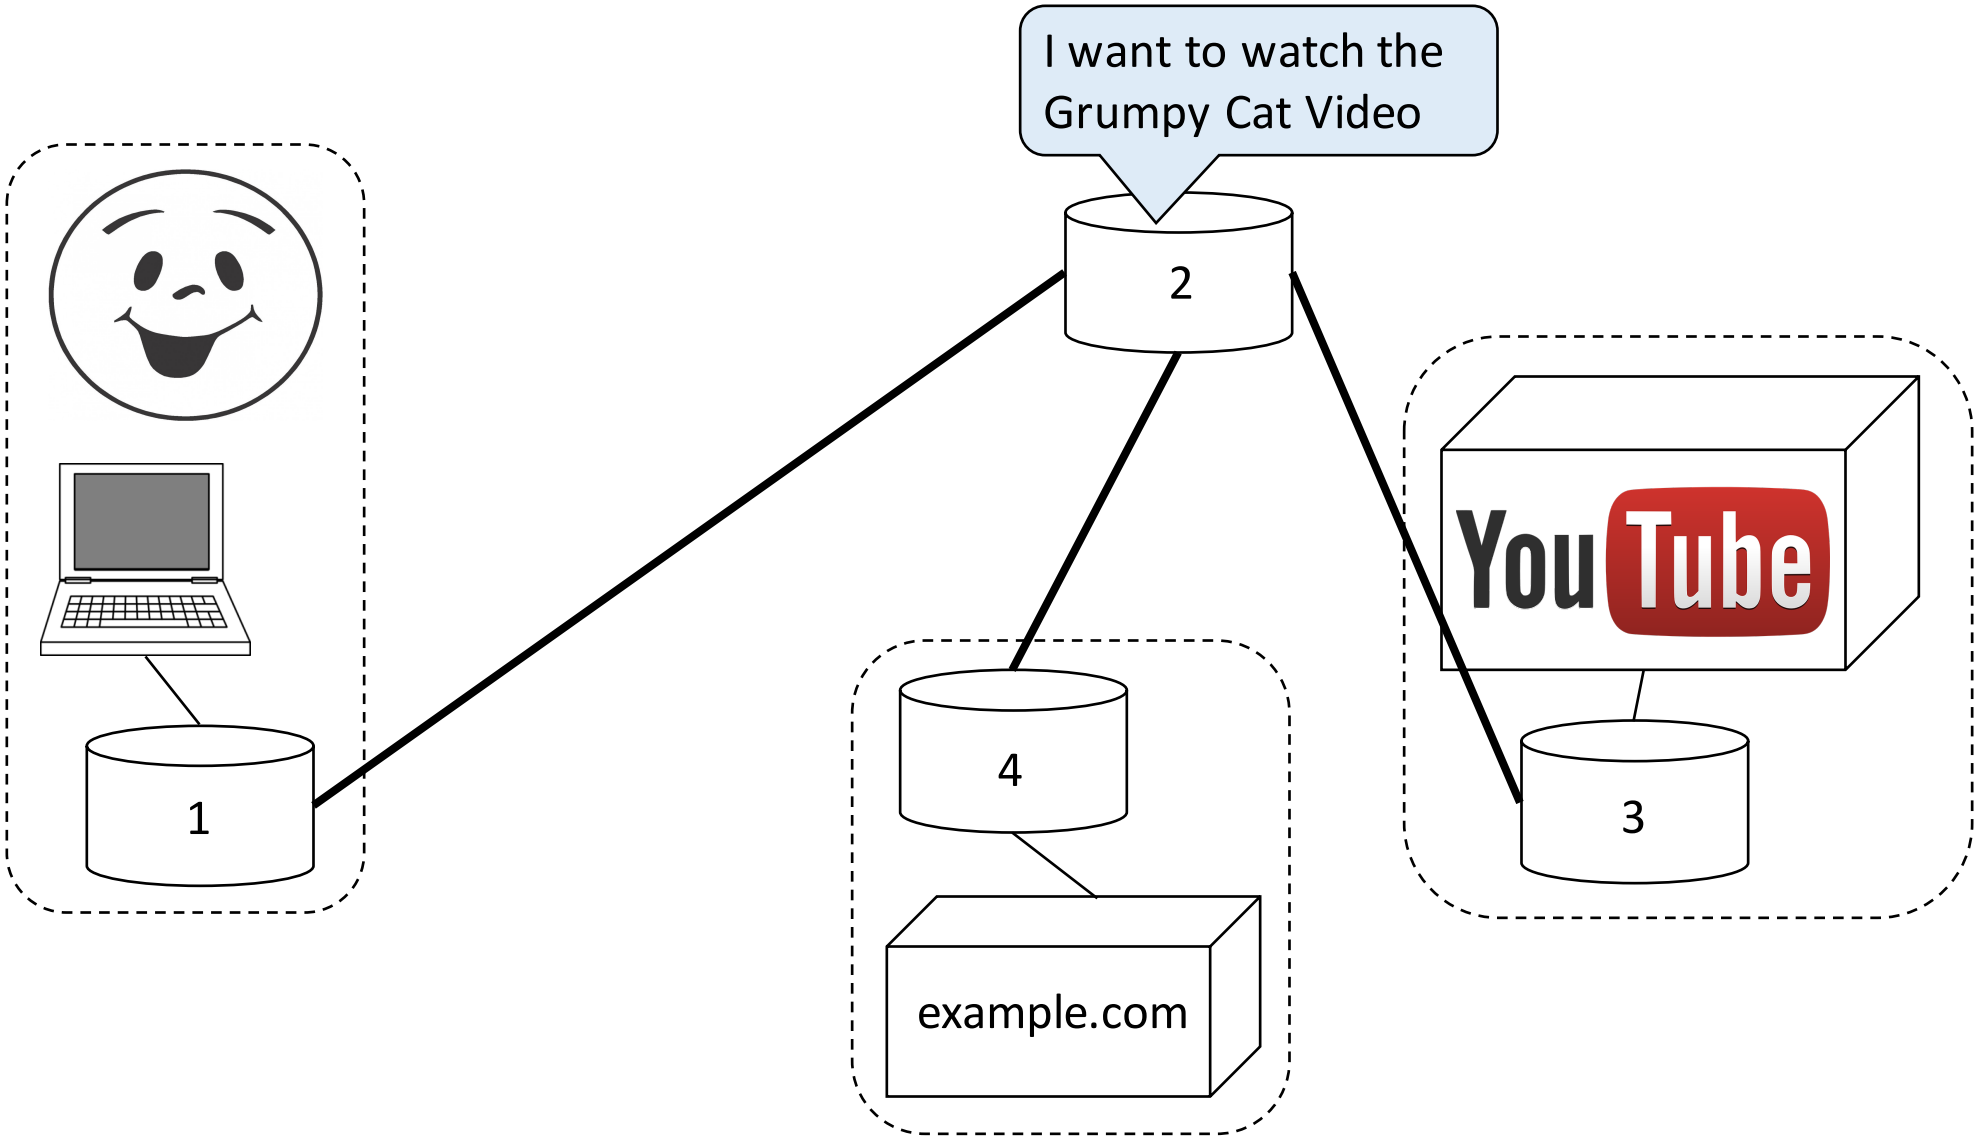
I want to watch the Grumpy Cat Video
2
1
4
example.com
3
YouTube

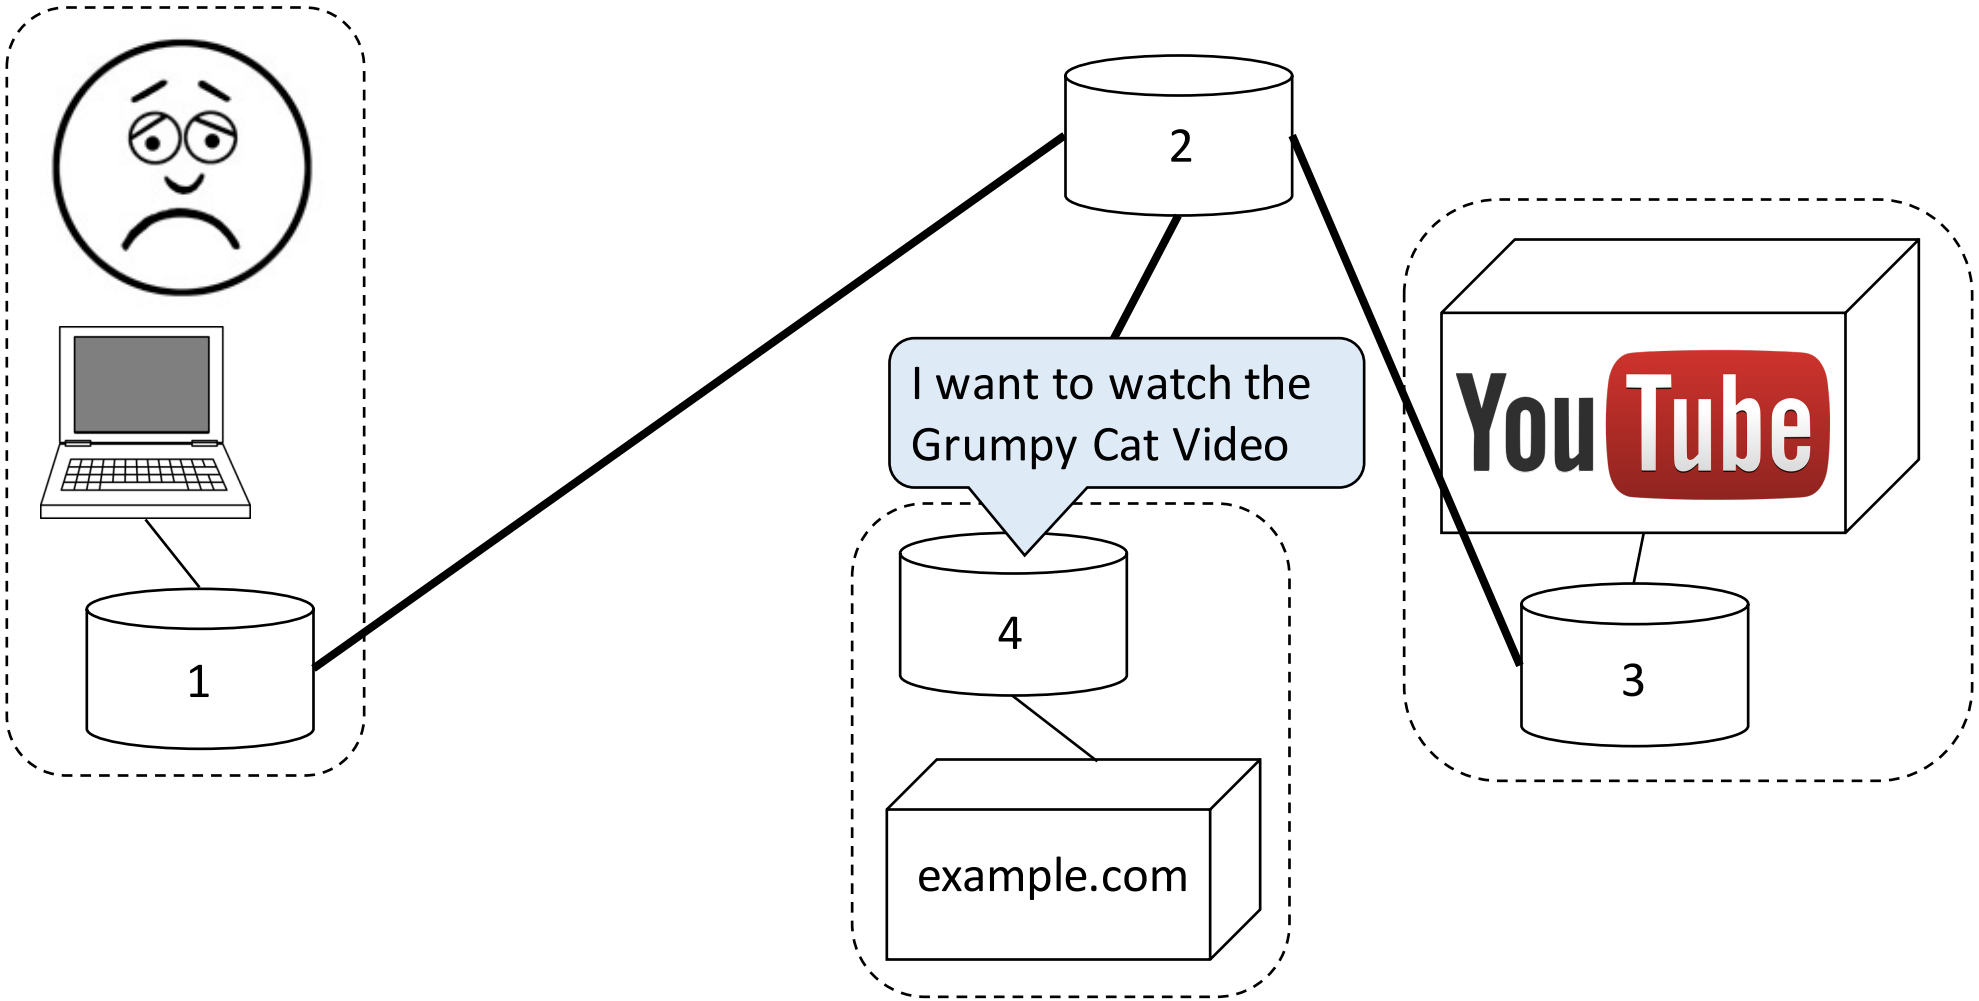
2
I want to watch the Grumpy Cat Video
YouTube
1
4
3
example.com

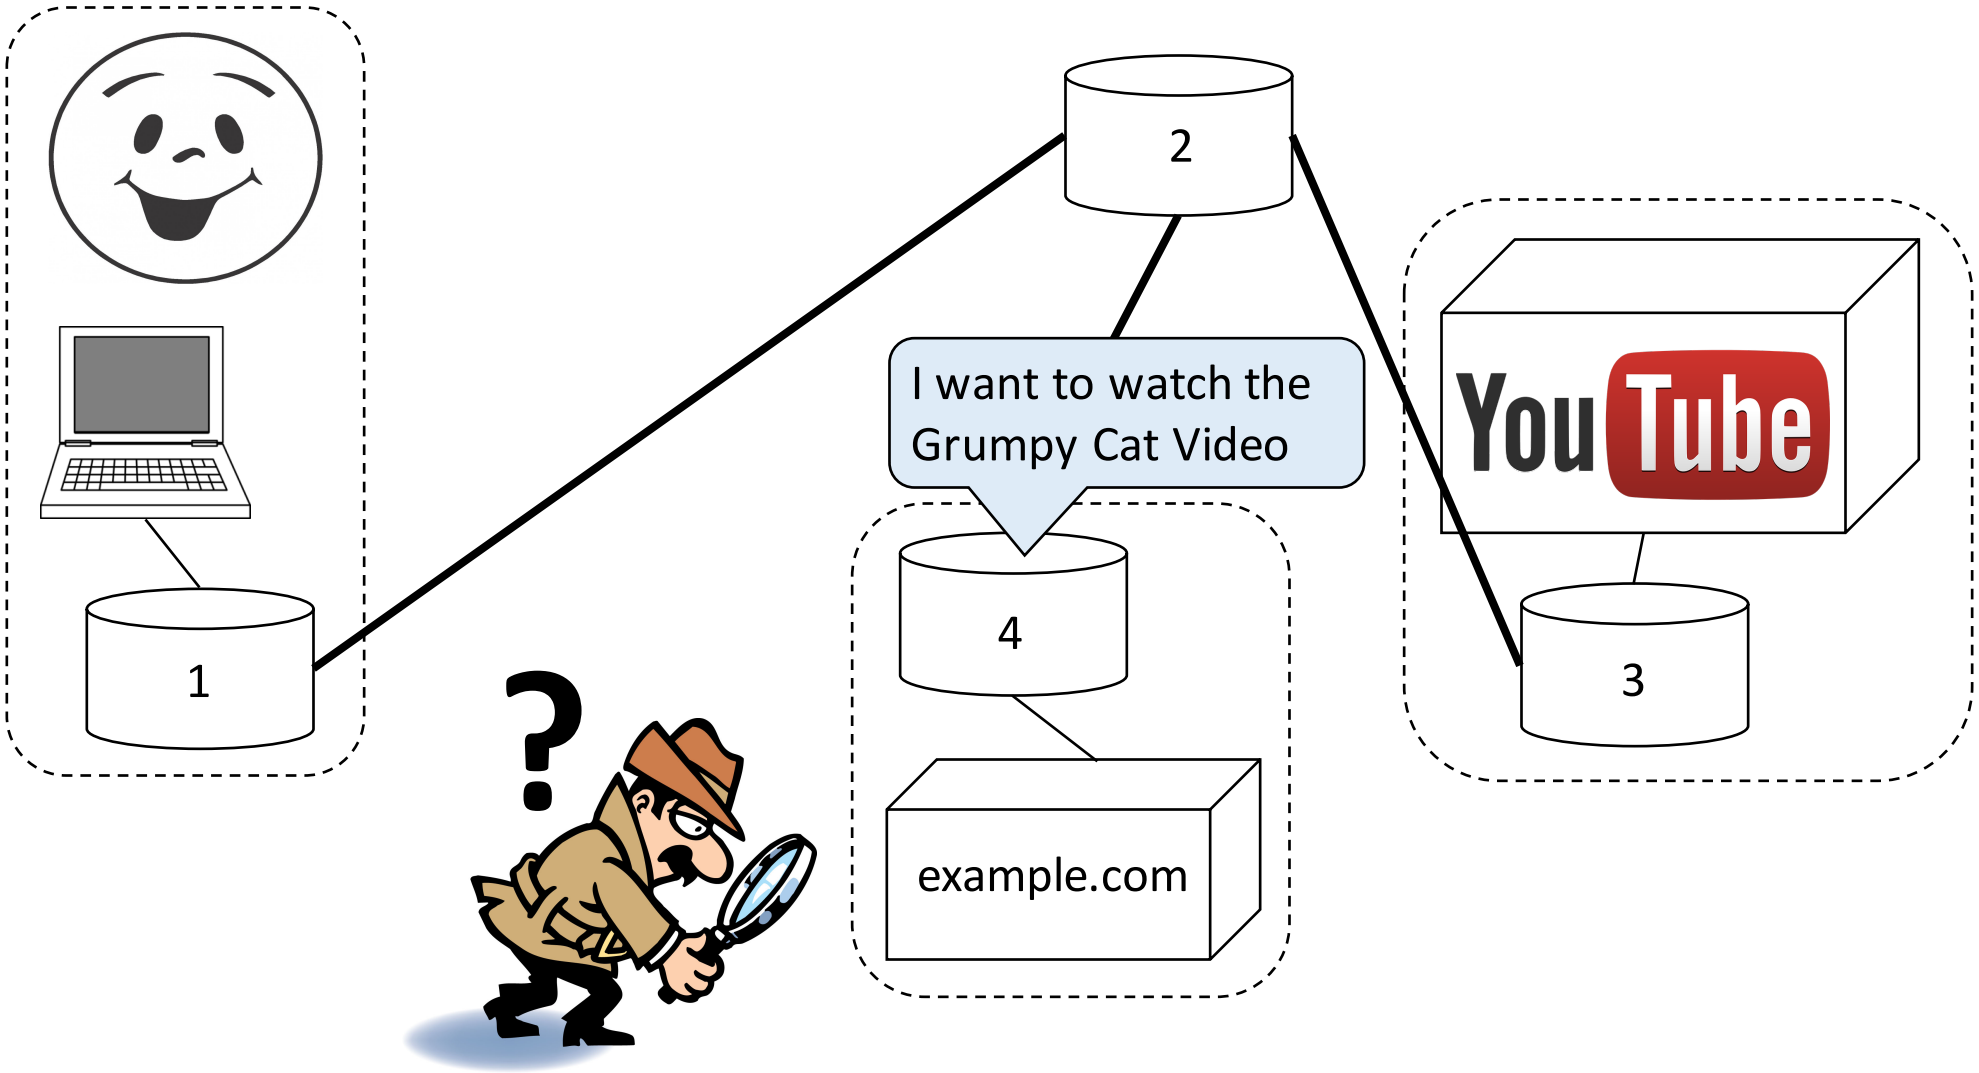
2
I want to watch the Grumpy Cat Video
YouTube
4
1
3
example.com
?

Network Provenance:
"The history and derivations of network state resulting from the execution of a network protocol."

Received packet from Router 1.
Routing table: Go to Router 4.

2

I want to watch the Grumpy Cat Video

4

example.com

Received packet from Router 2.

YouTube

3

1

Routing table: Go to Router 2.

Challenge

Large amount of storage needed to maintain network provenance at Internet-scale.

Our Solution

Identify and remove redundancy in network provenance before storing at runtime.

Roadmap

- ***Background***
- Key insights
- Our compression scheme
- Conclusion

Loo et al., *Declarative Networking*. SIGMOD'06.

Declarative Networks

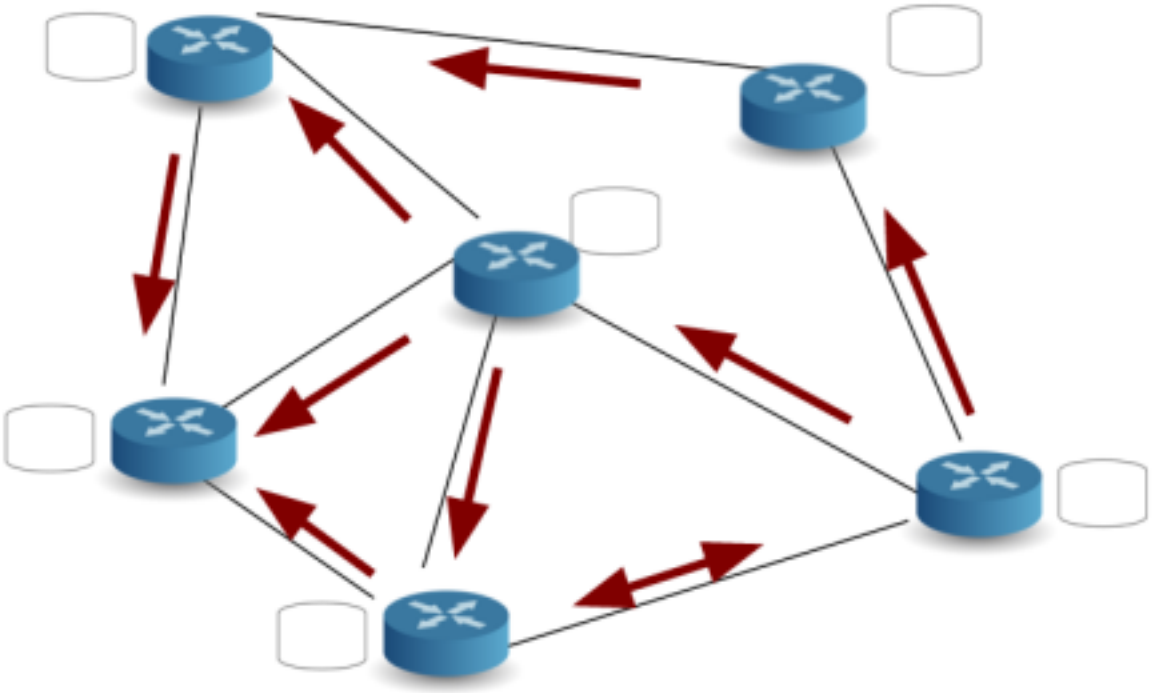

Traditional Network	Declarative Network
Network State • e.g. routing table	Distributed database • Info in routing table stored as tuples

Loo et al., *Declarative Networking*. SIGMOD'06.

Declarative Networks

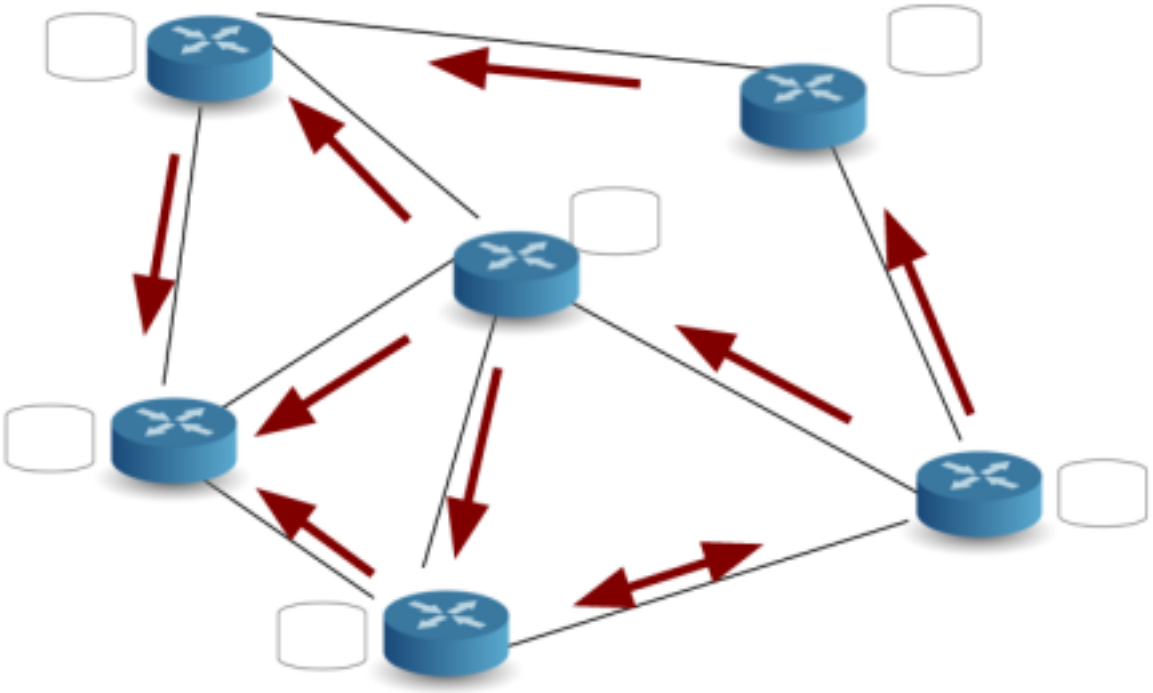

Traditional Network	Declarative Network
Network State • e.g. routing table	Distributed database • Info in routing table stored as ***tuples***
Network Protocol	***Network Datalog (NDLog)*** program

Loo et al., *Declarative Networking*. SIGMOD'06.

Declarative Networks

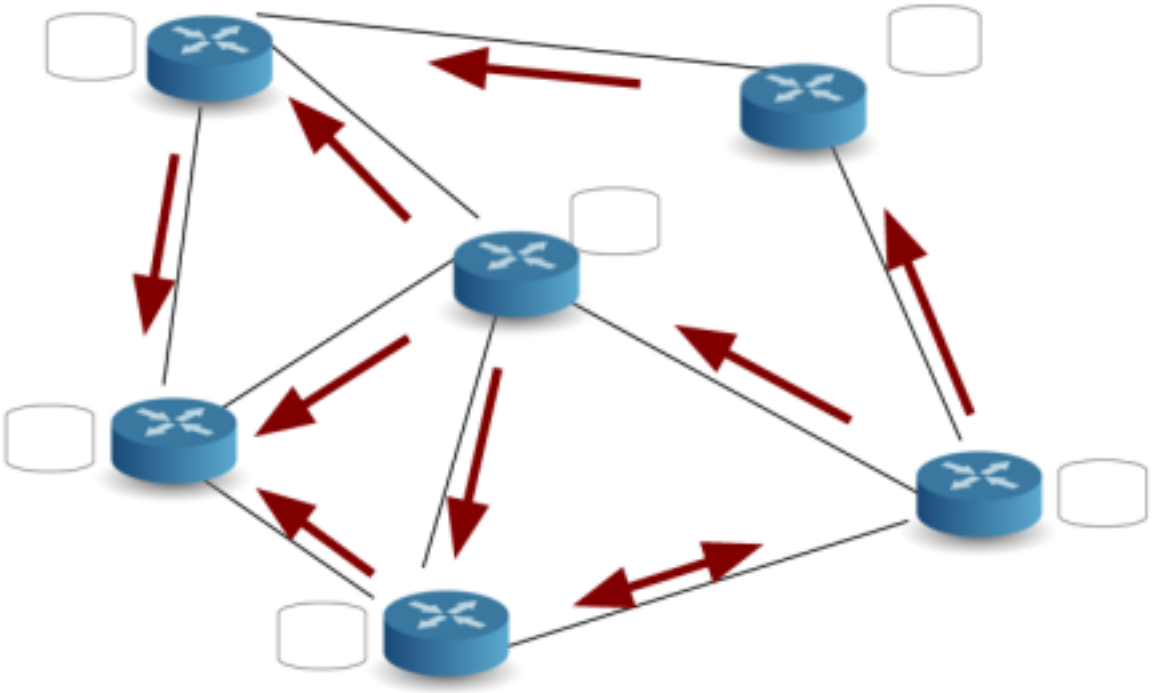

Q: Why ***NDLog***?
A: Concise

Traditional Network	Declarative Network
Network State • e.g. routing table	Distributed database • Info in routing table stored as ***tuples***
Network Protocol	***Network Datalog (NDLog)*** program

Loo et al., *Declarative Networking*. SIGMOD'06.

Review of Network Datalog (NDLog)

<result> :- <condition$_1$>, …, <condition$_N$>

Rule Head

Rule Body

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), Dst==Loc.

Loo et al., *Declarative Networking*. SIGMOD'06.

An example NDLog Program

Packet Forwarding

```
r1 packet(@Neigh,Src,Dst,Payload)
    :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).
r2 recv(@Loc,Src,Dst,Payload)
    :- packet(@Loc,Src,Dst,Payload), Dst==Loc.
```

Fast-changing

Updated automatically by program execution

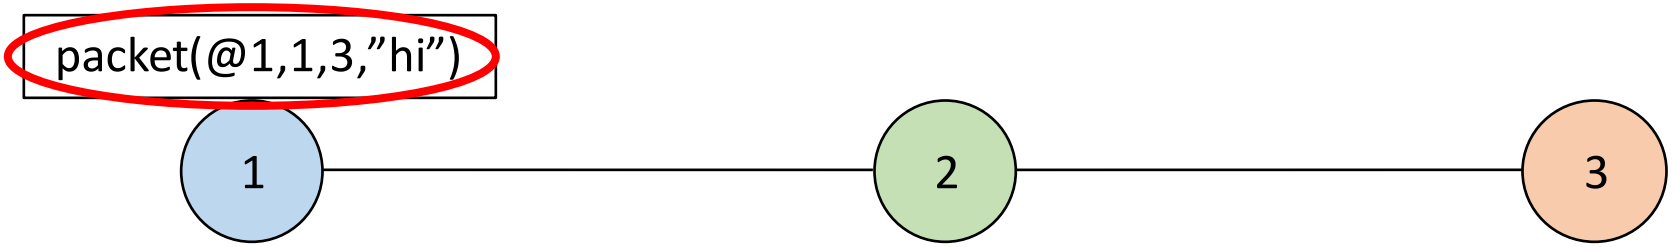

Loo et al., *Declarative Networking*. SIGMOD'06.

An example NDLog Program

Packet Forwarding

```
r1 packet(@Neigh,Src,Dst,Payload)
    :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).
r2 recv(@Loc,Src,Dst,Payload)
    :- packet(@Loc,Src,Dst,Payload),  Dst==Loc.
```

Location Specifier

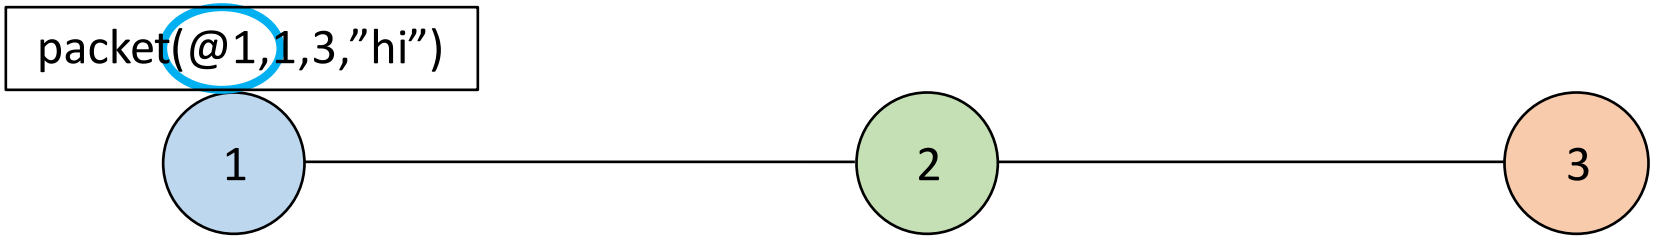

Loo et al., *Declarative Networking*. SIGMOD'06.

An example NDLog Program

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), Dst==Loc.

Slow-changing

Interval between updates is usually longer than the lifespan of program execution

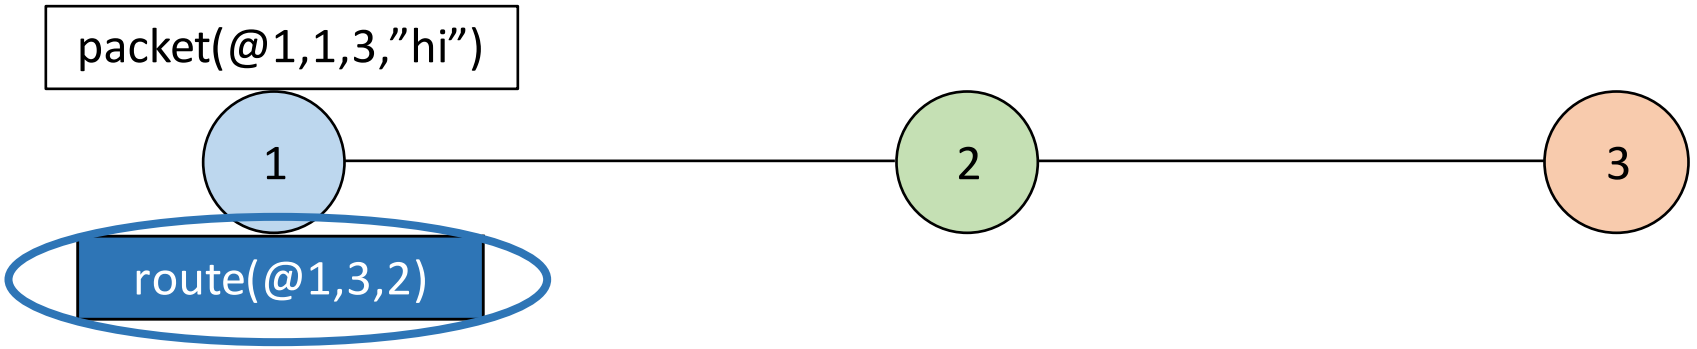

Loo et al., *Declarative Networking*. SIGMOD'06.

An example NDLog Program

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload)
:- packet(@Loc,Src,Dst,Payload), Dst==Loc.

Fast-changing

Updated automatically by program execution

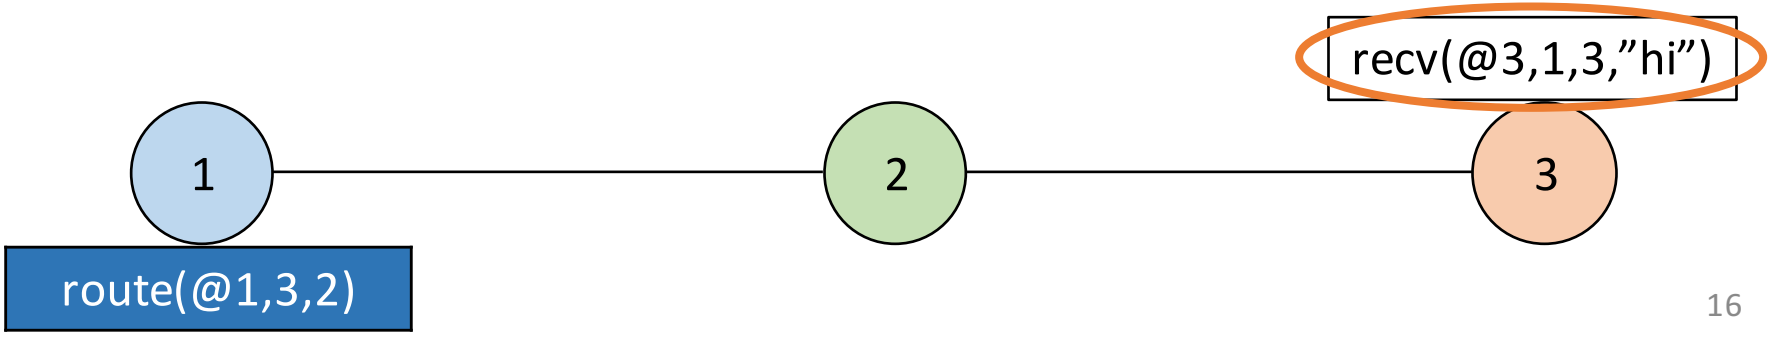

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

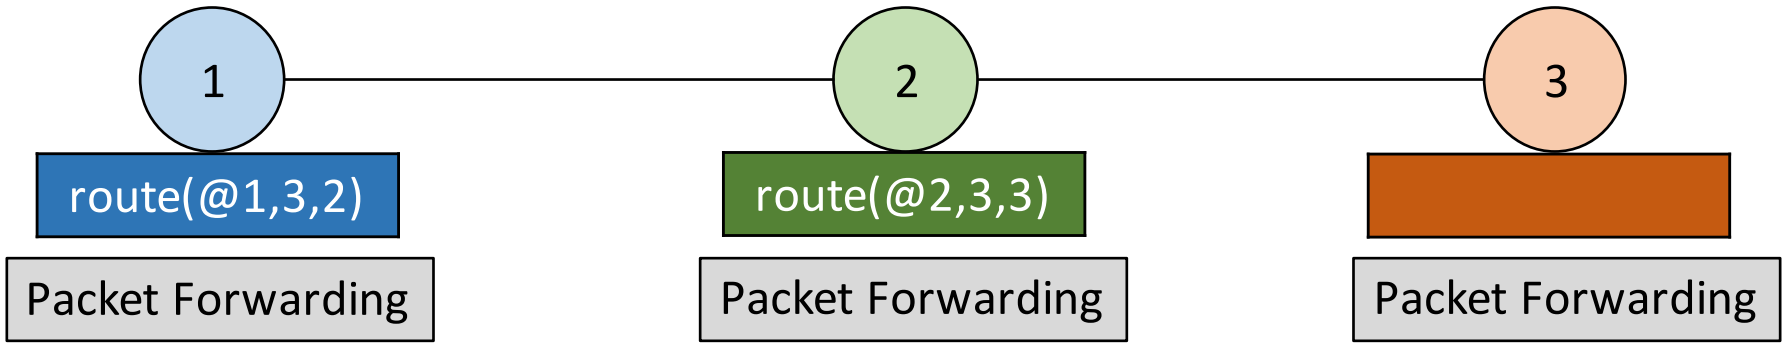

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

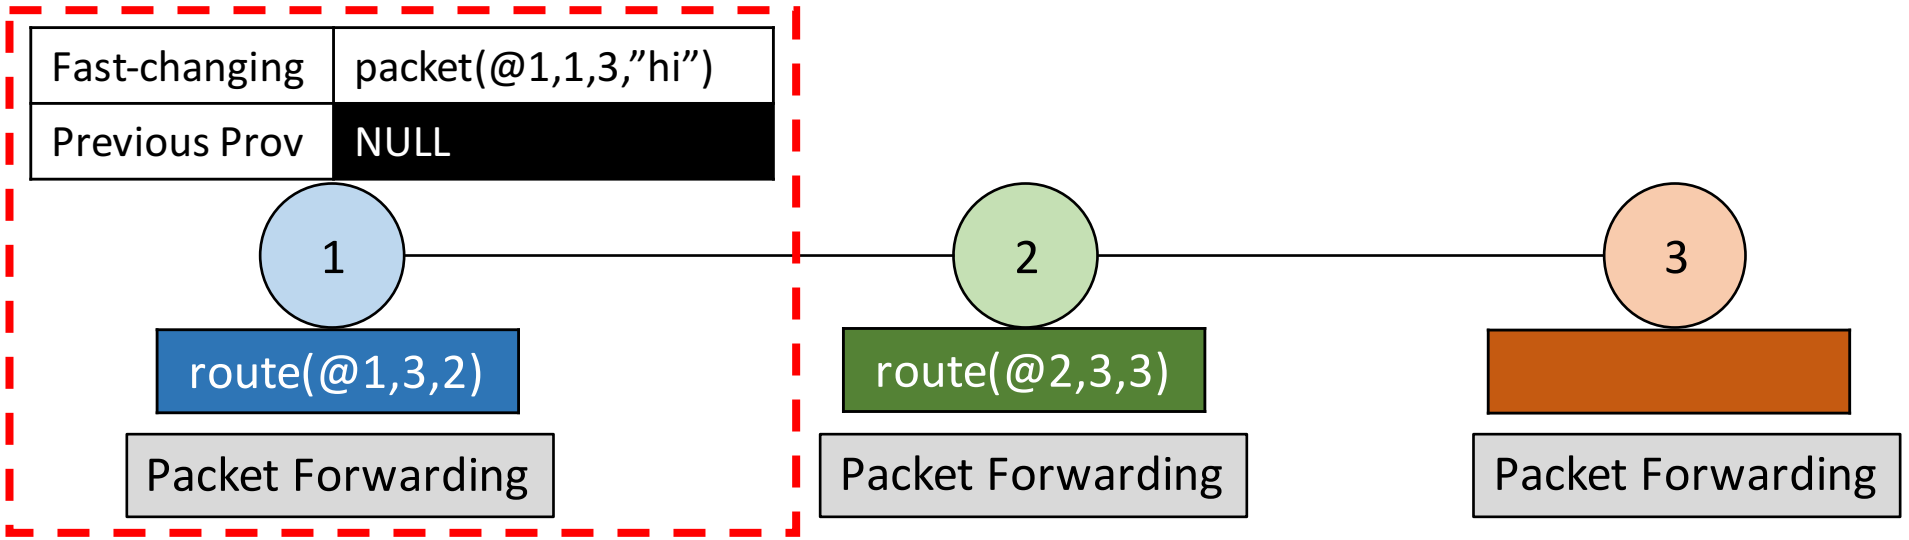

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

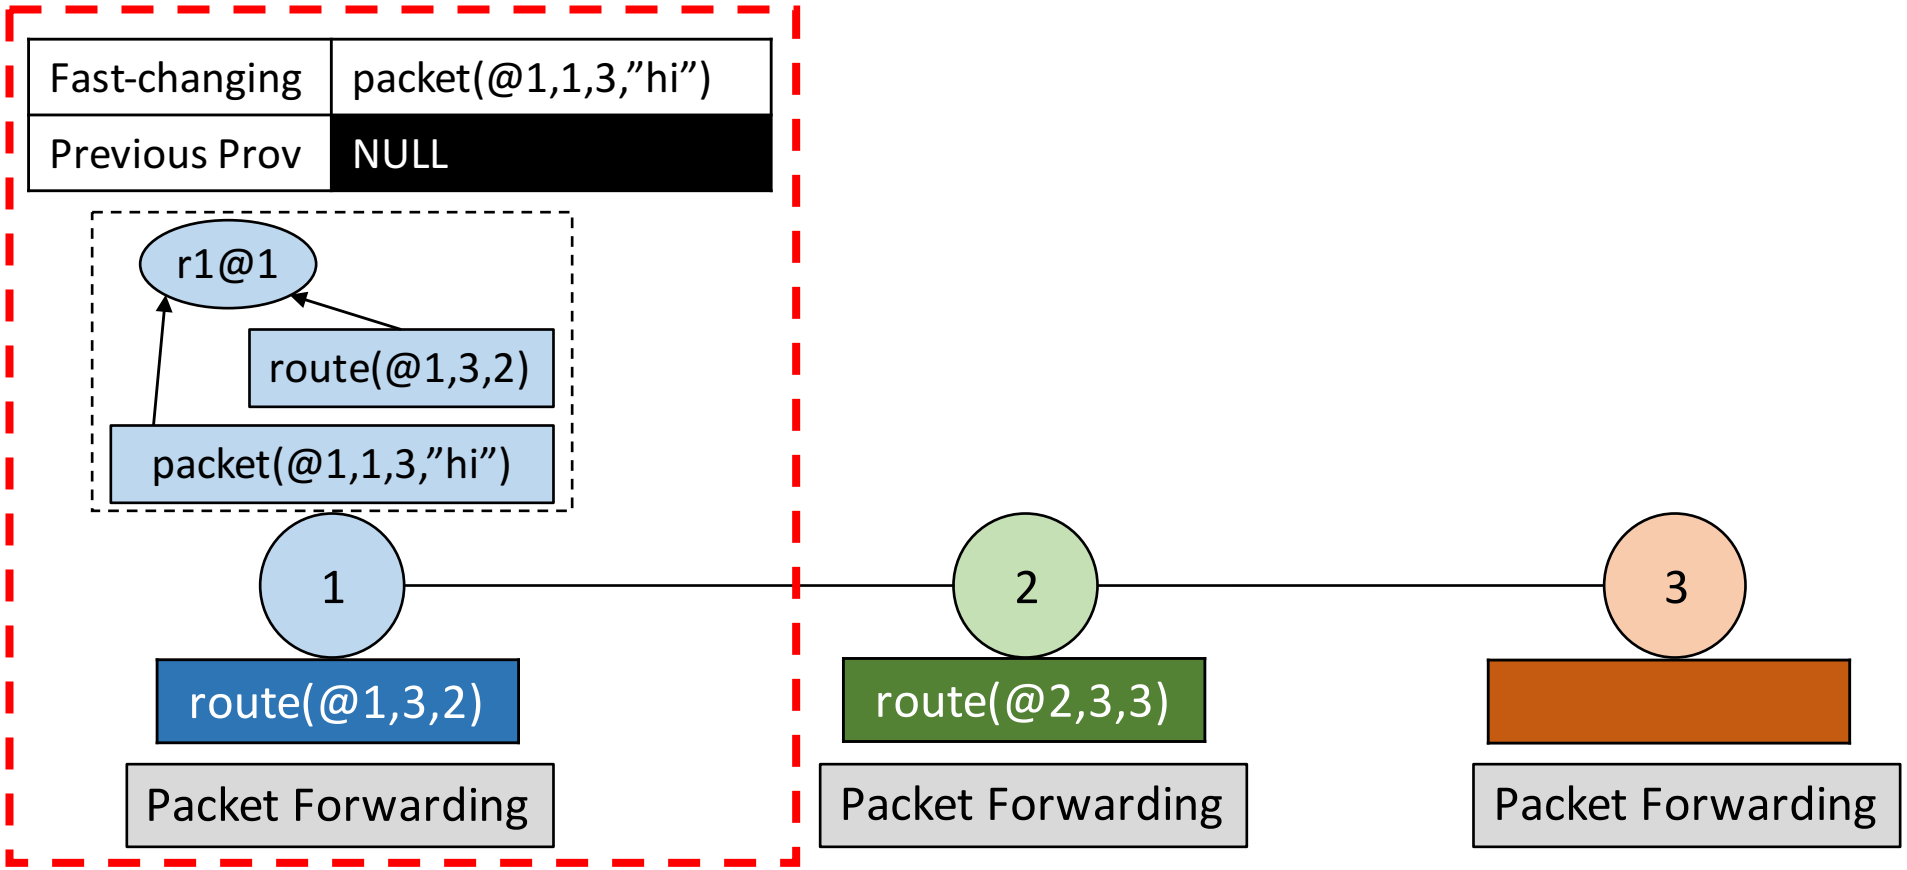

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

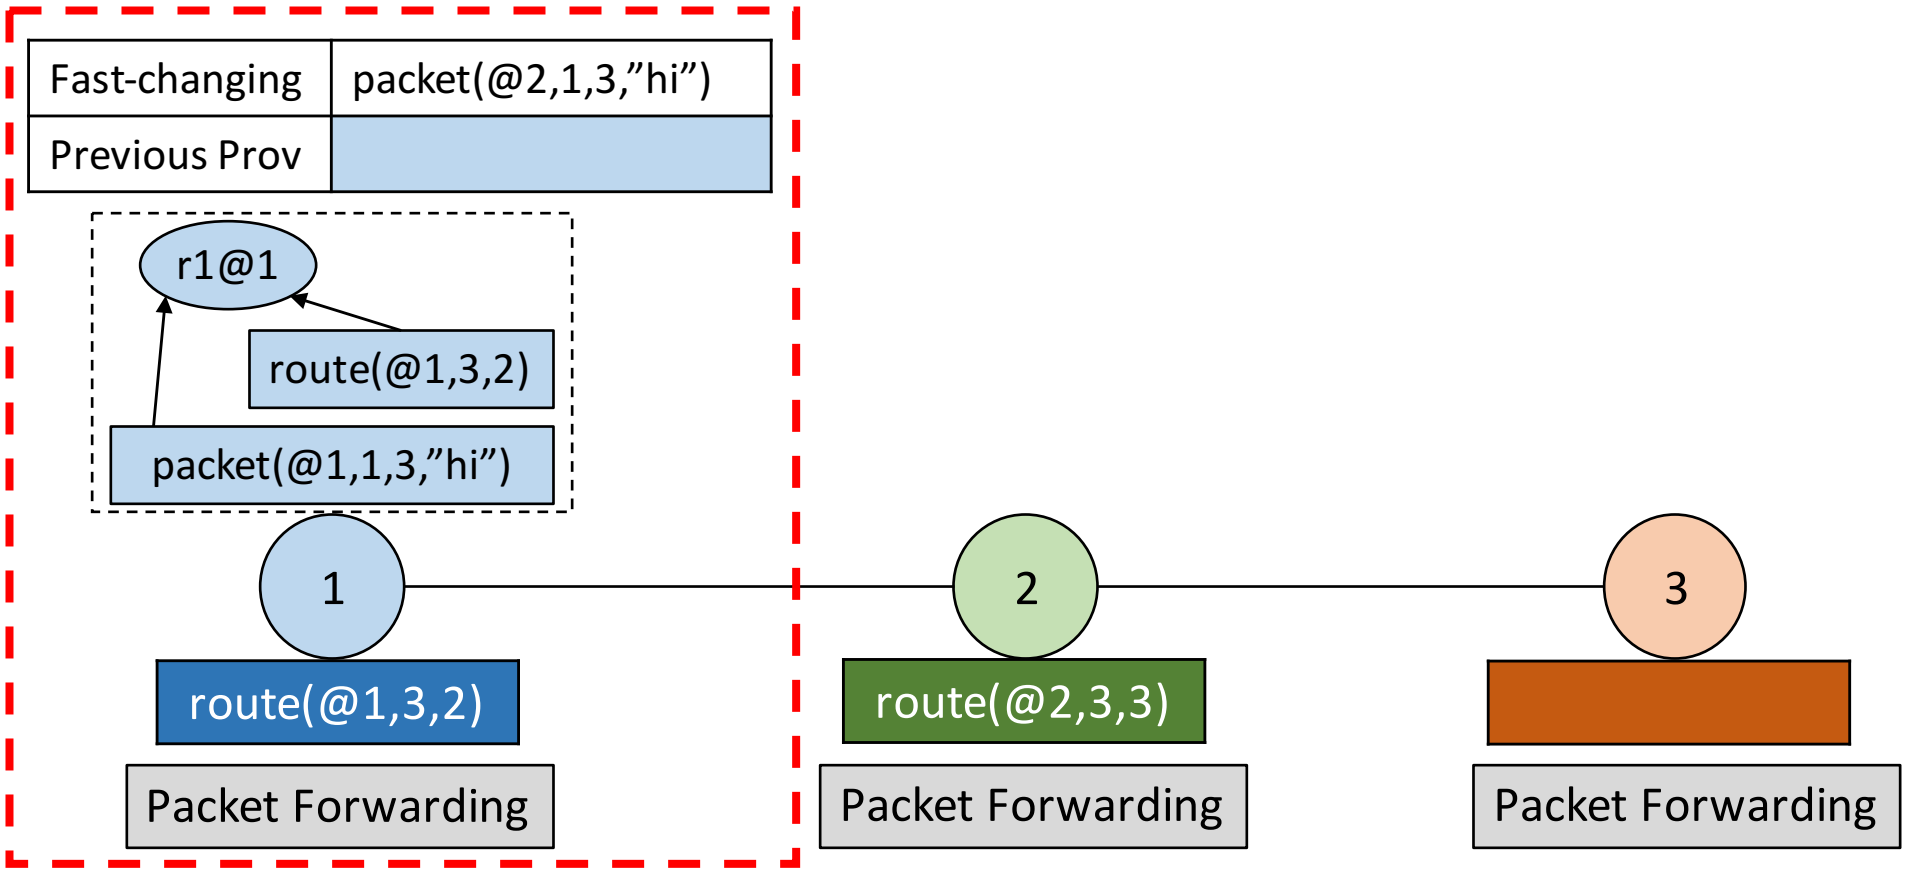

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

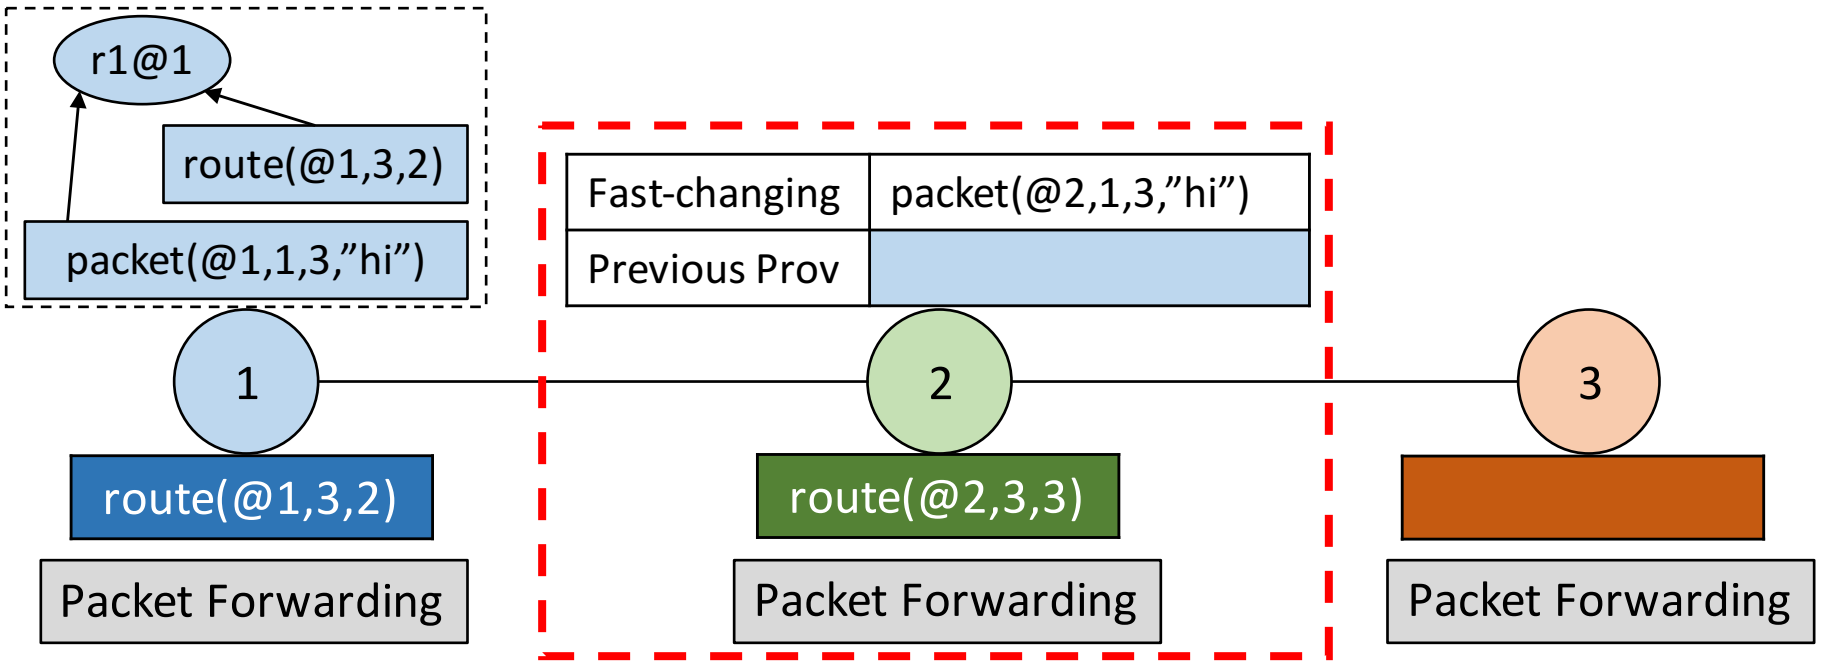

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

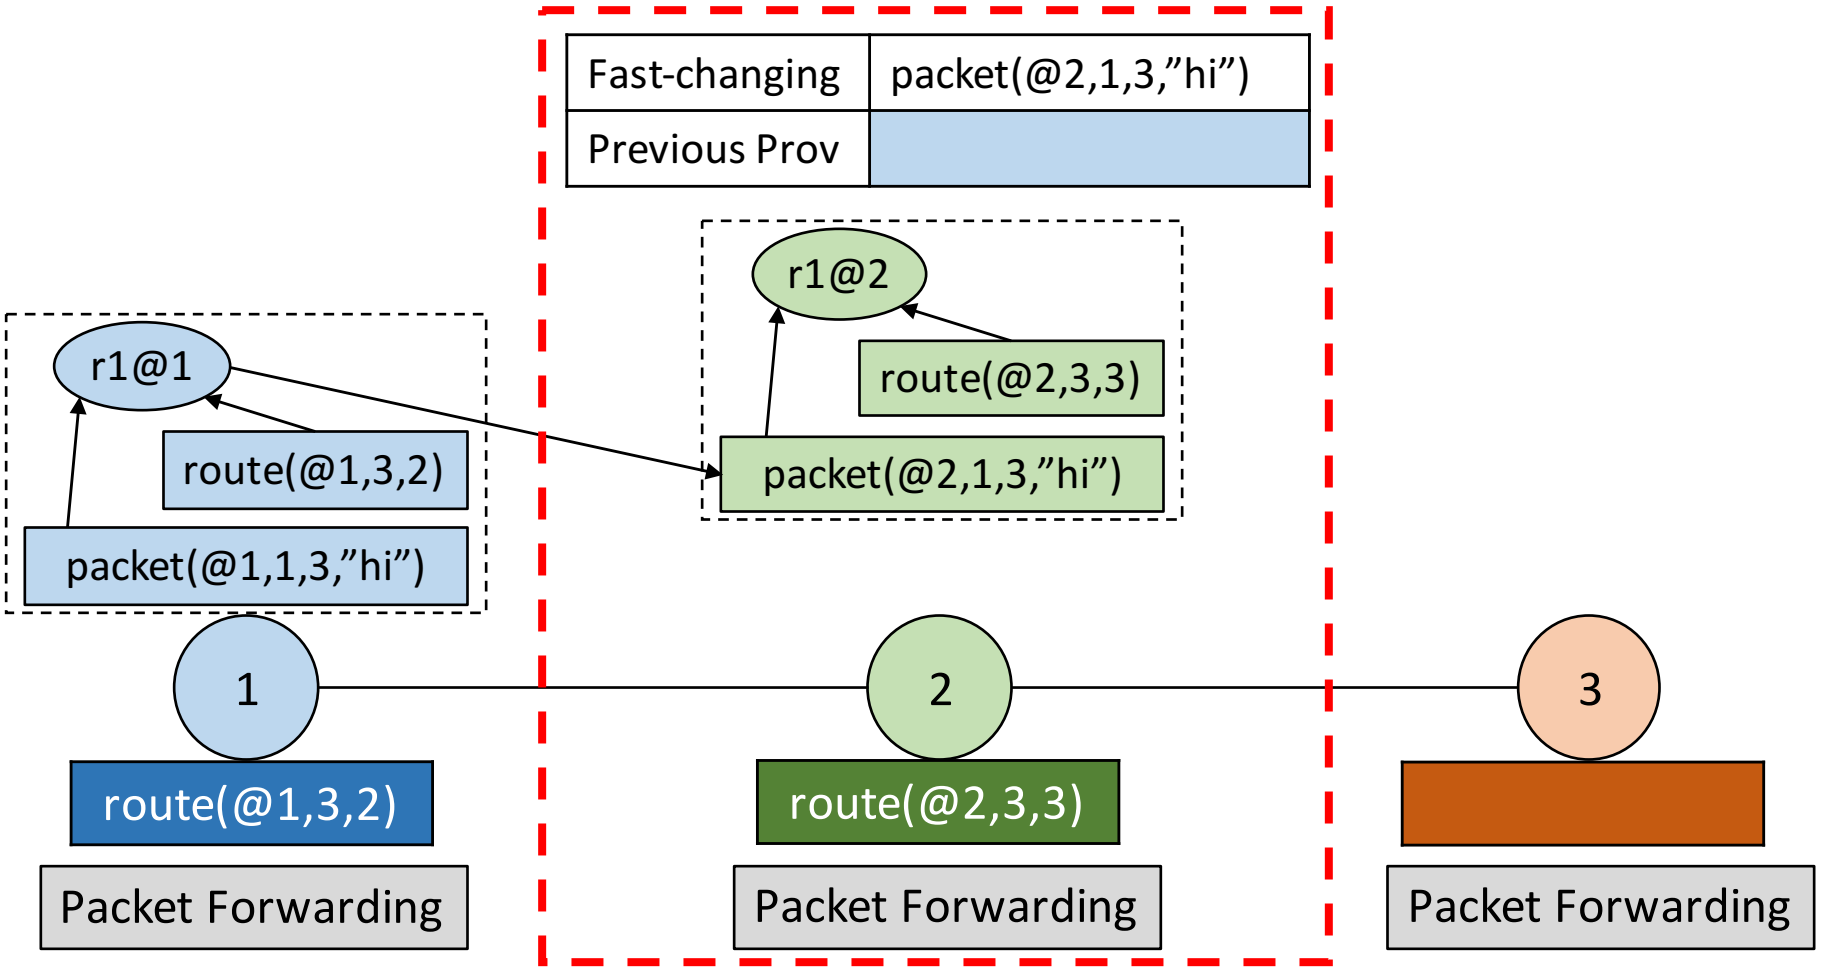

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

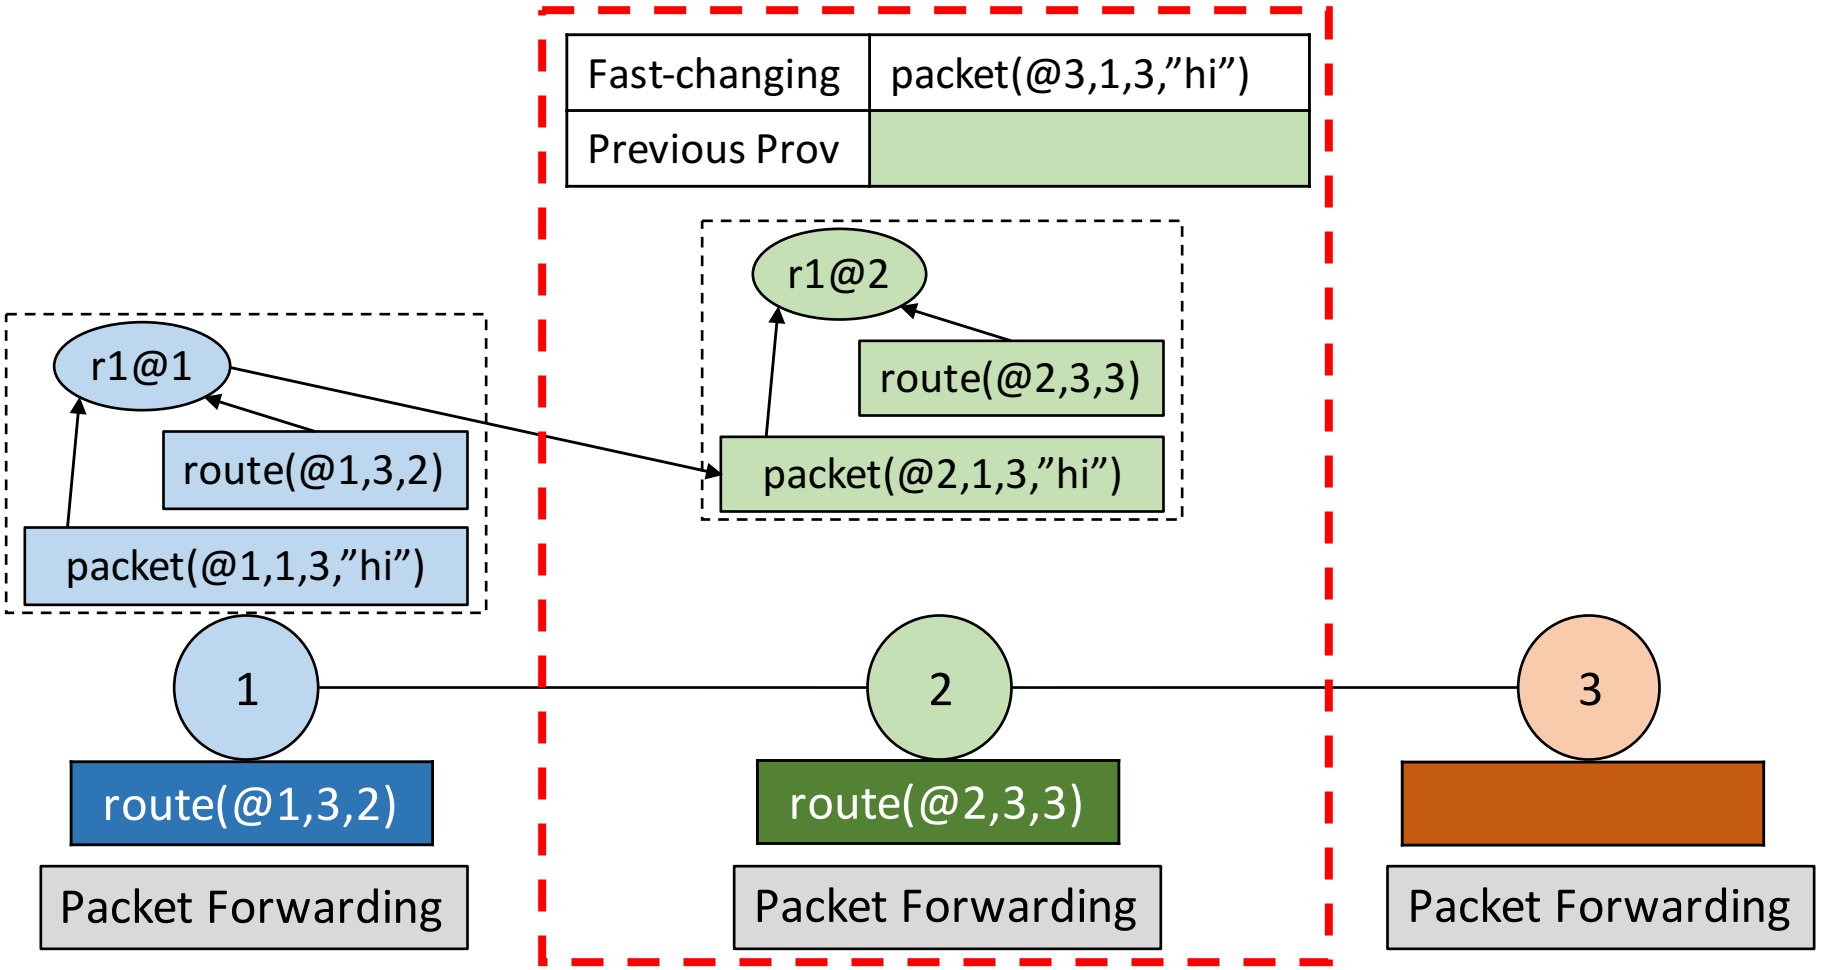

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

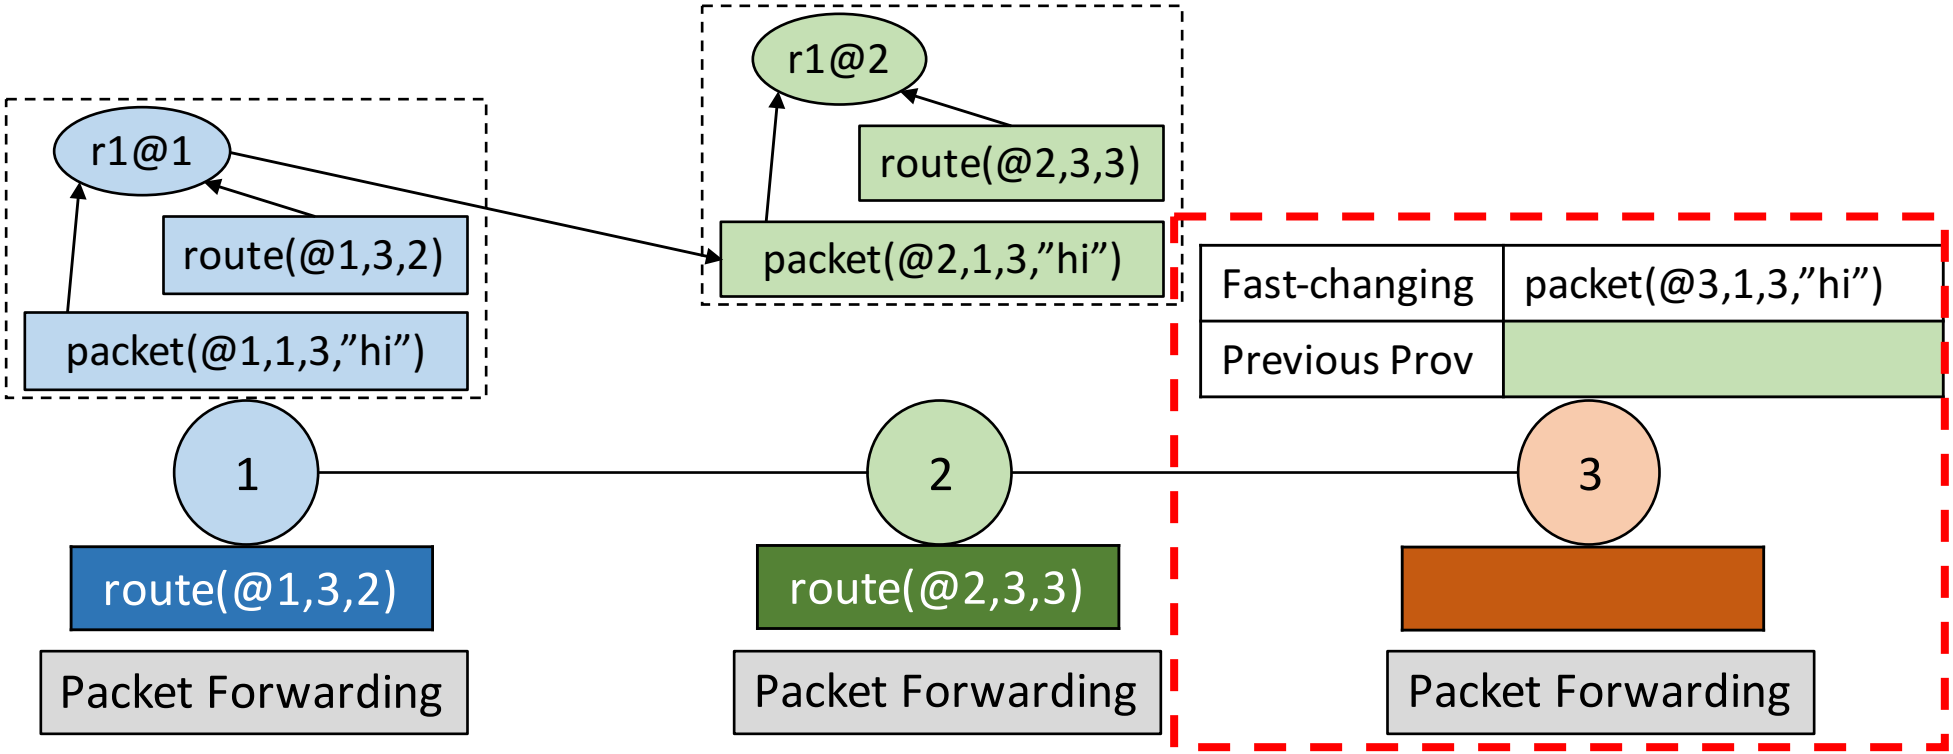

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

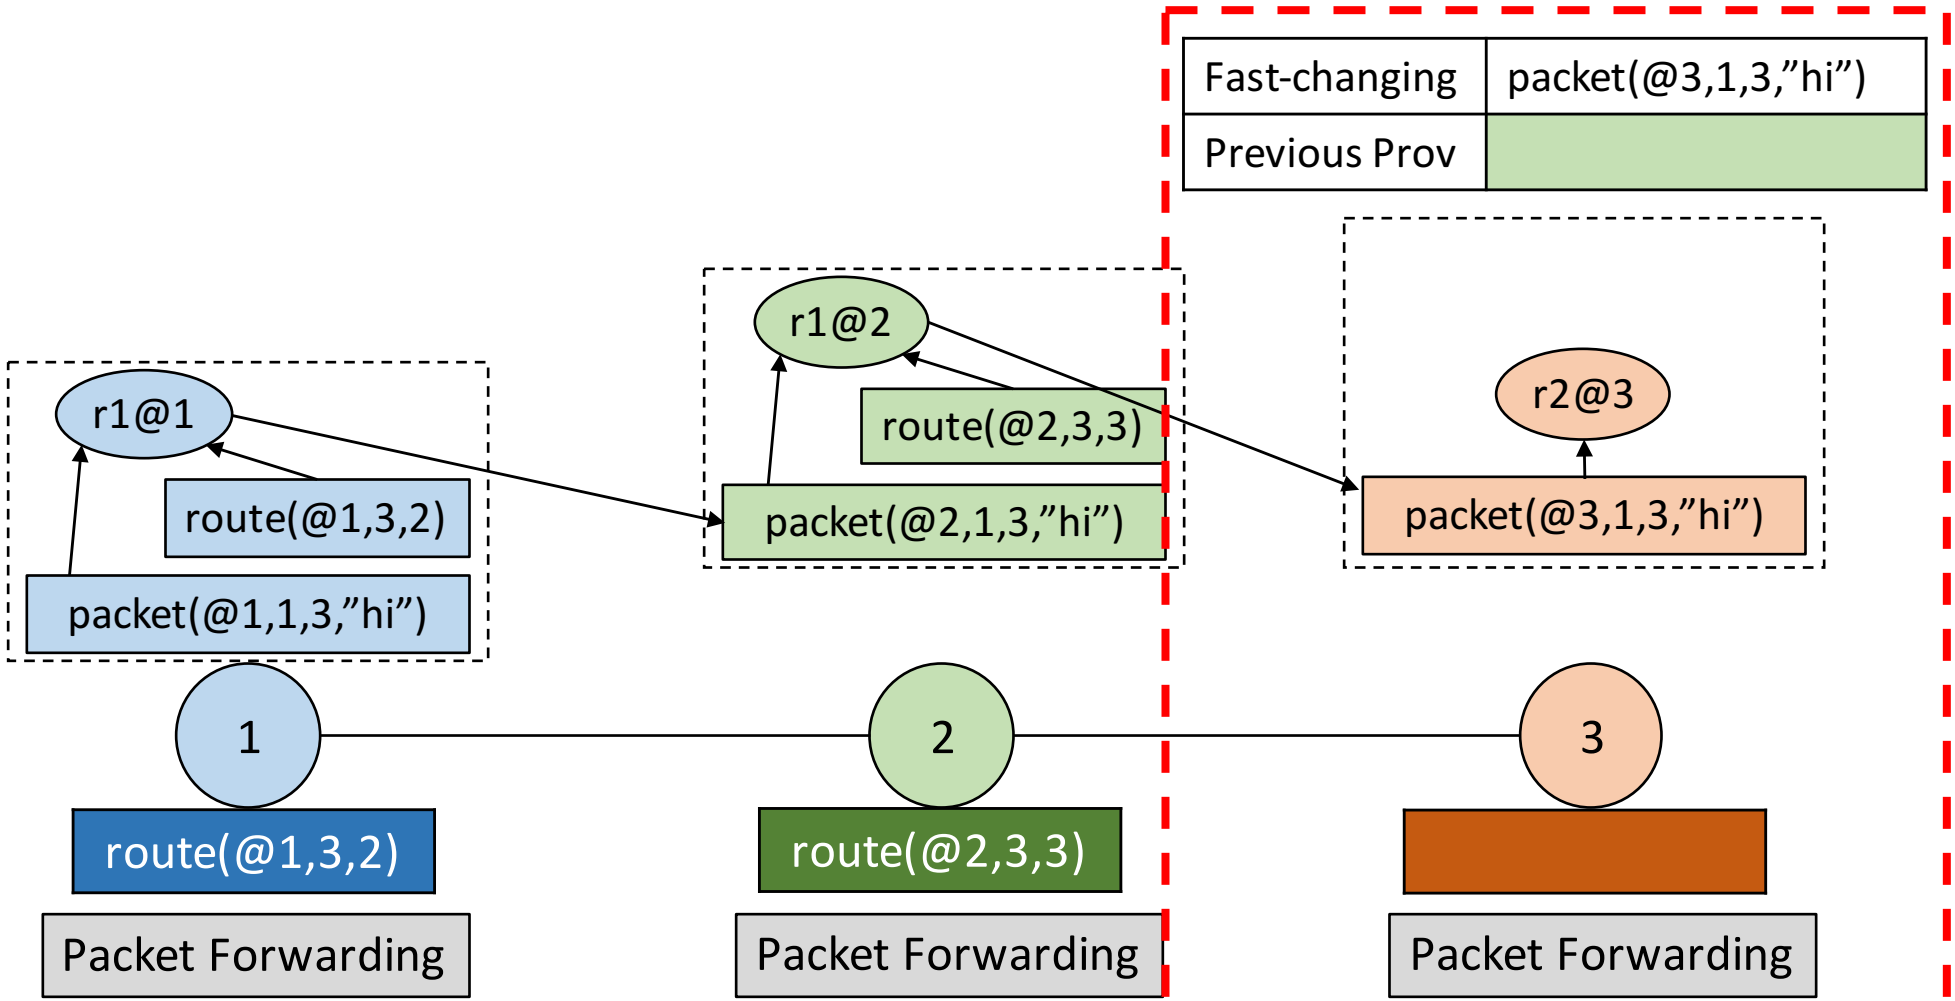

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

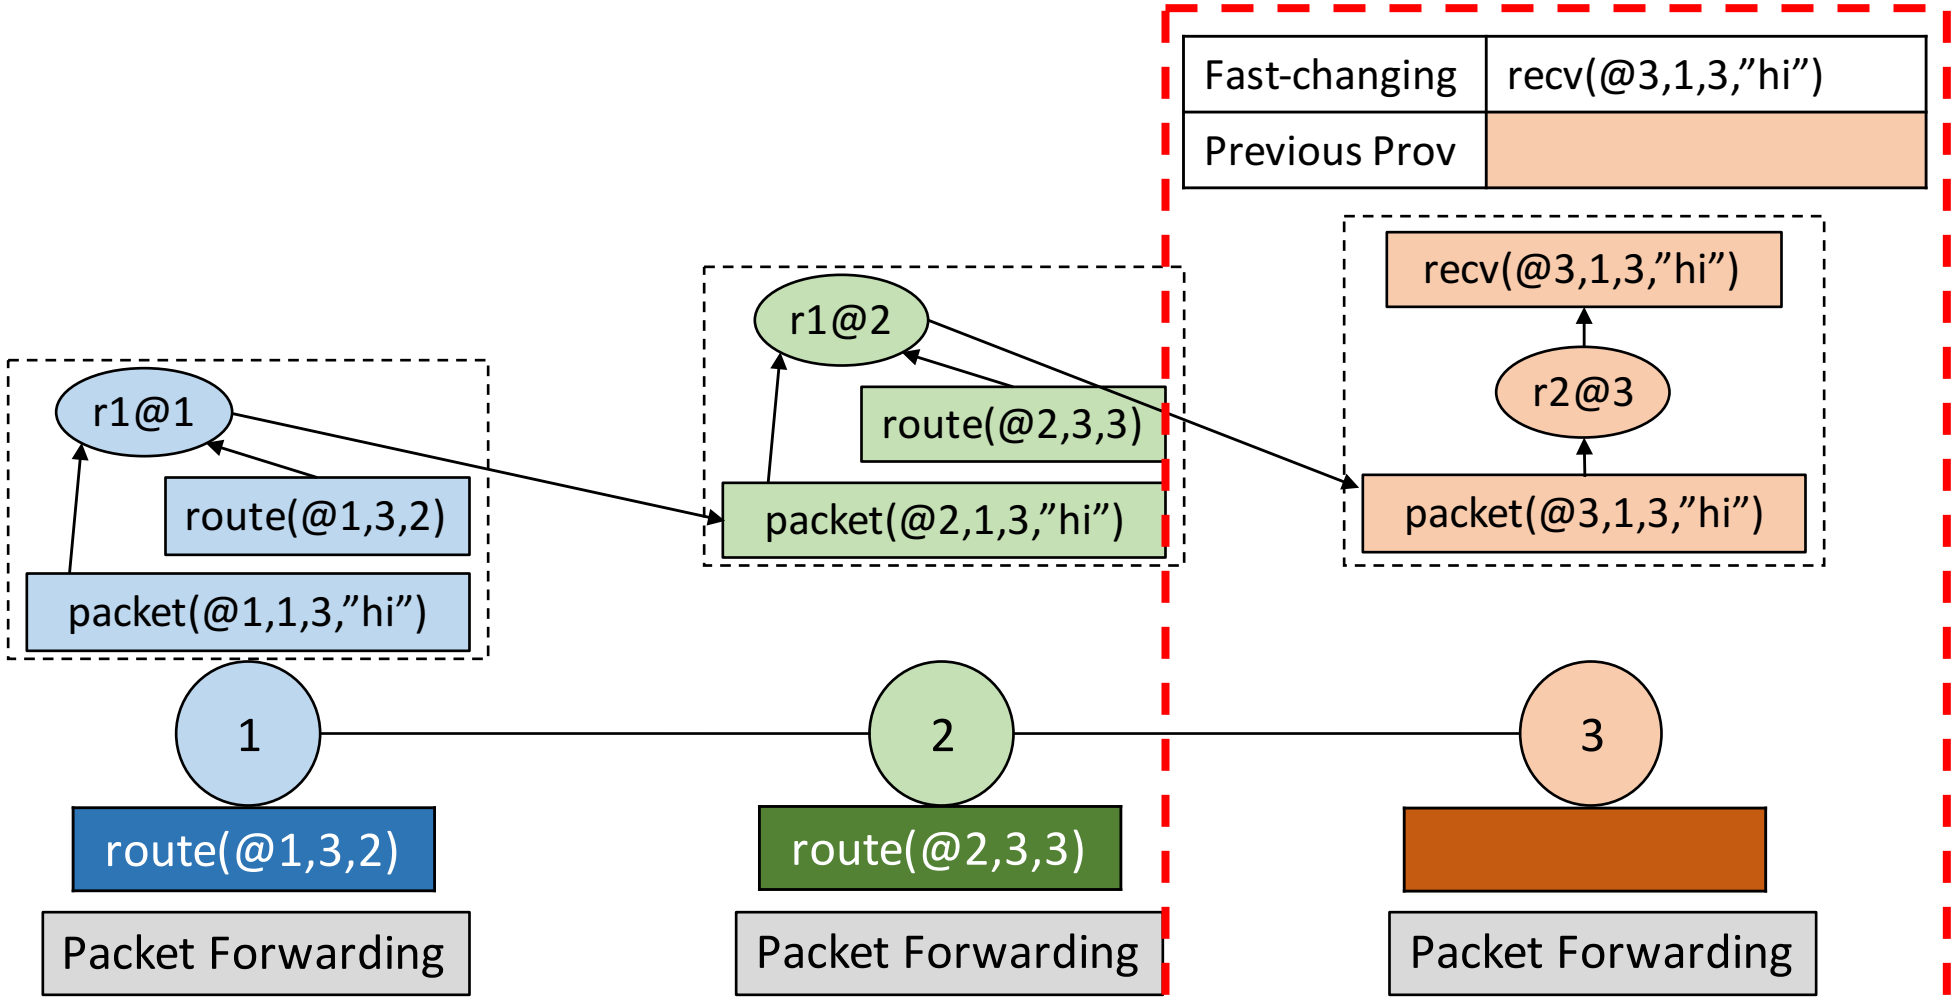

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

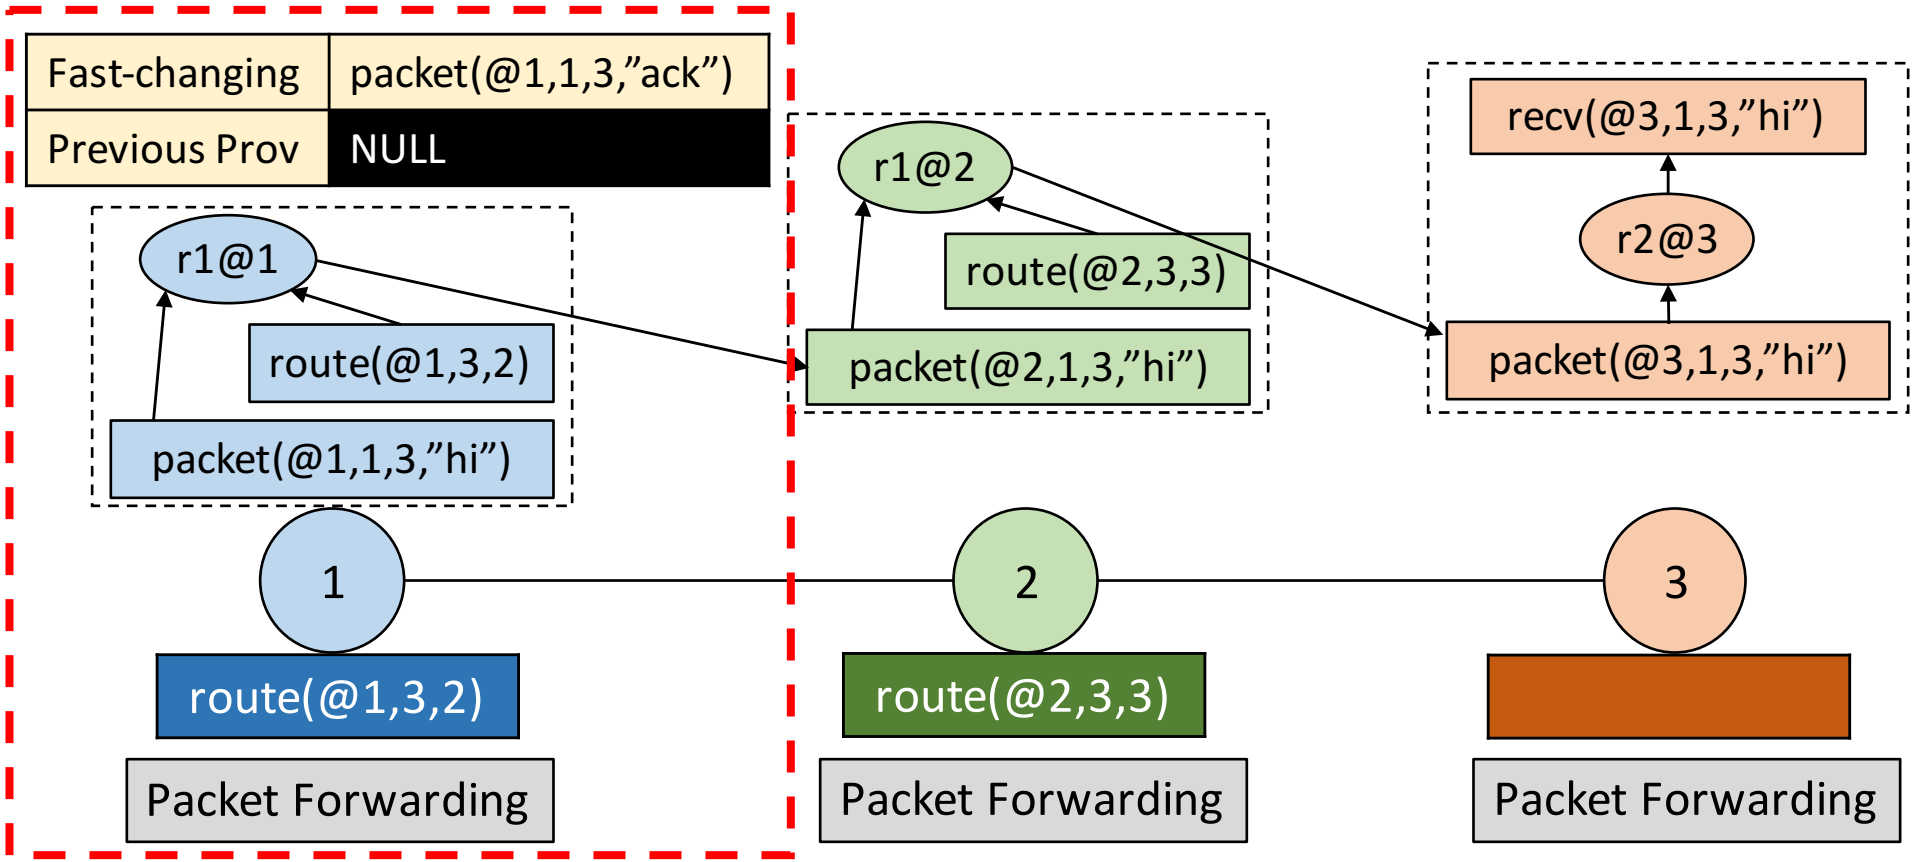

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

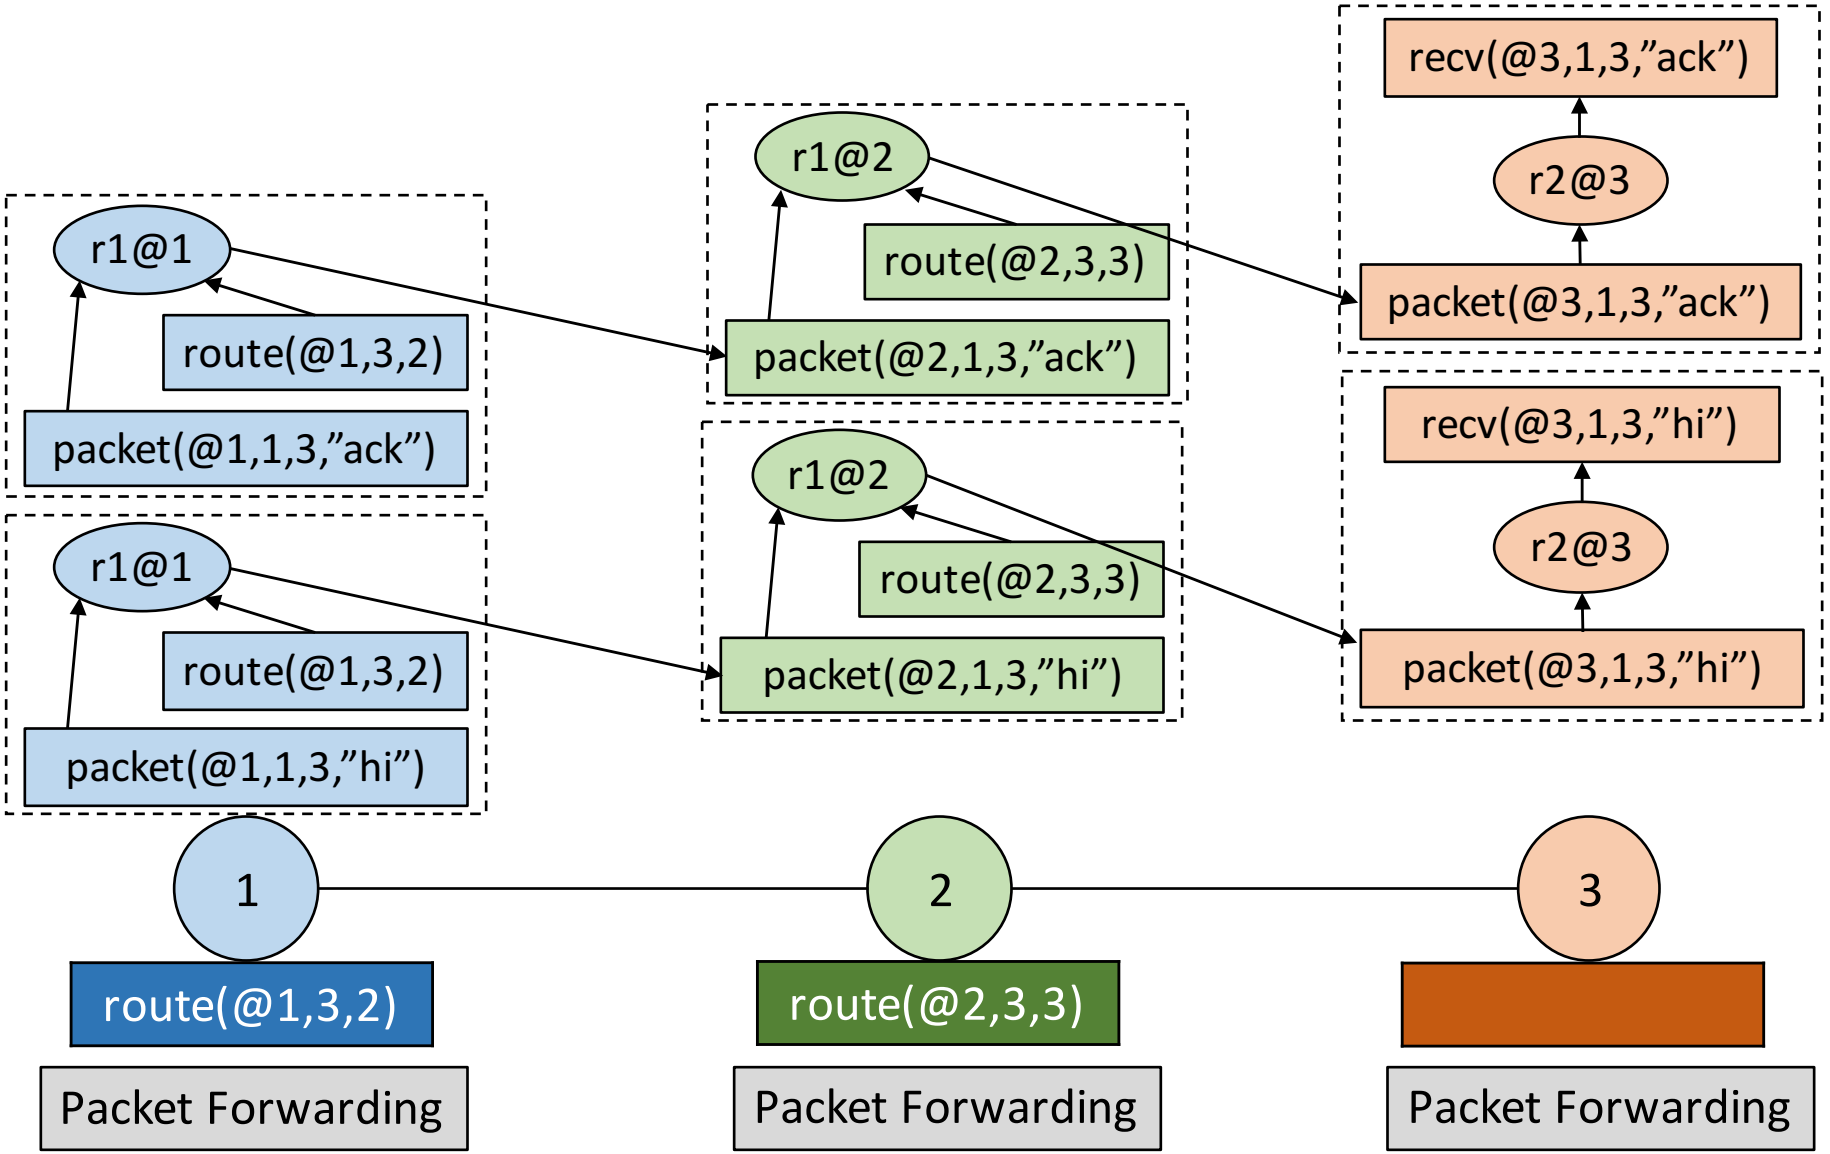

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Packet Forwarding

r1 packet(@Neigh,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), route(@Loc,Dst,Neigh).

r2 recv(@Loc,Src,Dst,Payload) :- packet(@Loc,Src,Dst,Payload), Dst==Loc.

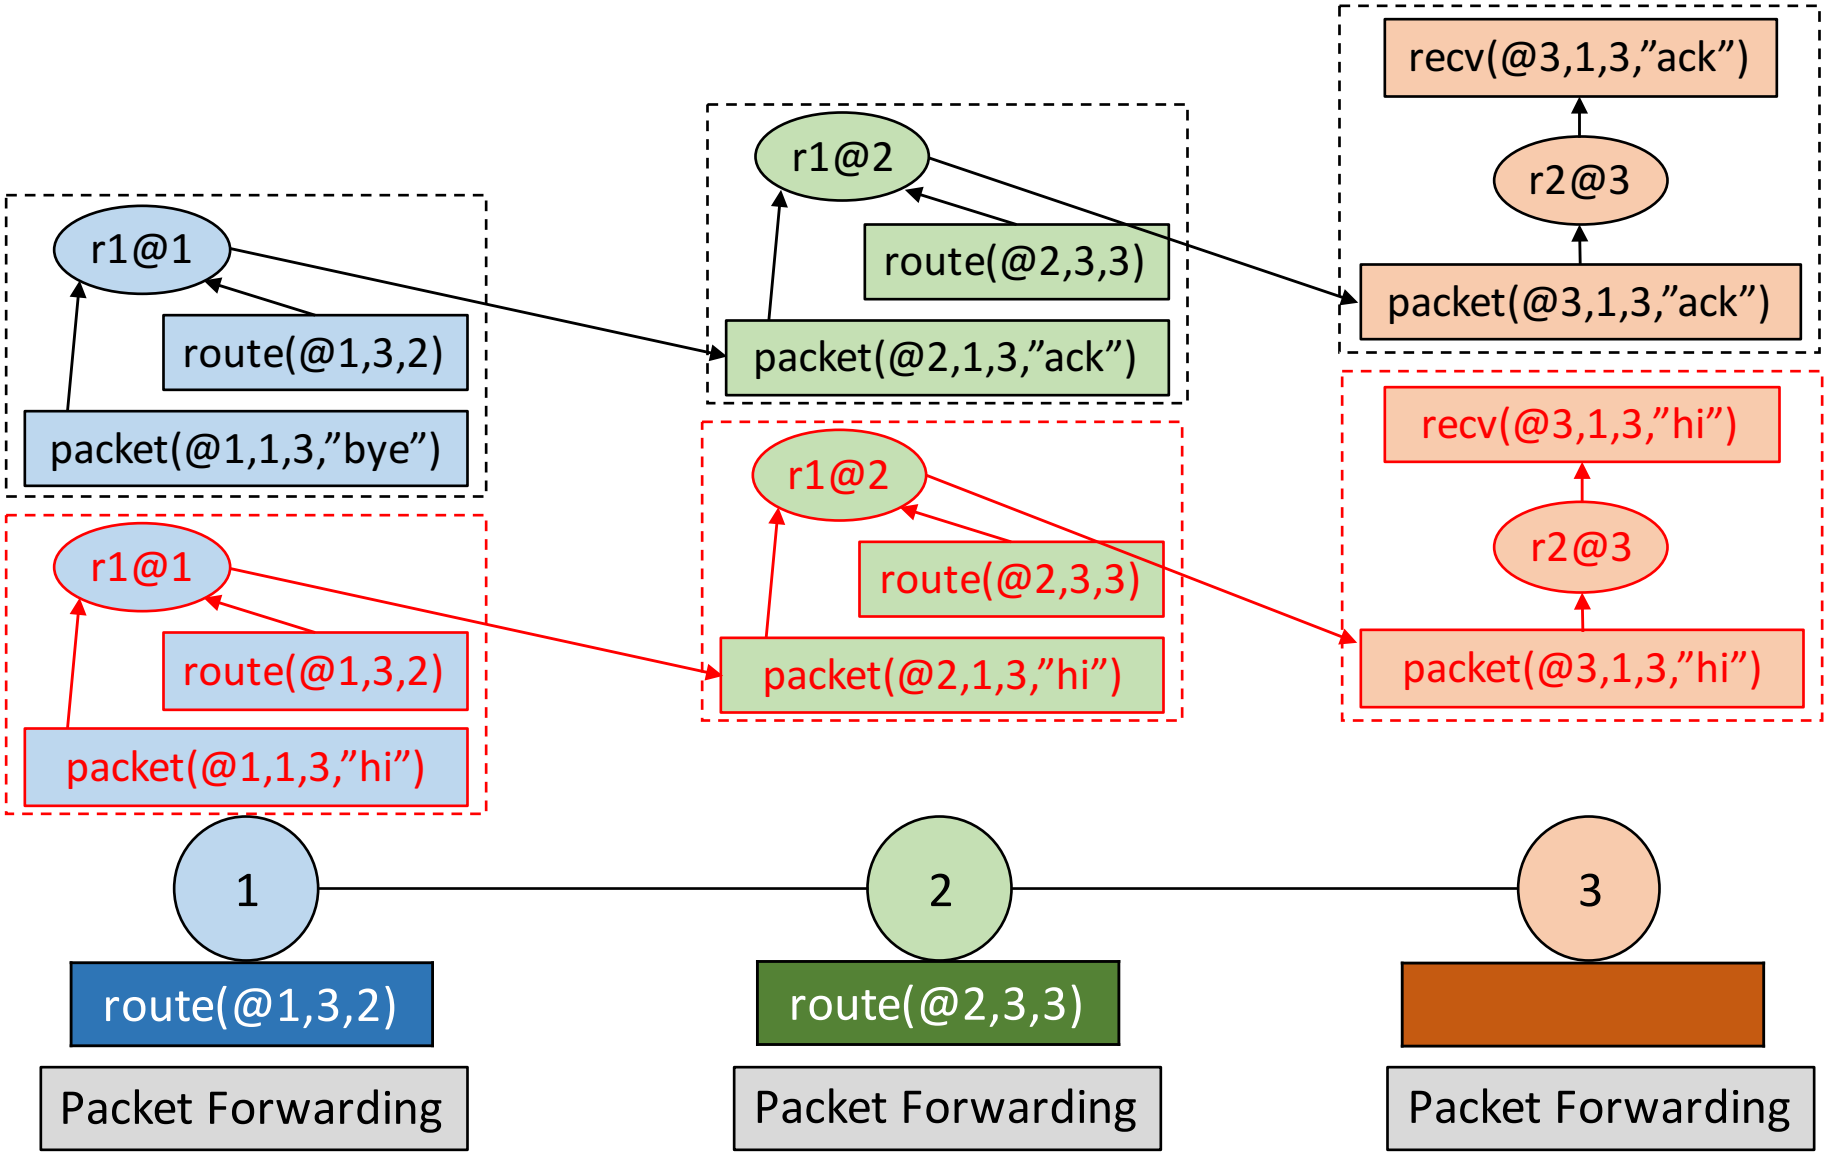

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

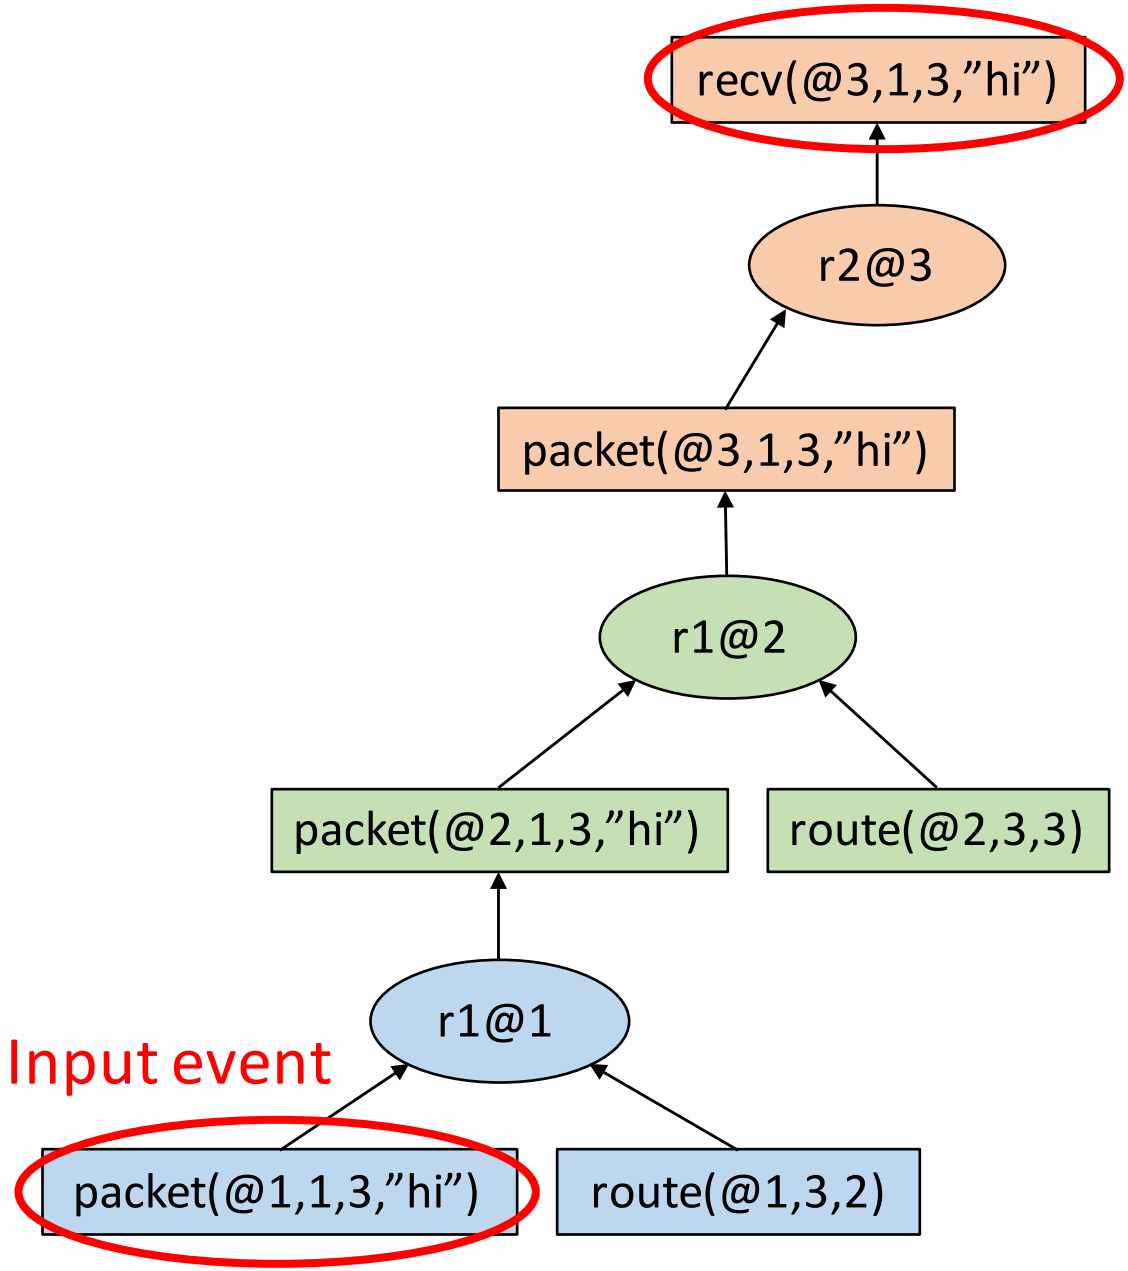

Provenance of a derived tuple can be represented as a **Provenance Tree**

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale*. SIGMOD'10

Each tuple of the output relation *recv* has a provenance tree representing its derivation

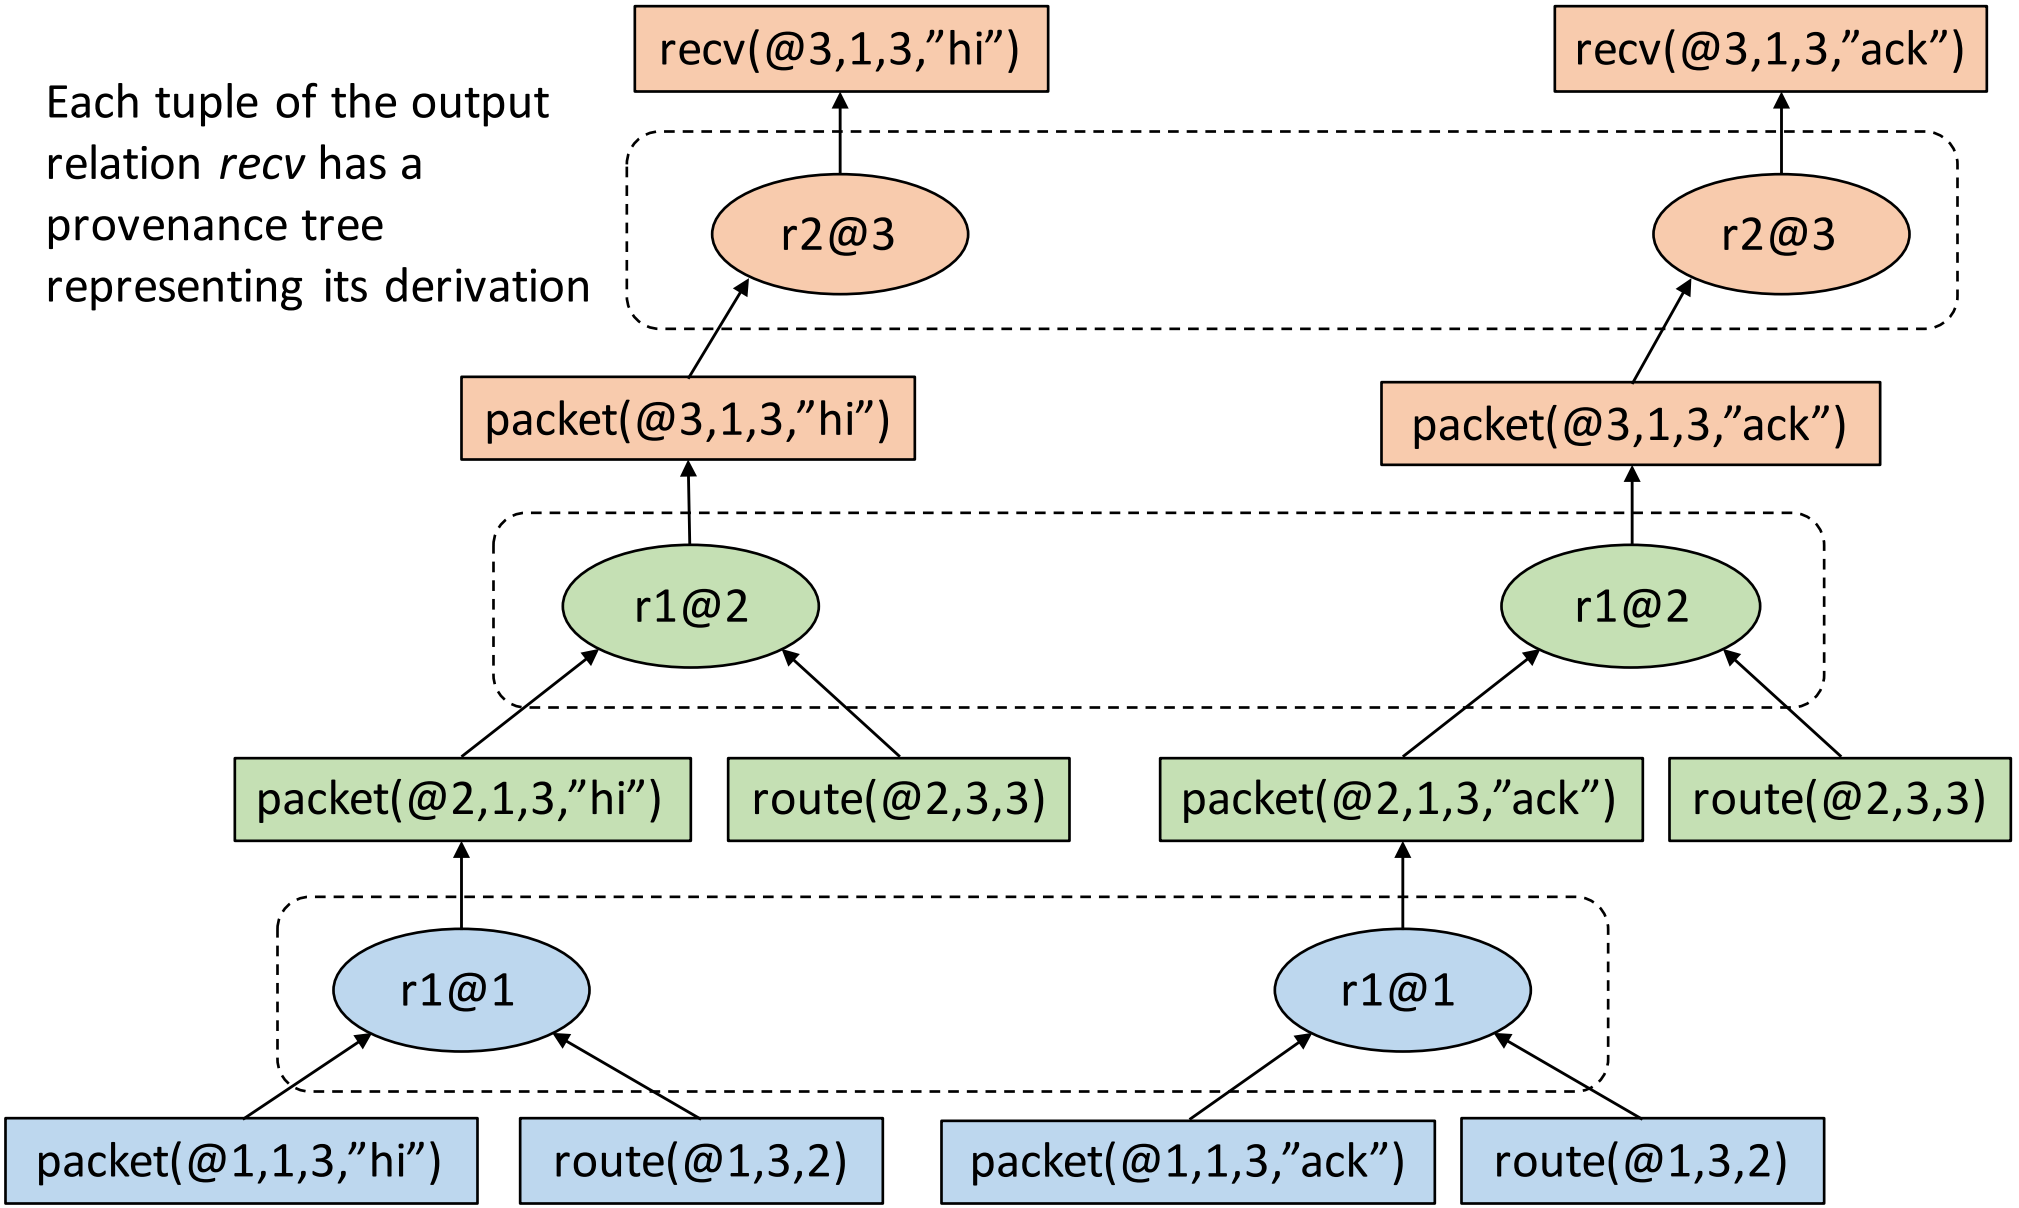

Roadmap

- Background
- ***Key insights***
- Our compression scheme
- Conclusion

Insight #1: Remove Redundancy

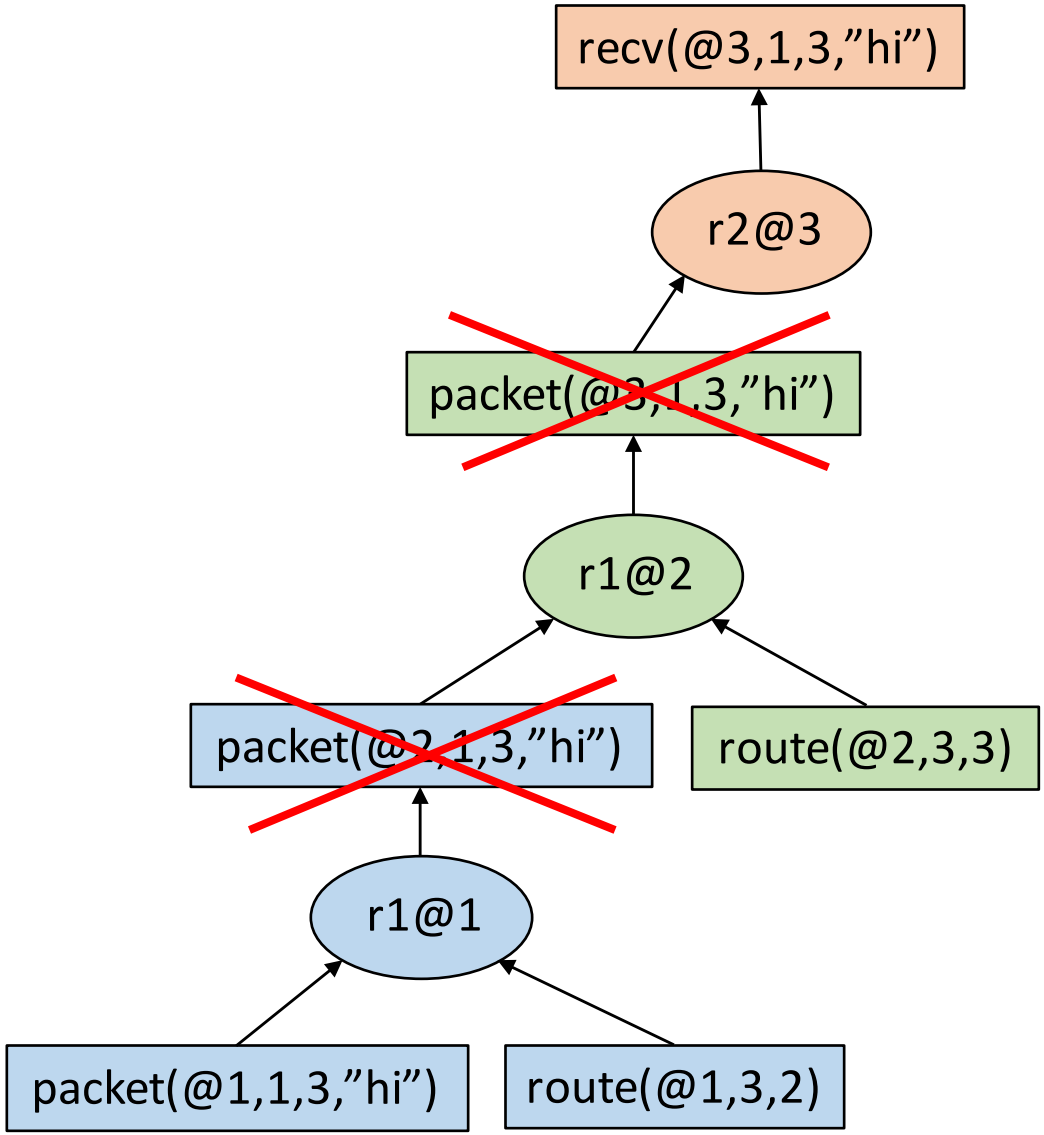

Insight #1: Remove Redundancy

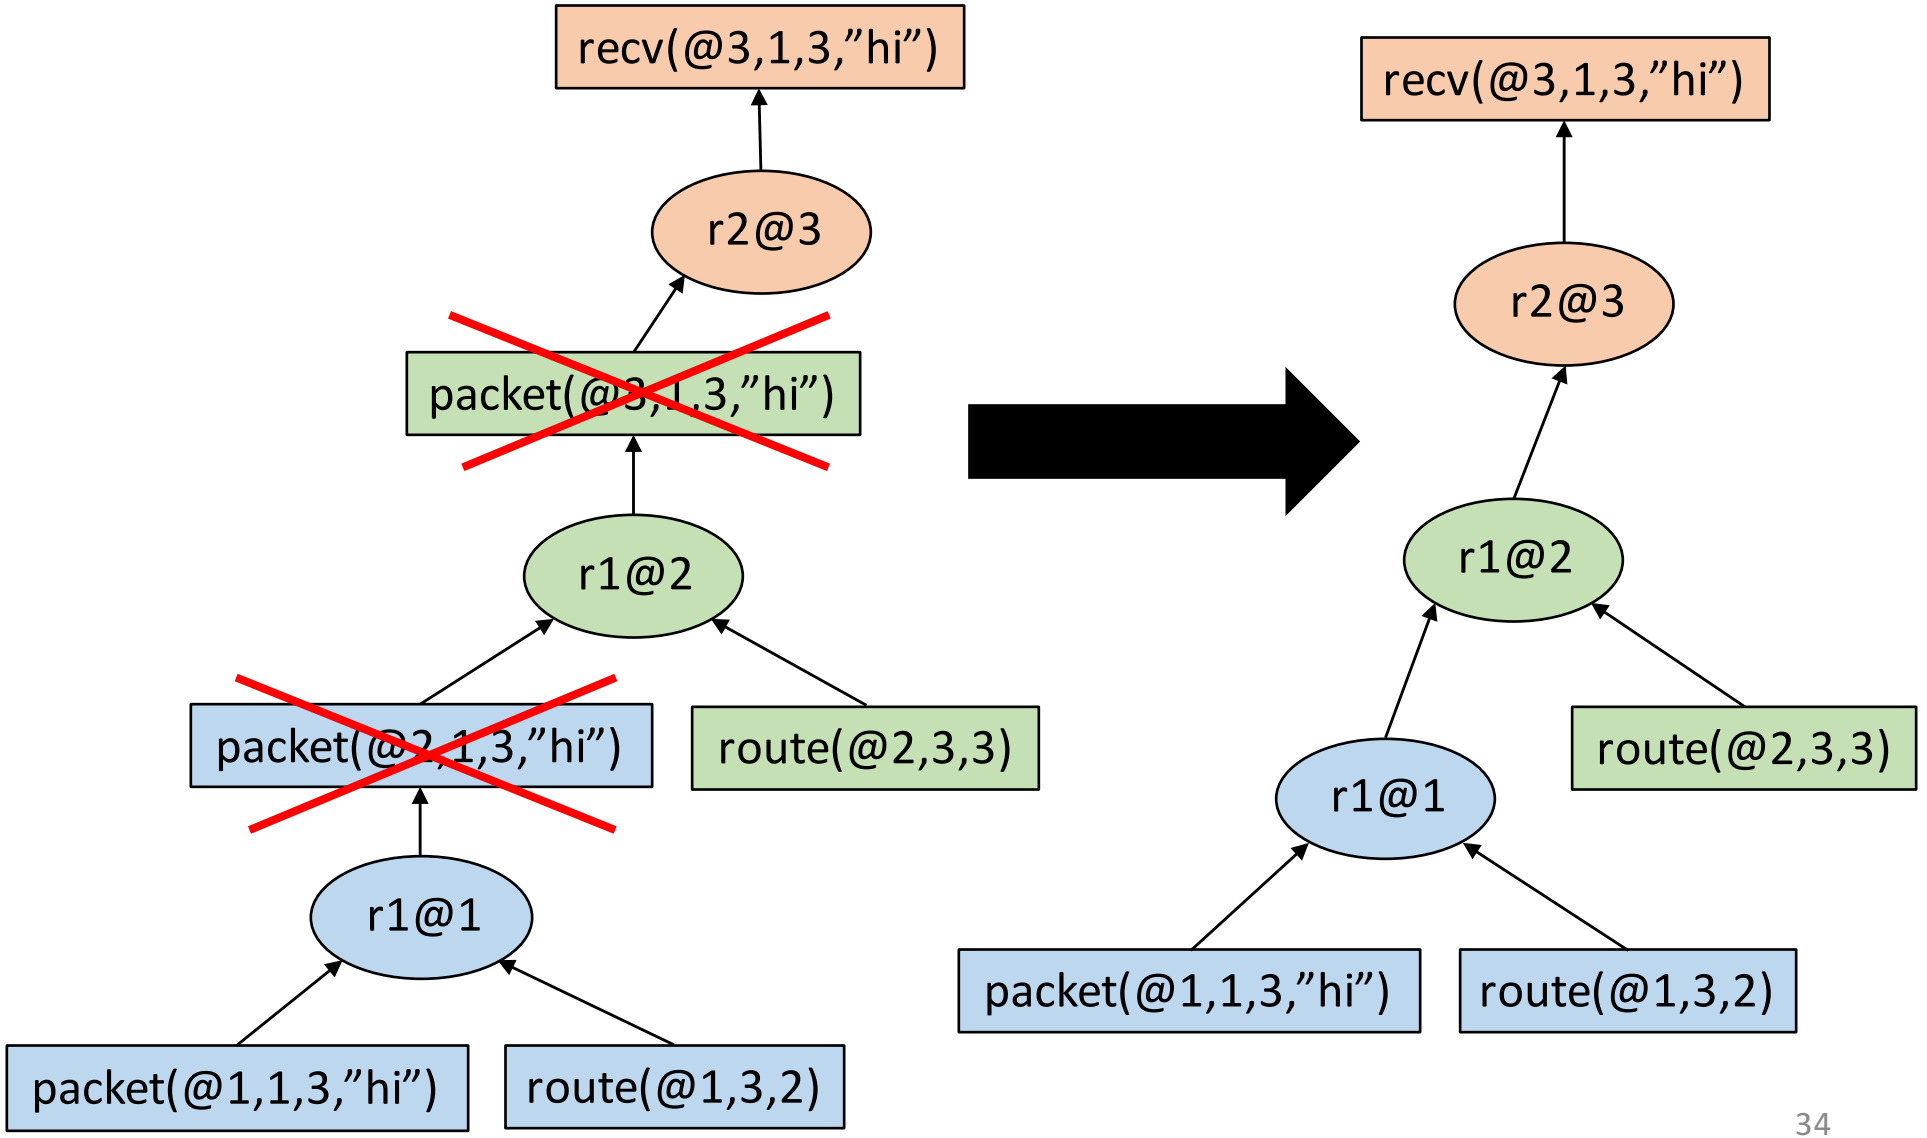

Insight #2:
Different packets may follow the same path

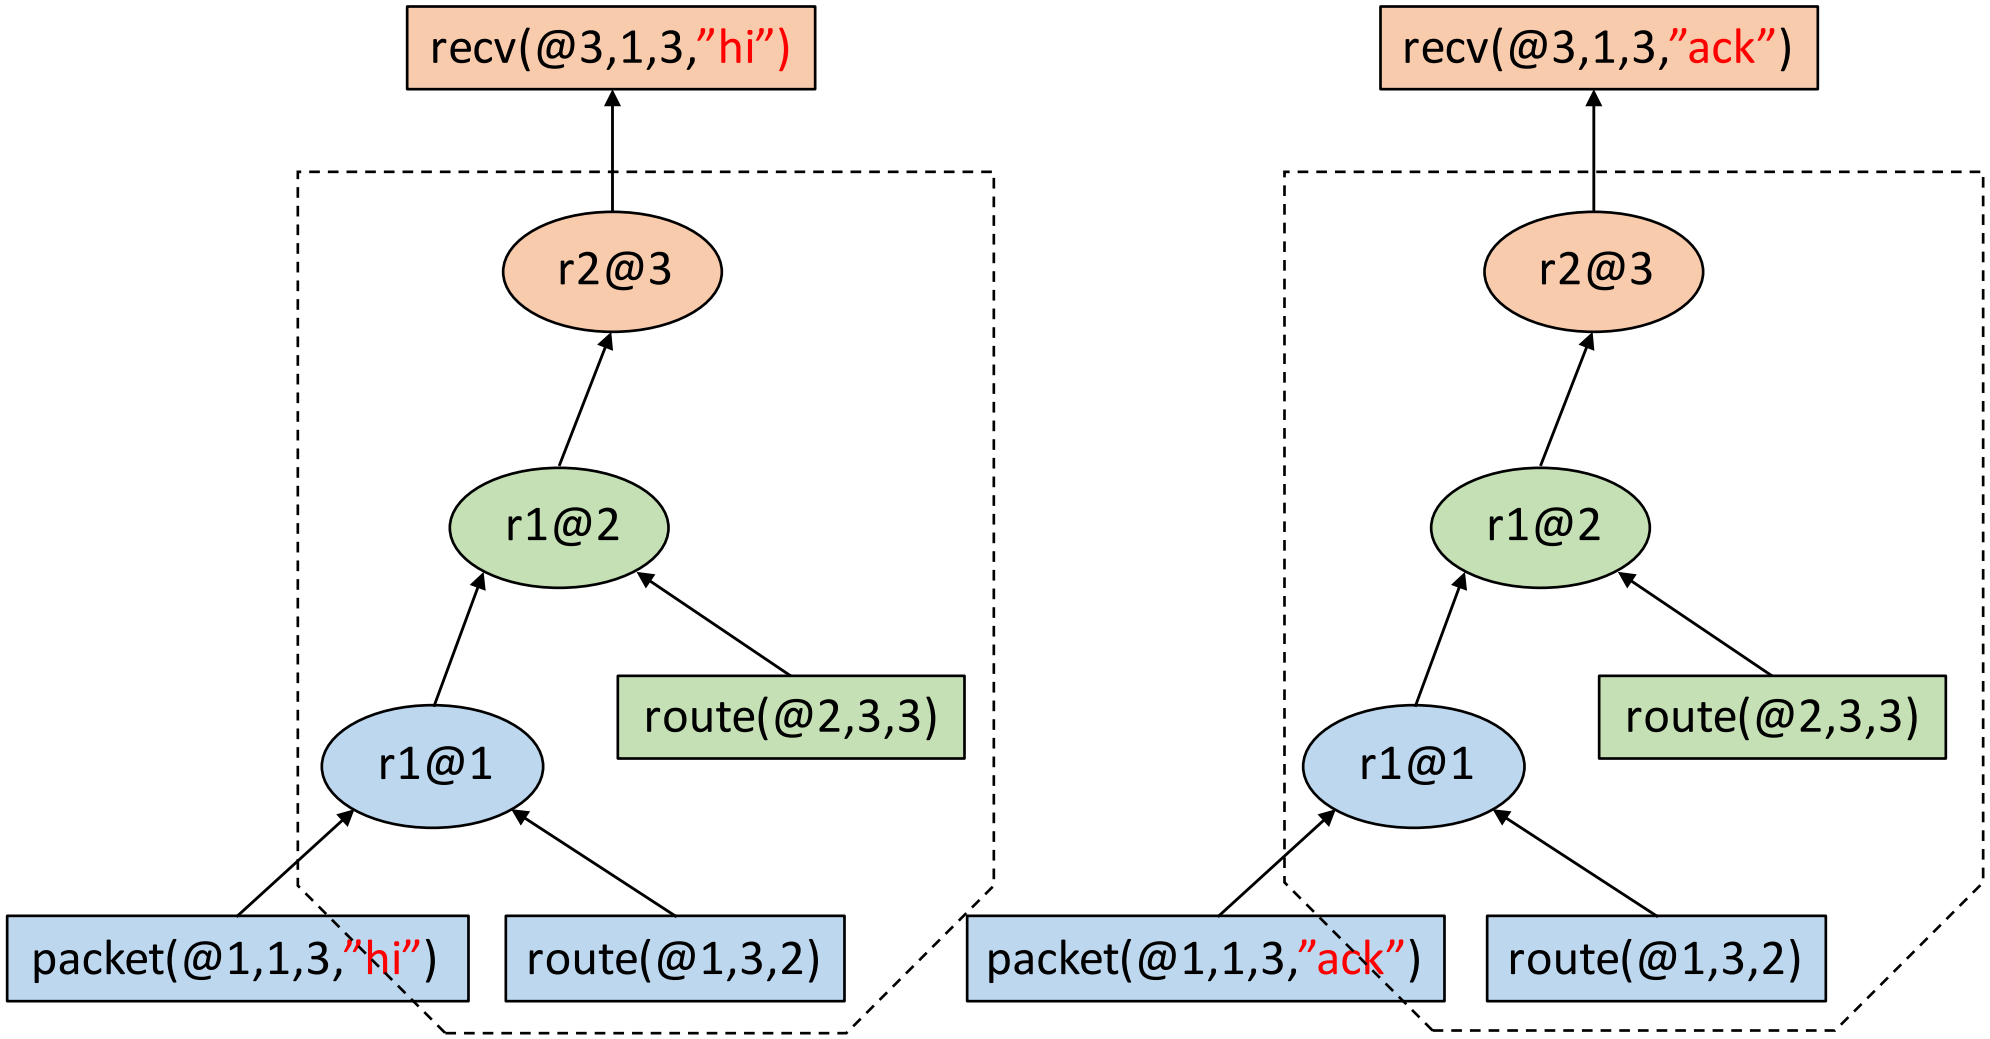

Simplifying Assumption #1

One fast-changing relation per NDLog rule

<result> :- <condition$_1$>, <condition$_2$>, … , <condition$_N$>

Fast-changing

Slow-changing

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

Simplifying Assumption #2

The set of slow changing tuples is ***constant***

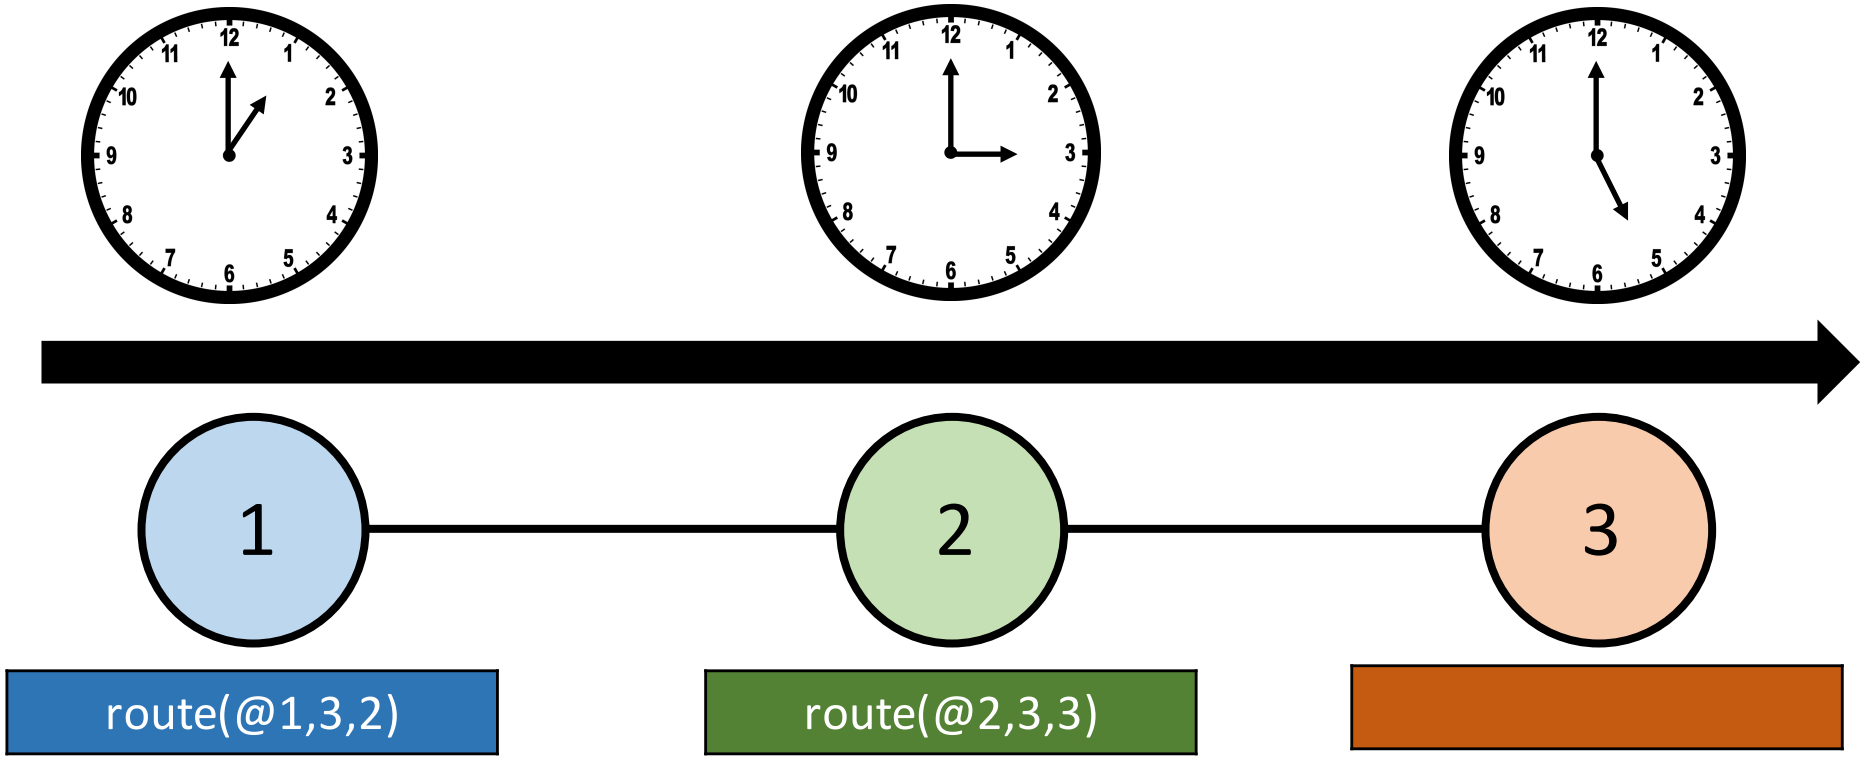

Roadmap

- Background
- Key insights
- ***Our compression scheme***
- Conclusion

Key Idea: Group provenance trees into ***Equivalence Classes***

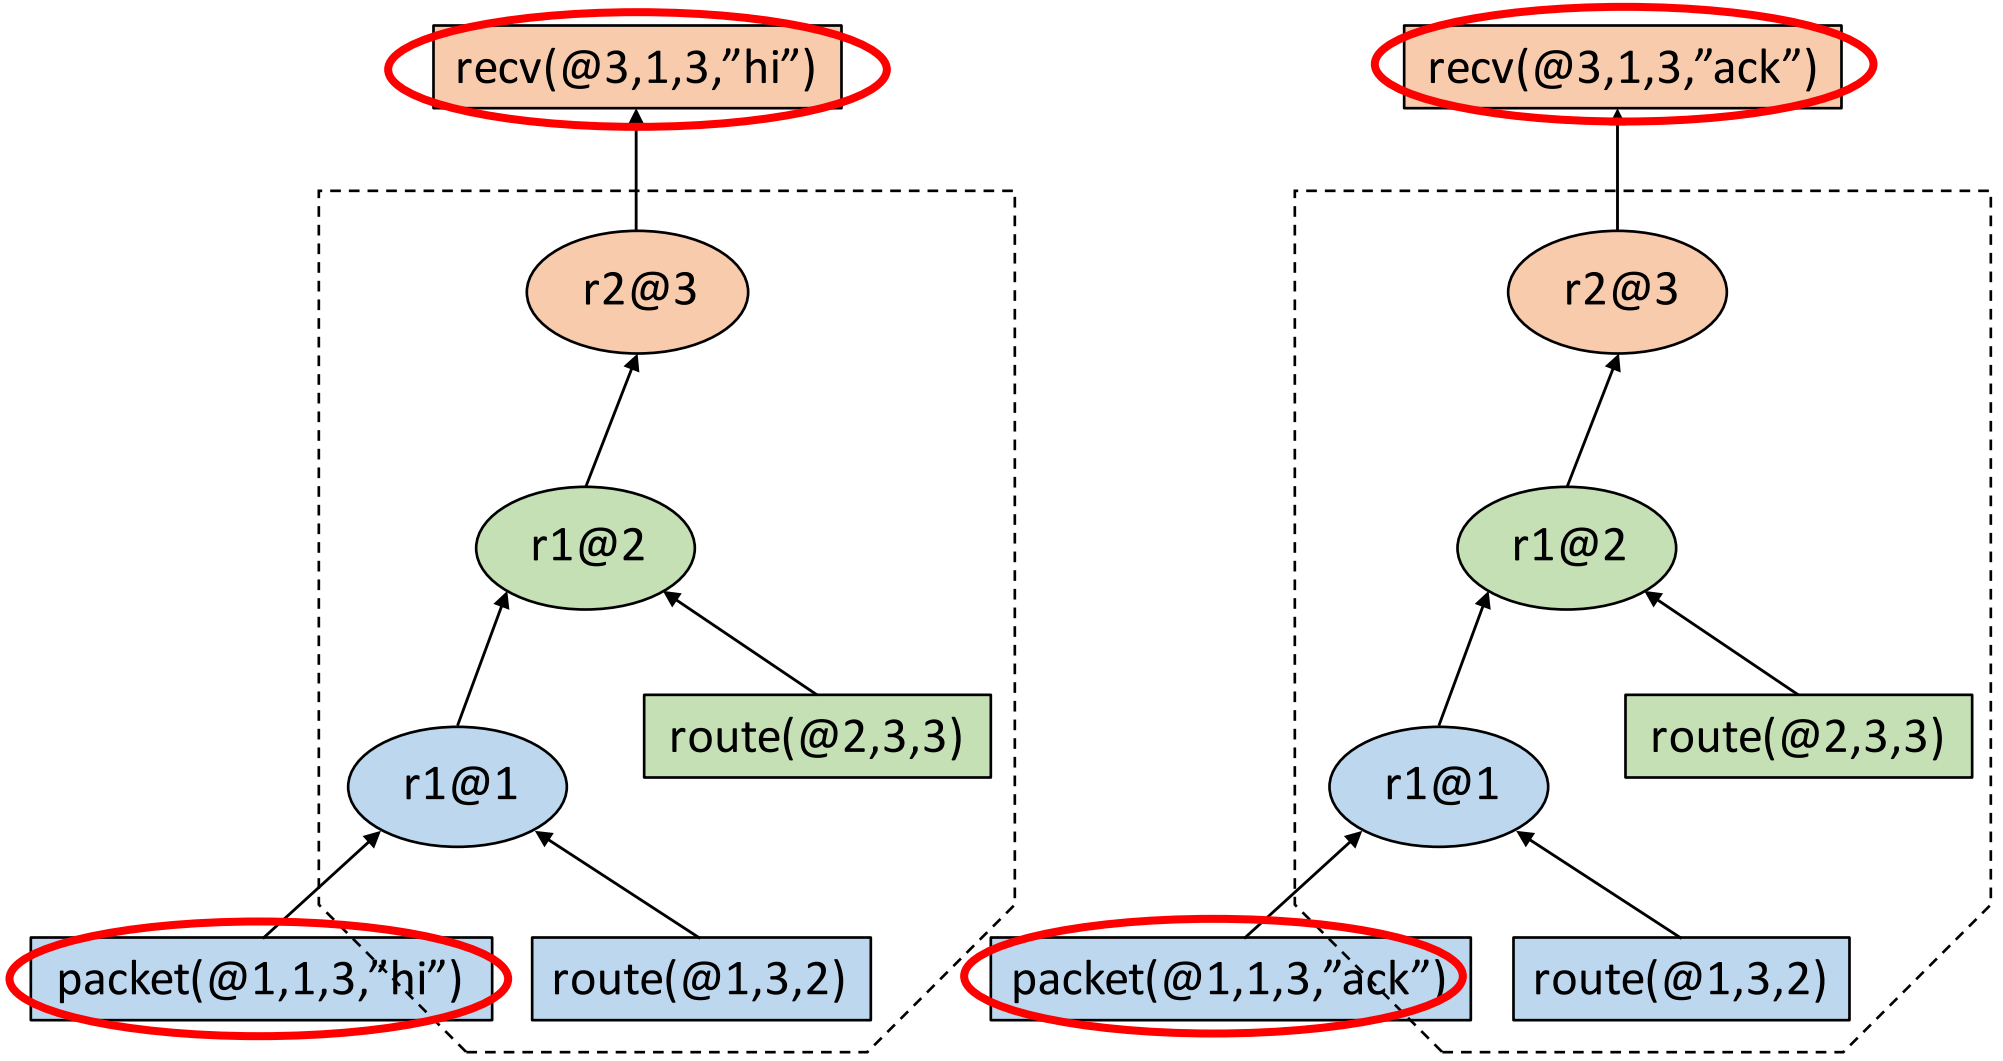

Collect provenance trees in ***one location***

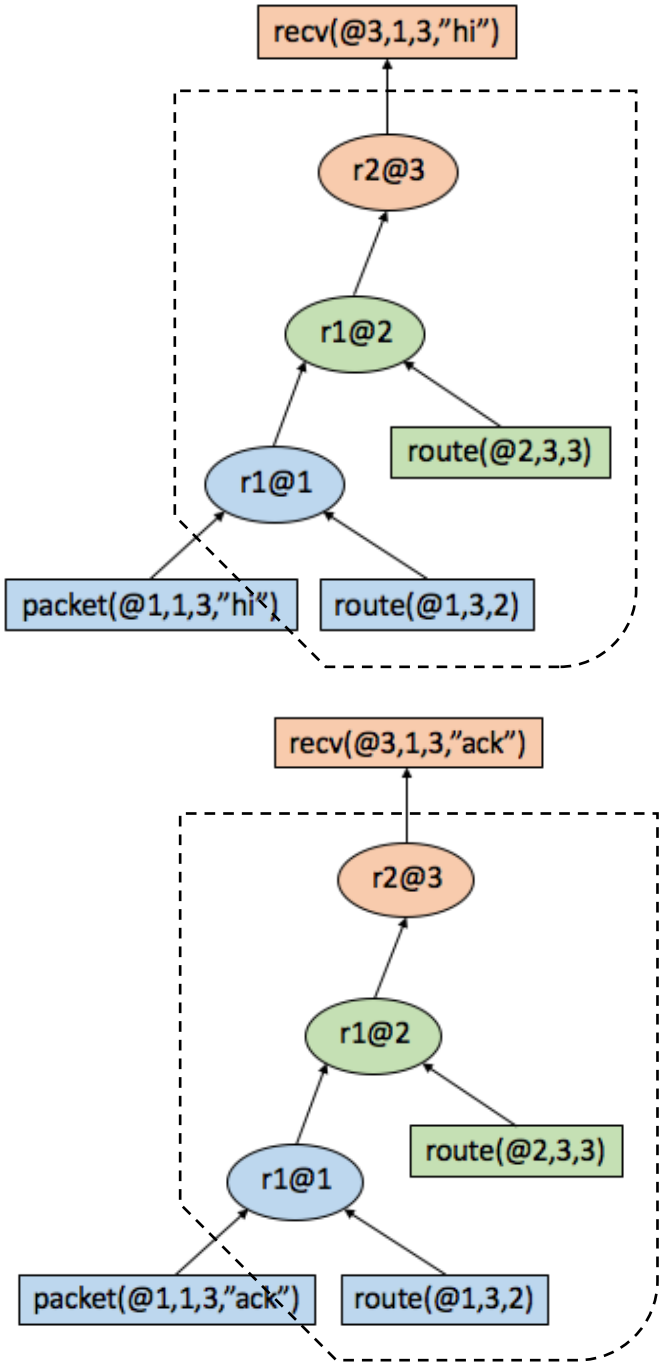

First attempt at compression

Collect provenance trees in ***one location***

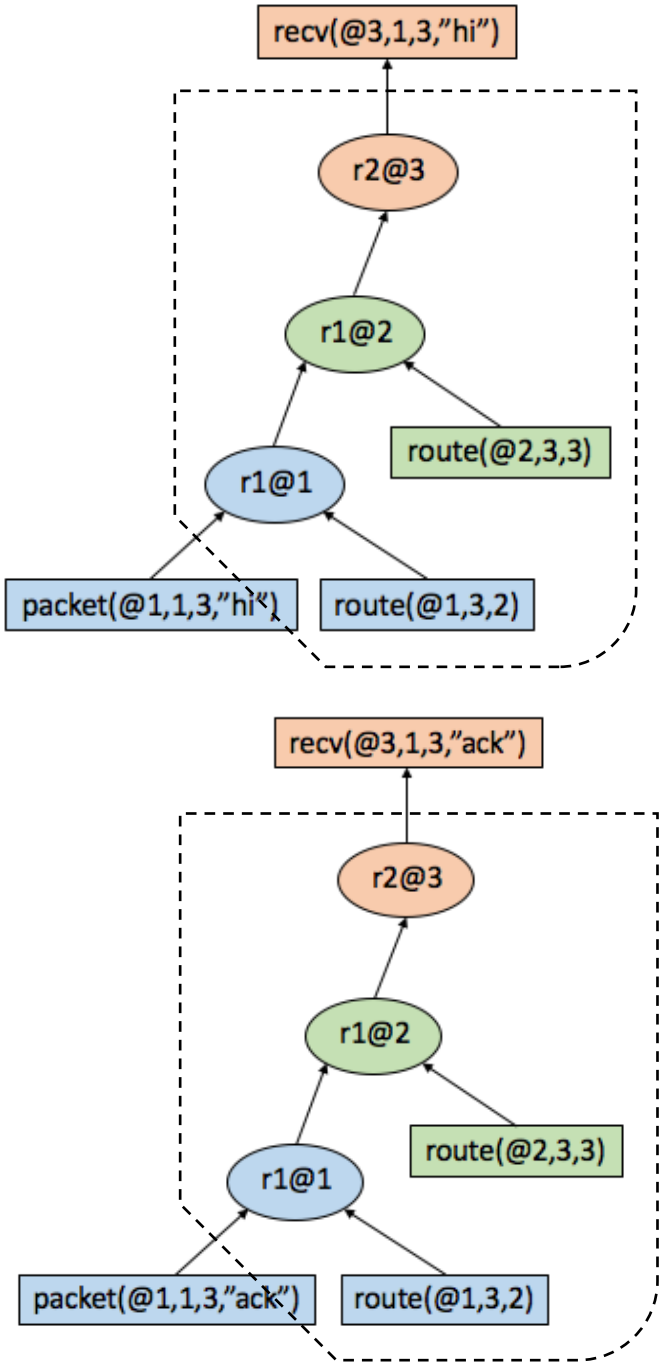

First attempt at compression

Chapman et al., *Efficient Provenance Storage*. SIGMOD'08

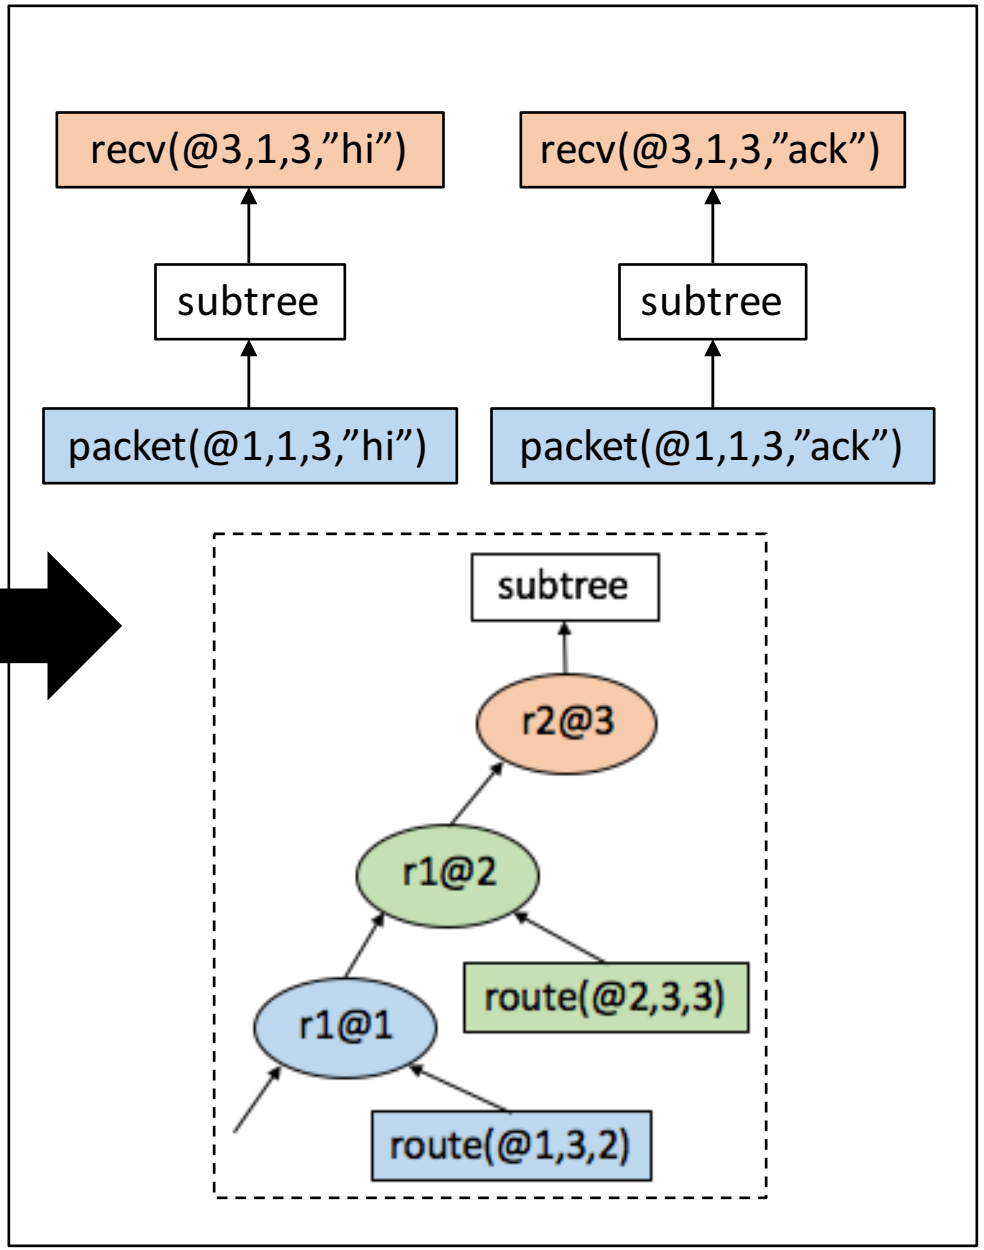

Zhou et al., *Efficient Querying & Maintenance of Network Provenance at Internet-Scale.* SIGMOD'10

Chapman et al., *Efficient Provenance Storage*. SIGMOD'08

Why not collect provenance in a centralized server?

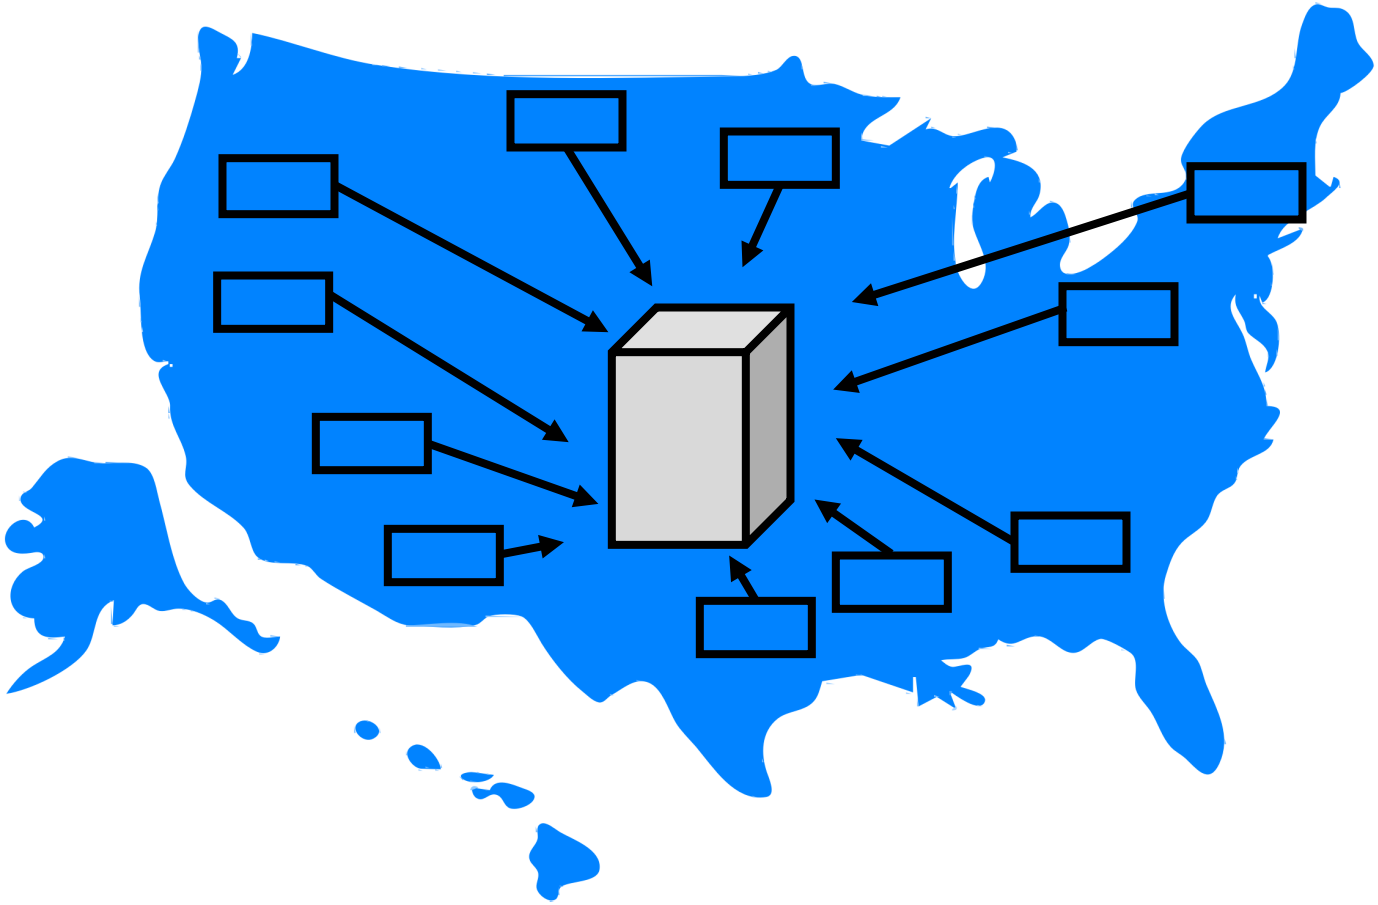

1. Bottleneck at the centralized server
2. High bandwidth utilization
3. Feasibility – entities in the network may span large geographic locations

Our Compression Scheme

- Identifies how to share storage ***before*** program execution.
- Stores compressed provenances in a ***distributed*** setting ***at runtime***.

Workflow for our Compression Scheme

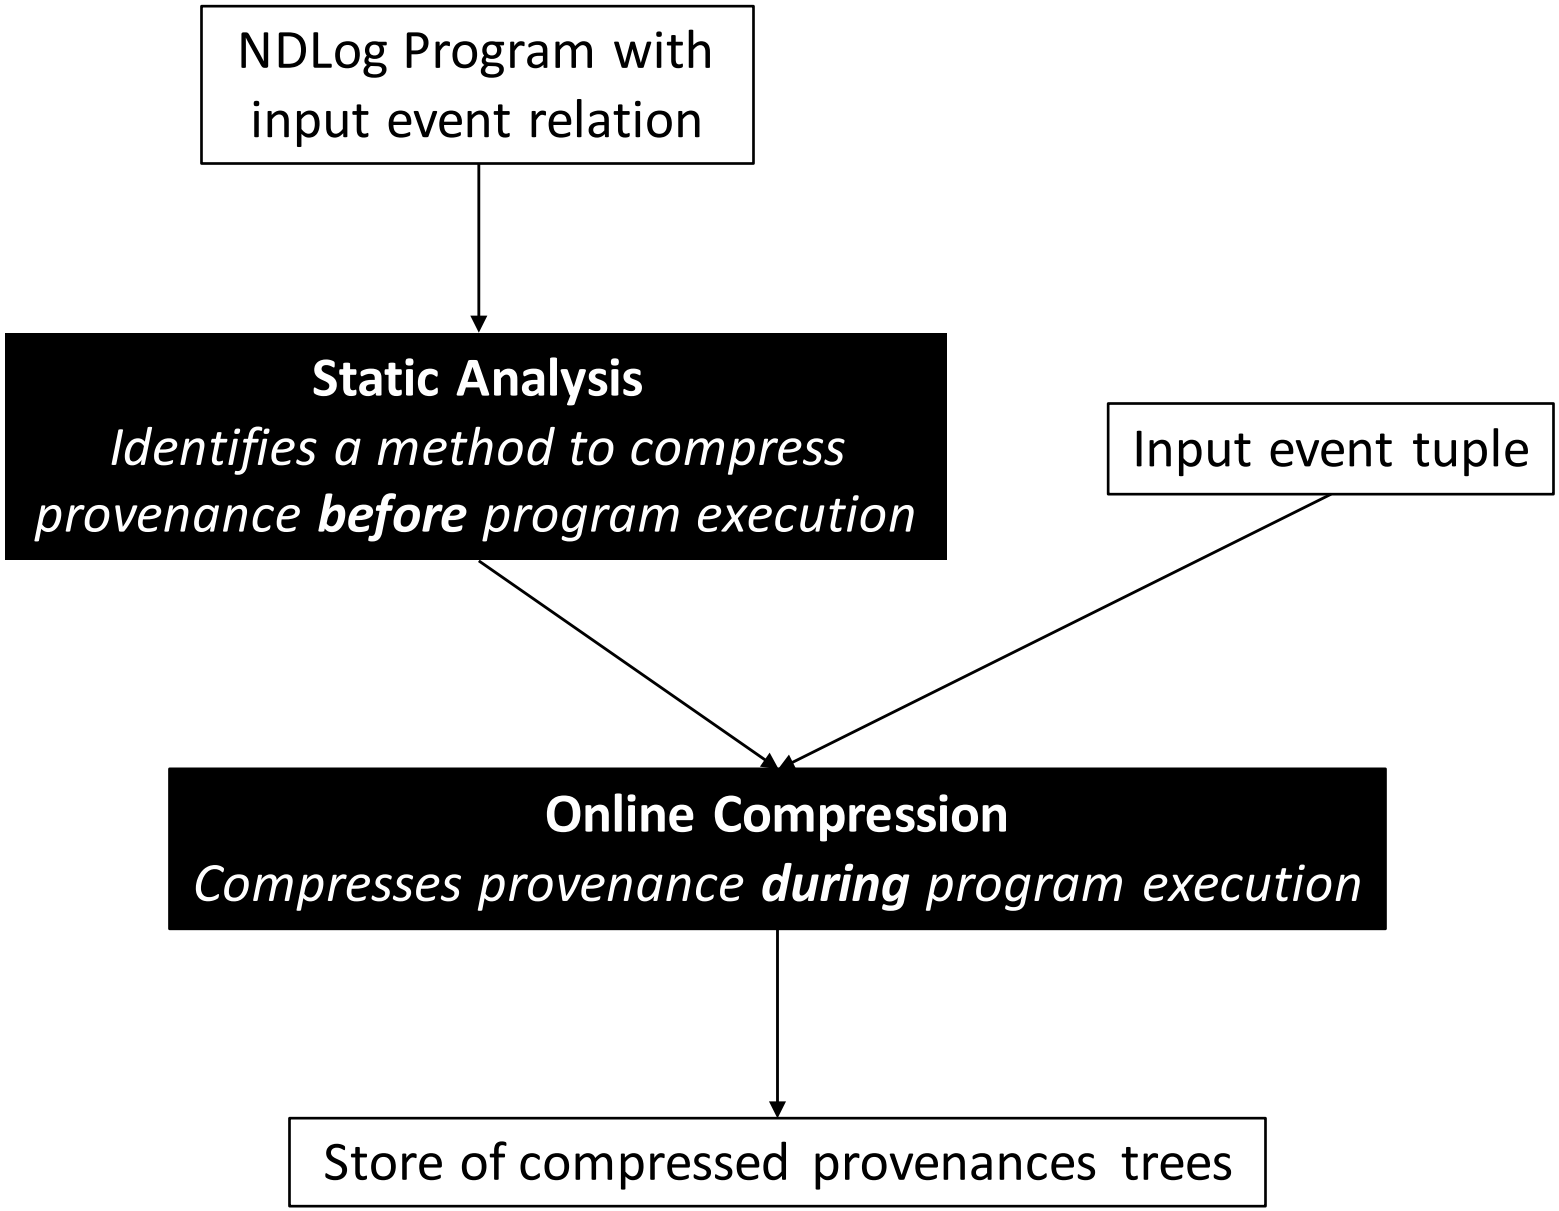

Workflow of Static Analysis

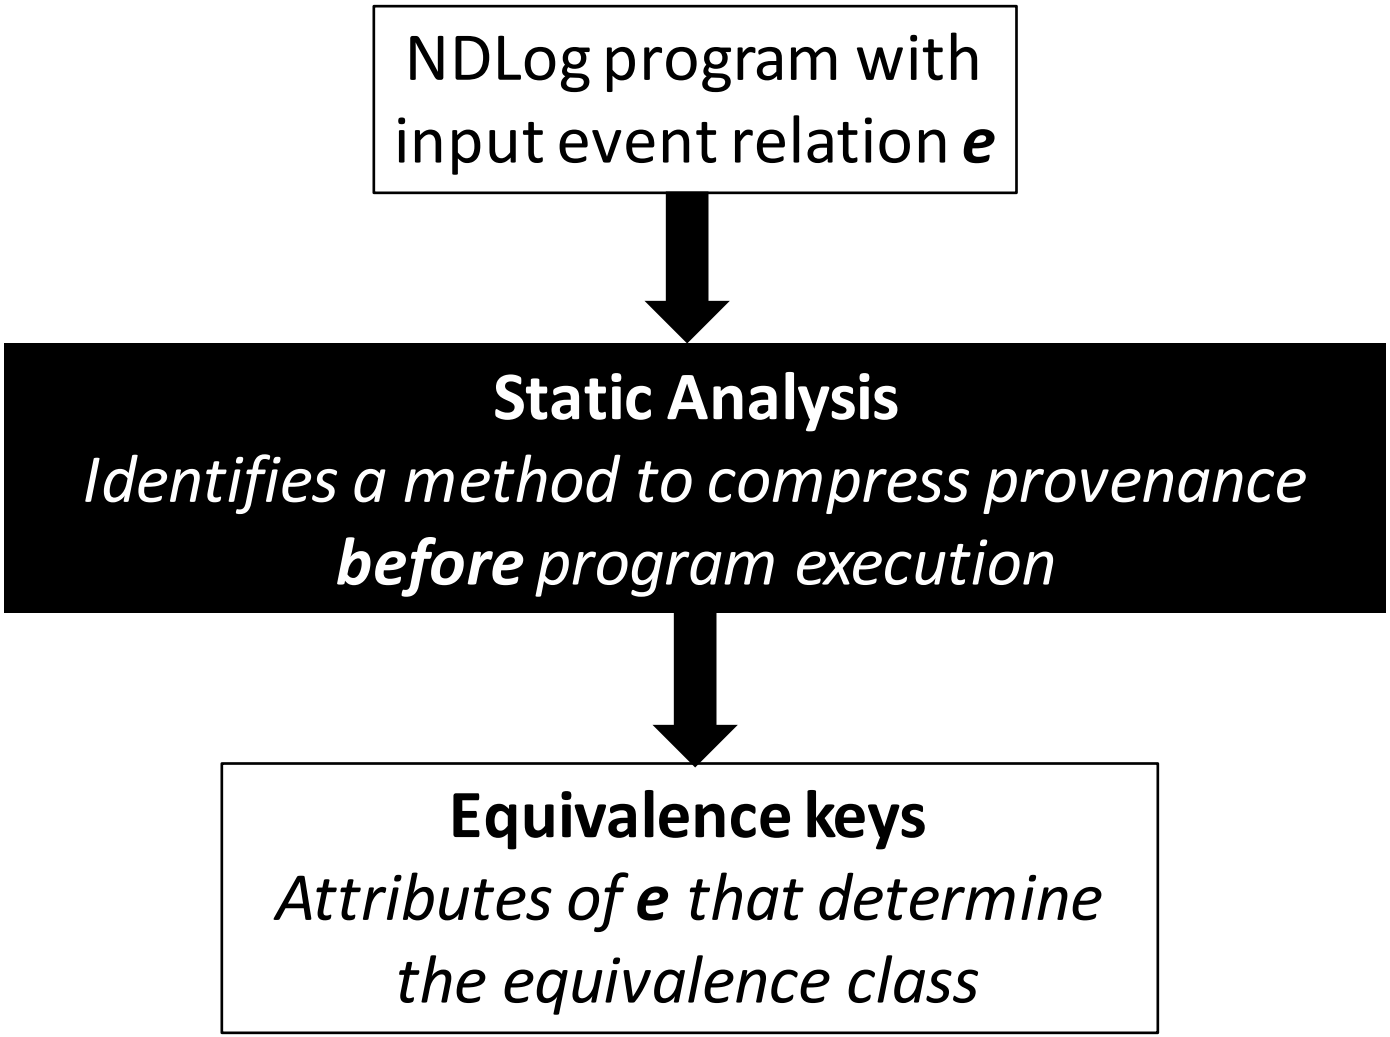

Which attributes of ***packet*** affect the provenance tree generated?

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

packet(@1,1,3,”hi”)

Provenance trees in the same equivalence class

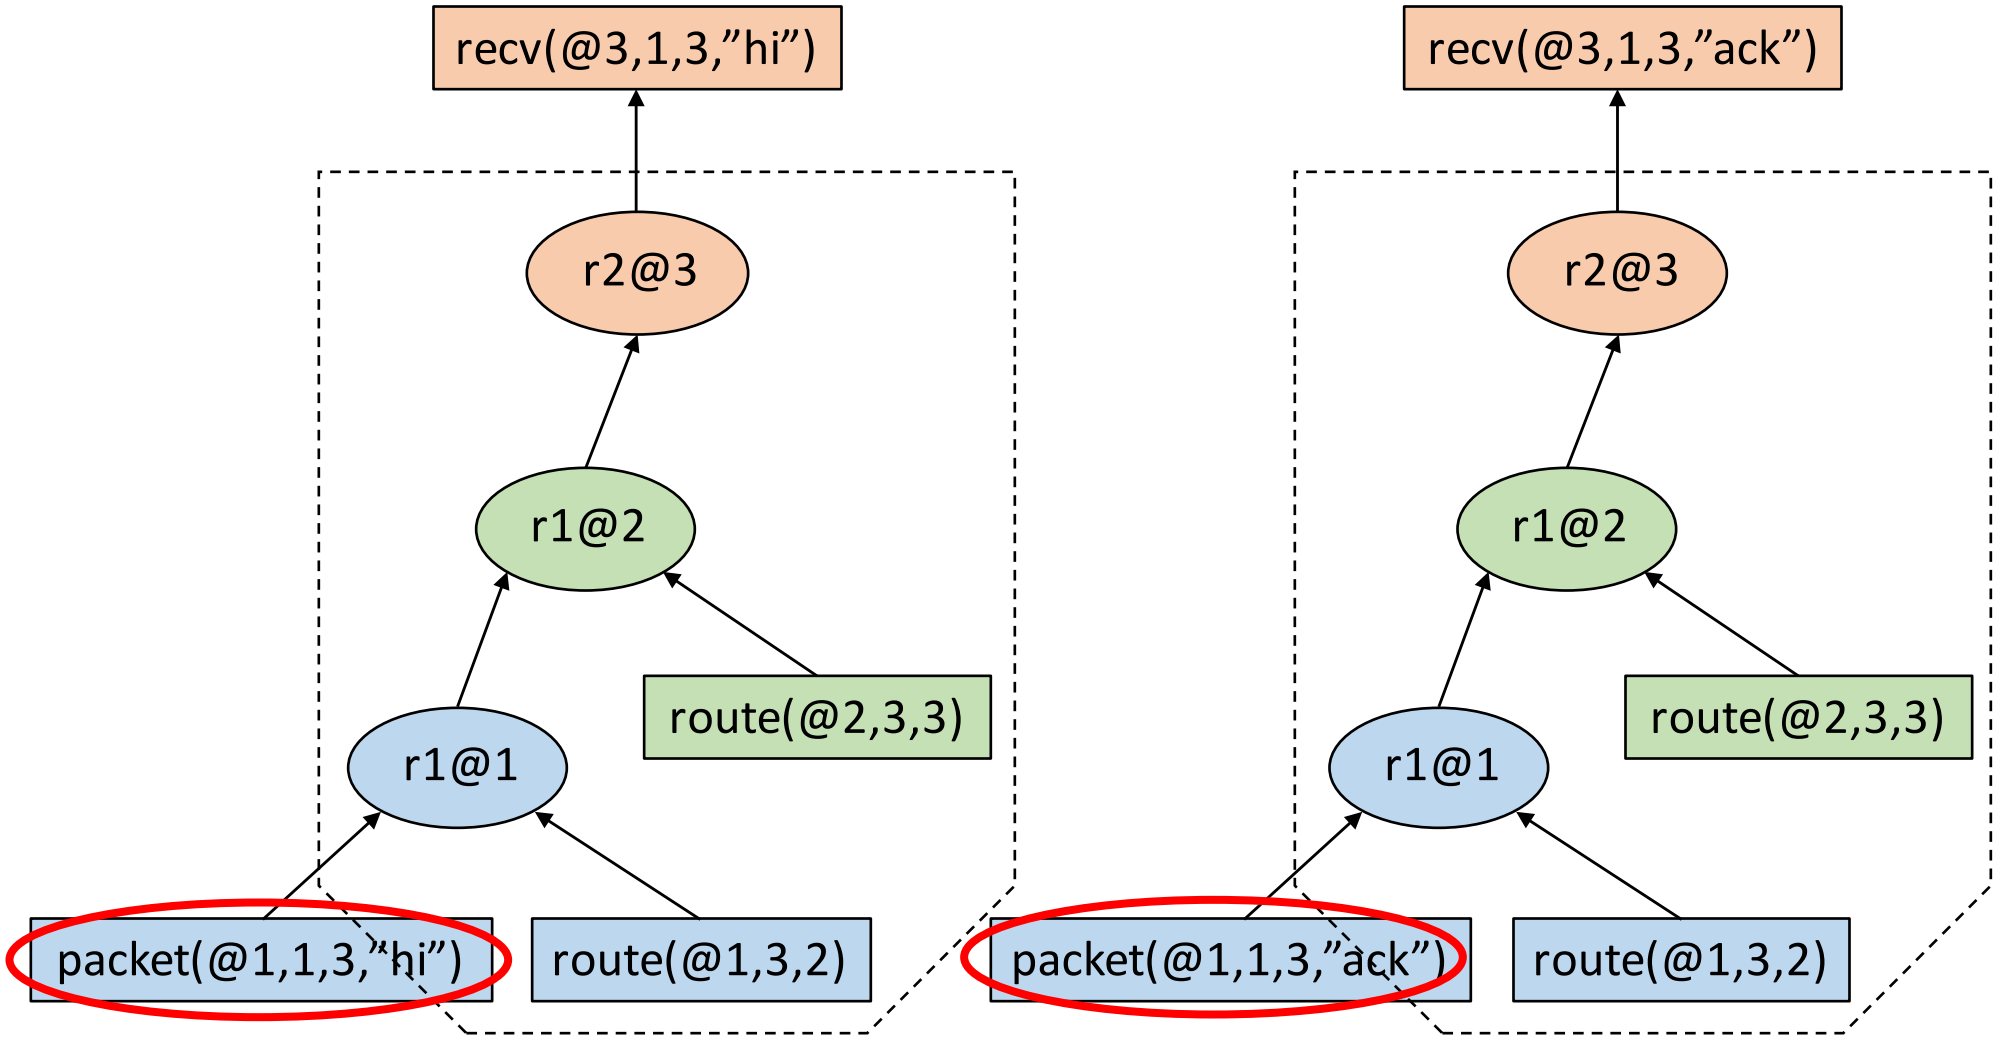

Identifying an equivalence key

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

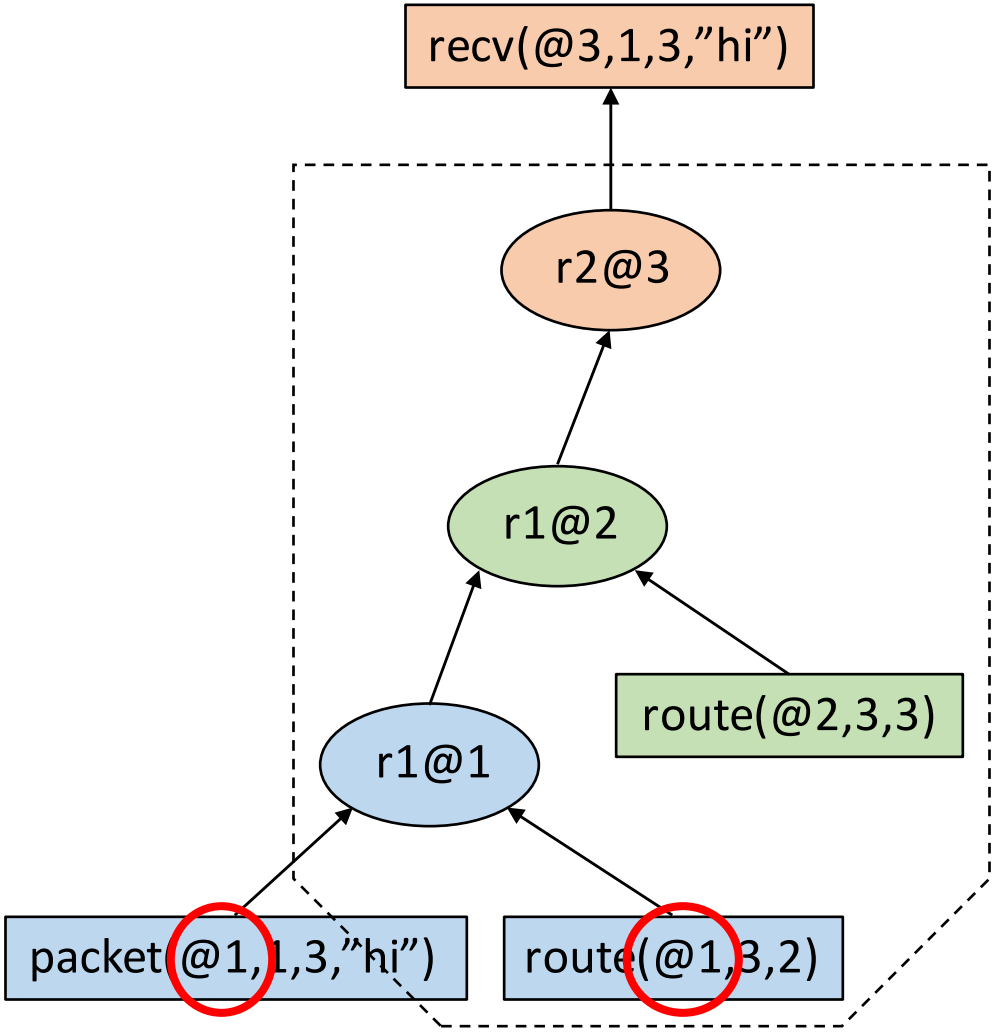

Identifying an equivalence key

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

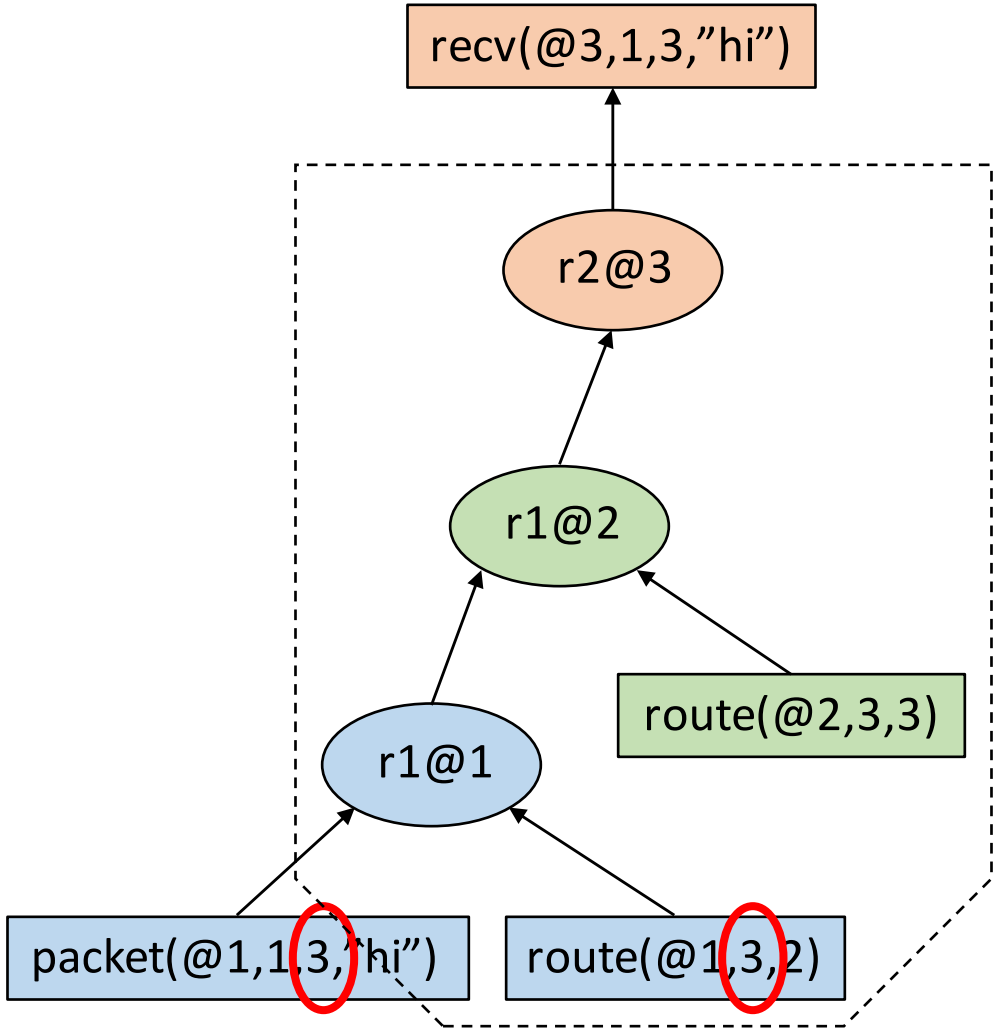

Not an equivalence key

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

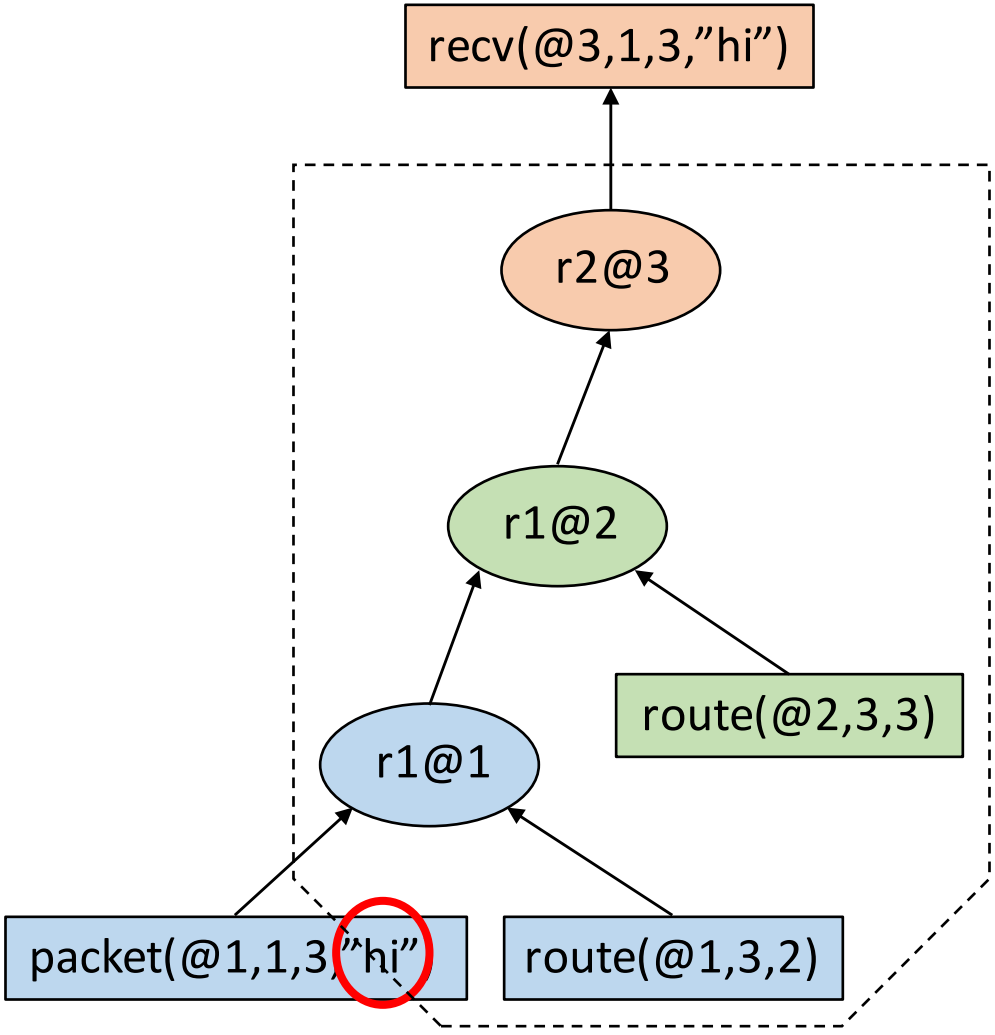

Correctness of Static Analysis

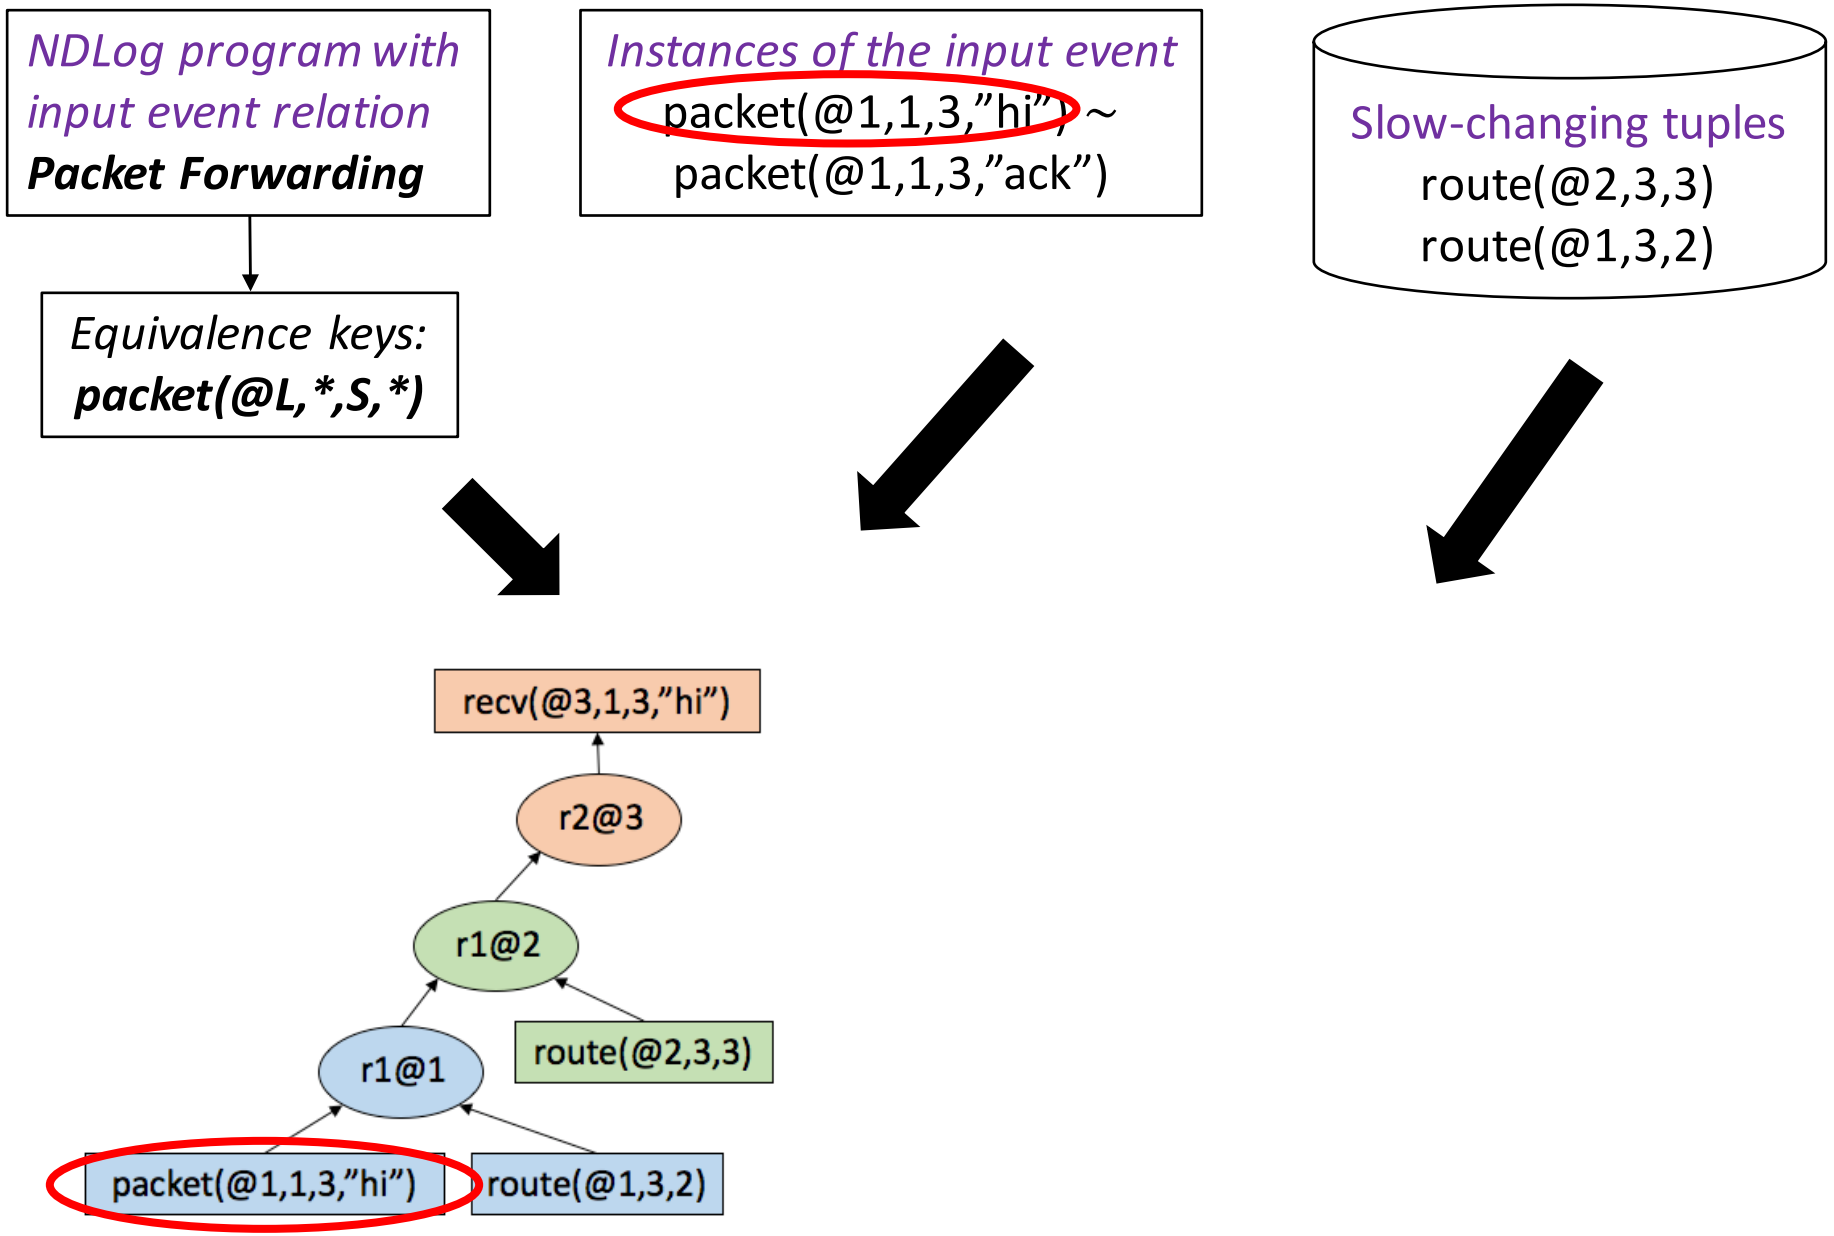

Correctness of Static Analysis

NDLog program with input event relation ***Packet Forwarding***

Equivalence keys: ***packet(@L,*,S,*)***

Instances of the input event

packet(@1,1,3,"hi") ~
packet(@1,1,3,"ack")

Slow-changing tuples

route(@2,3,3)

route(@1,3,2)

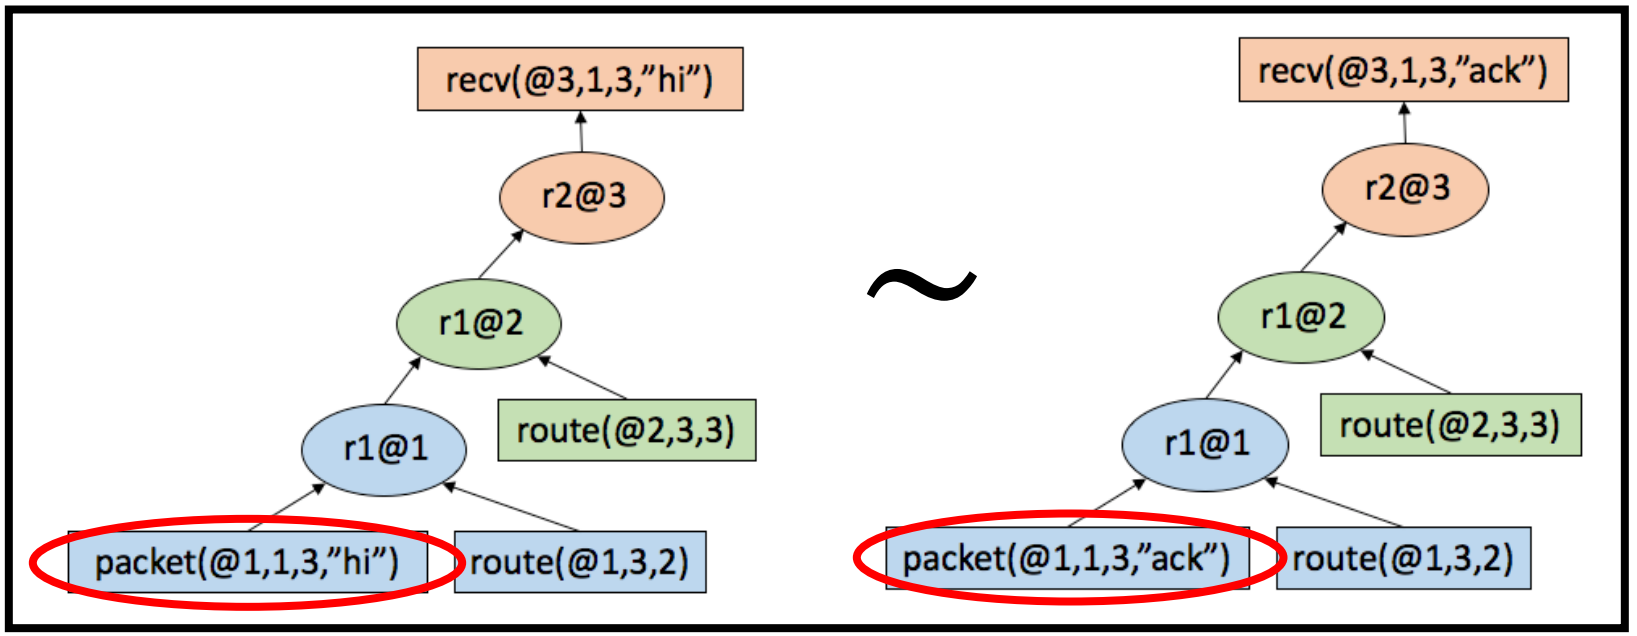

Workflow of Online Compression

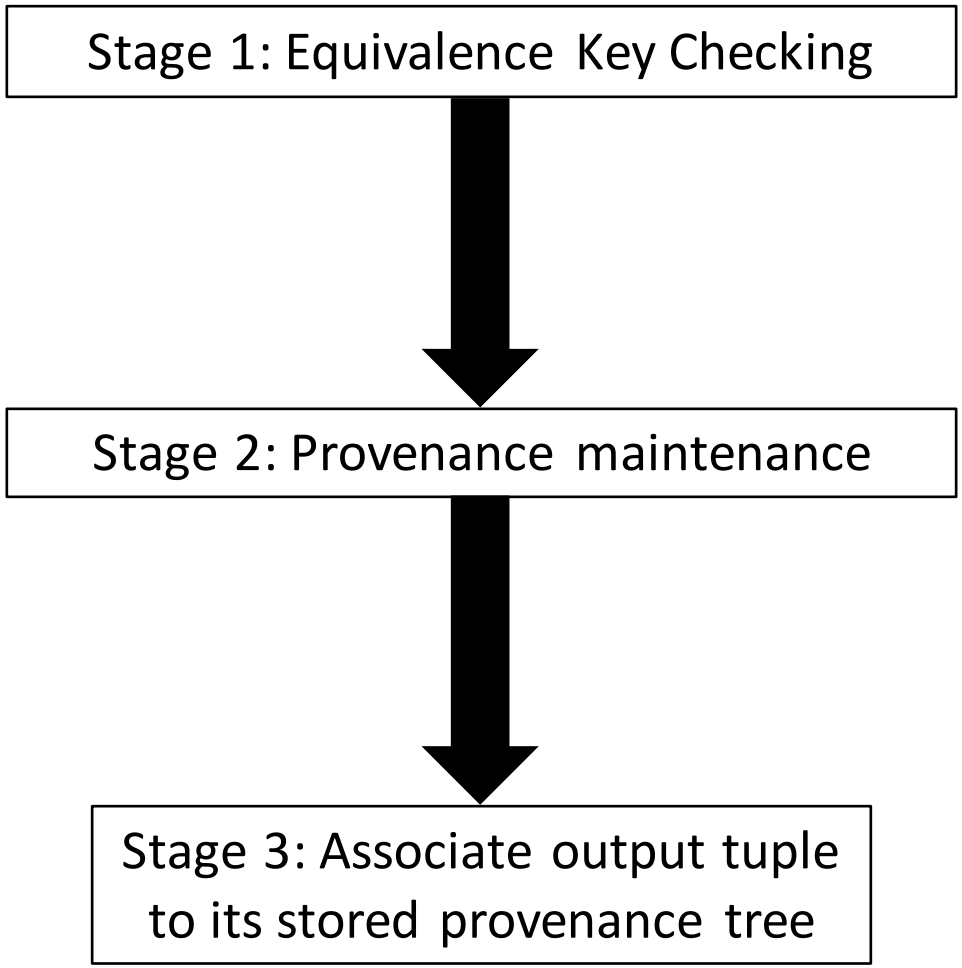

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

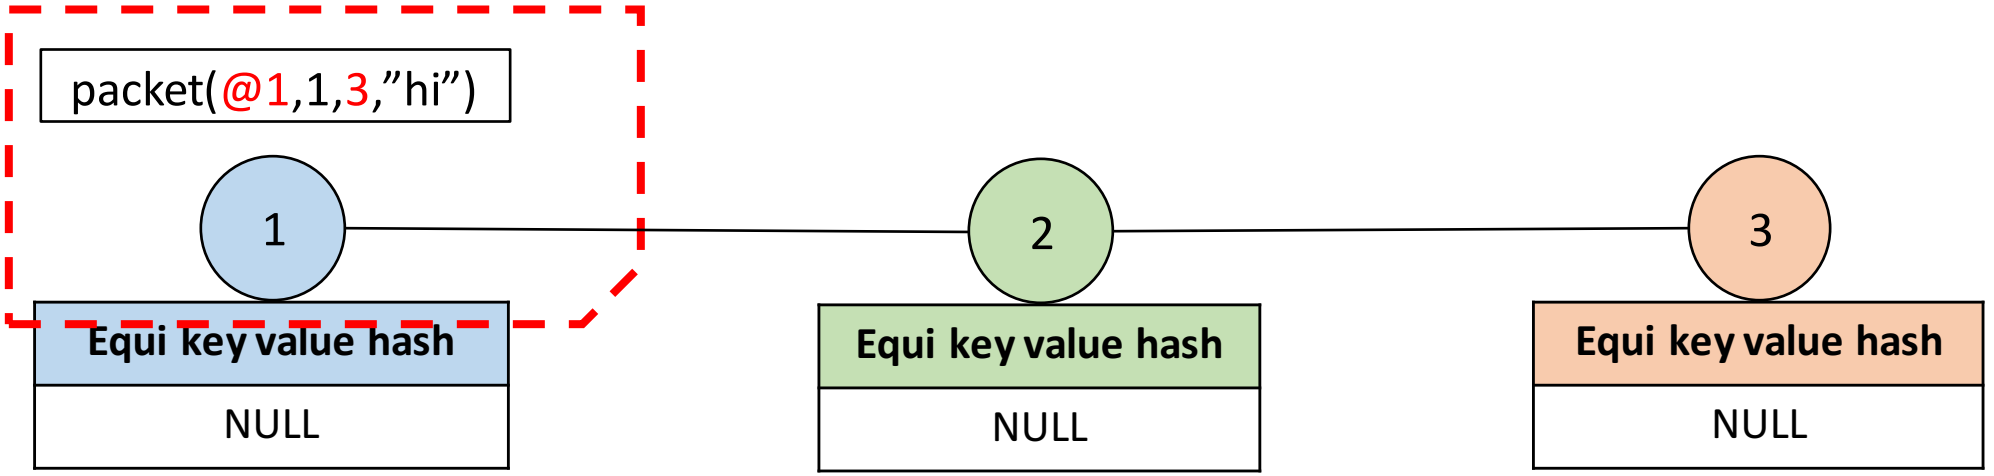

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

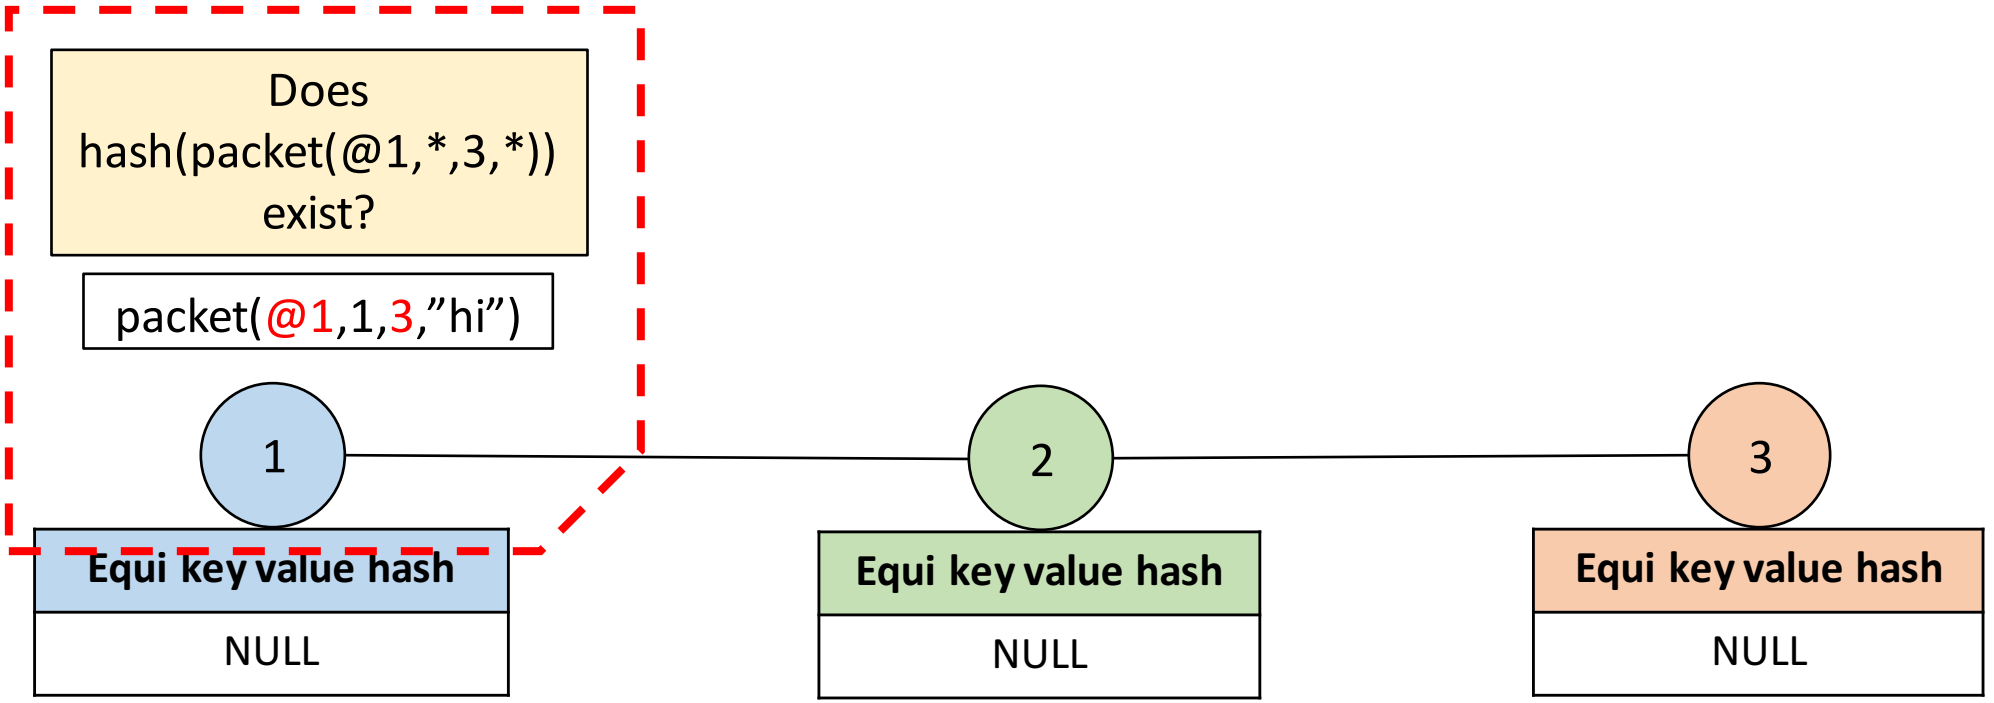

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

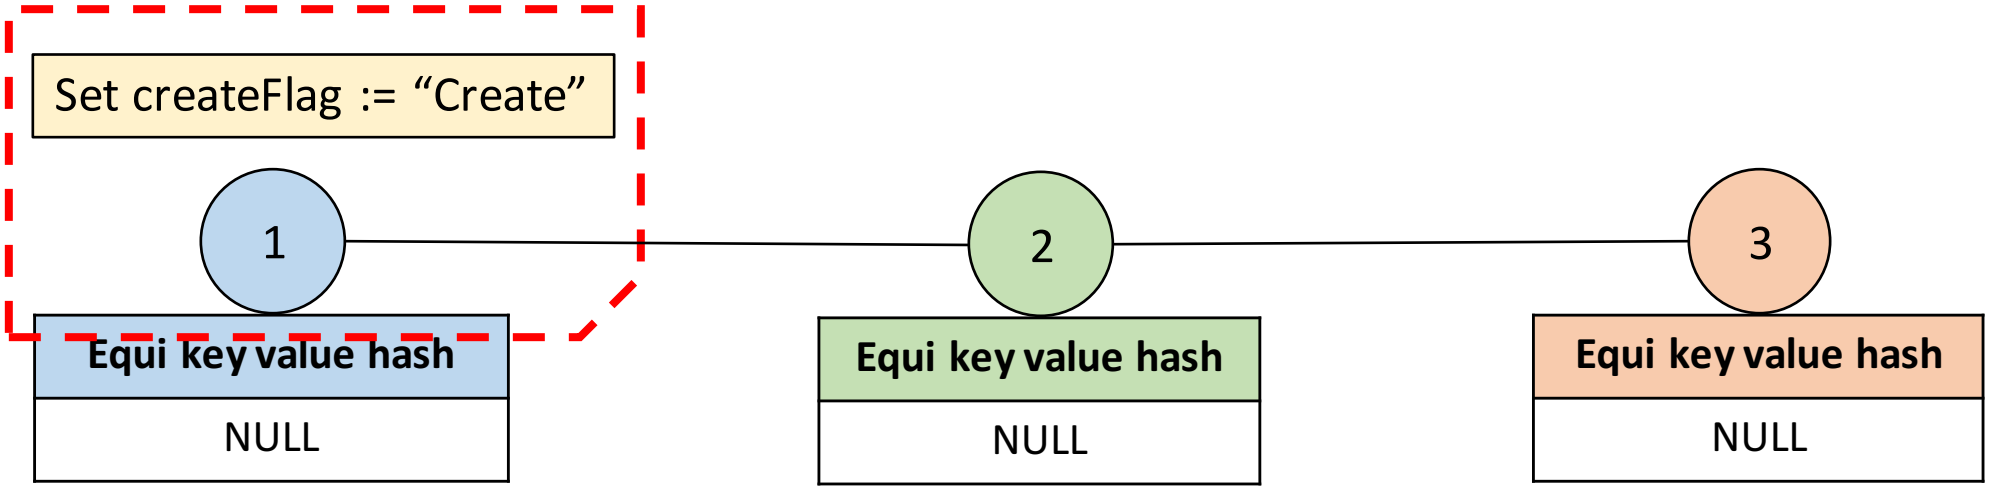

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

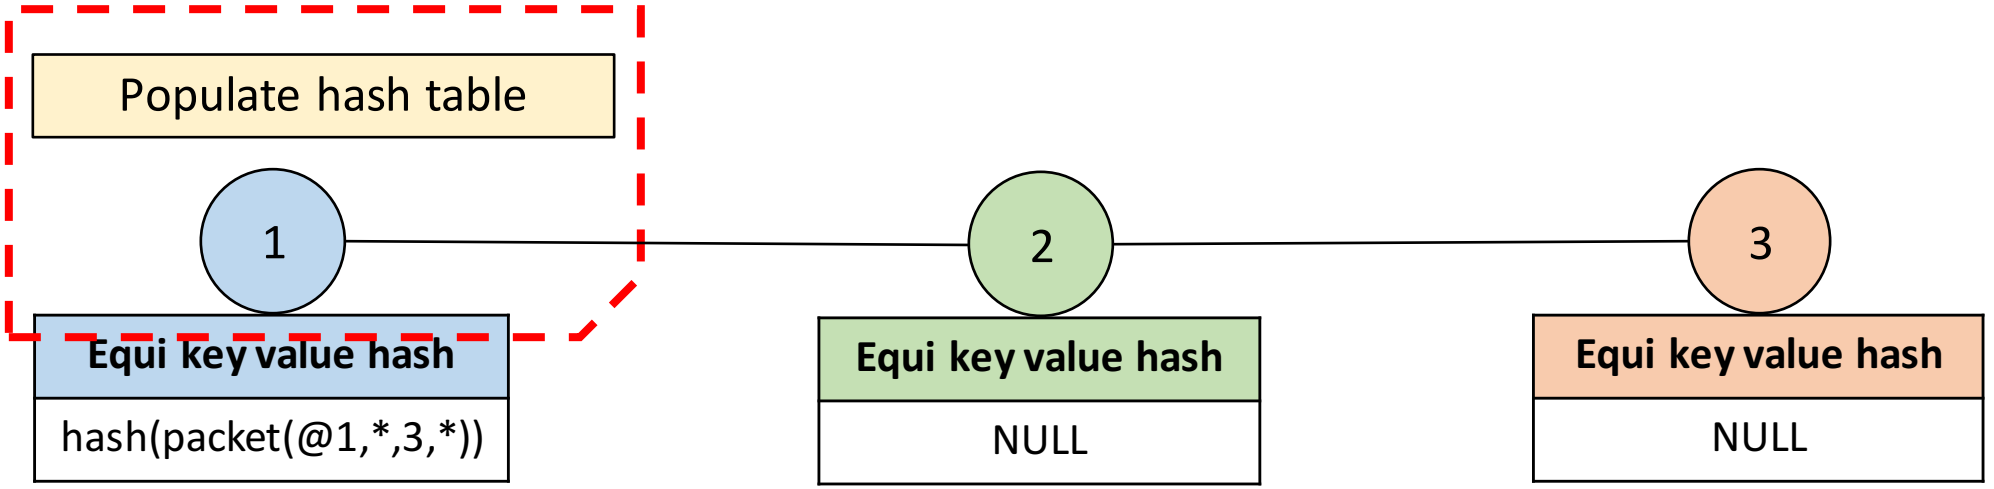

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

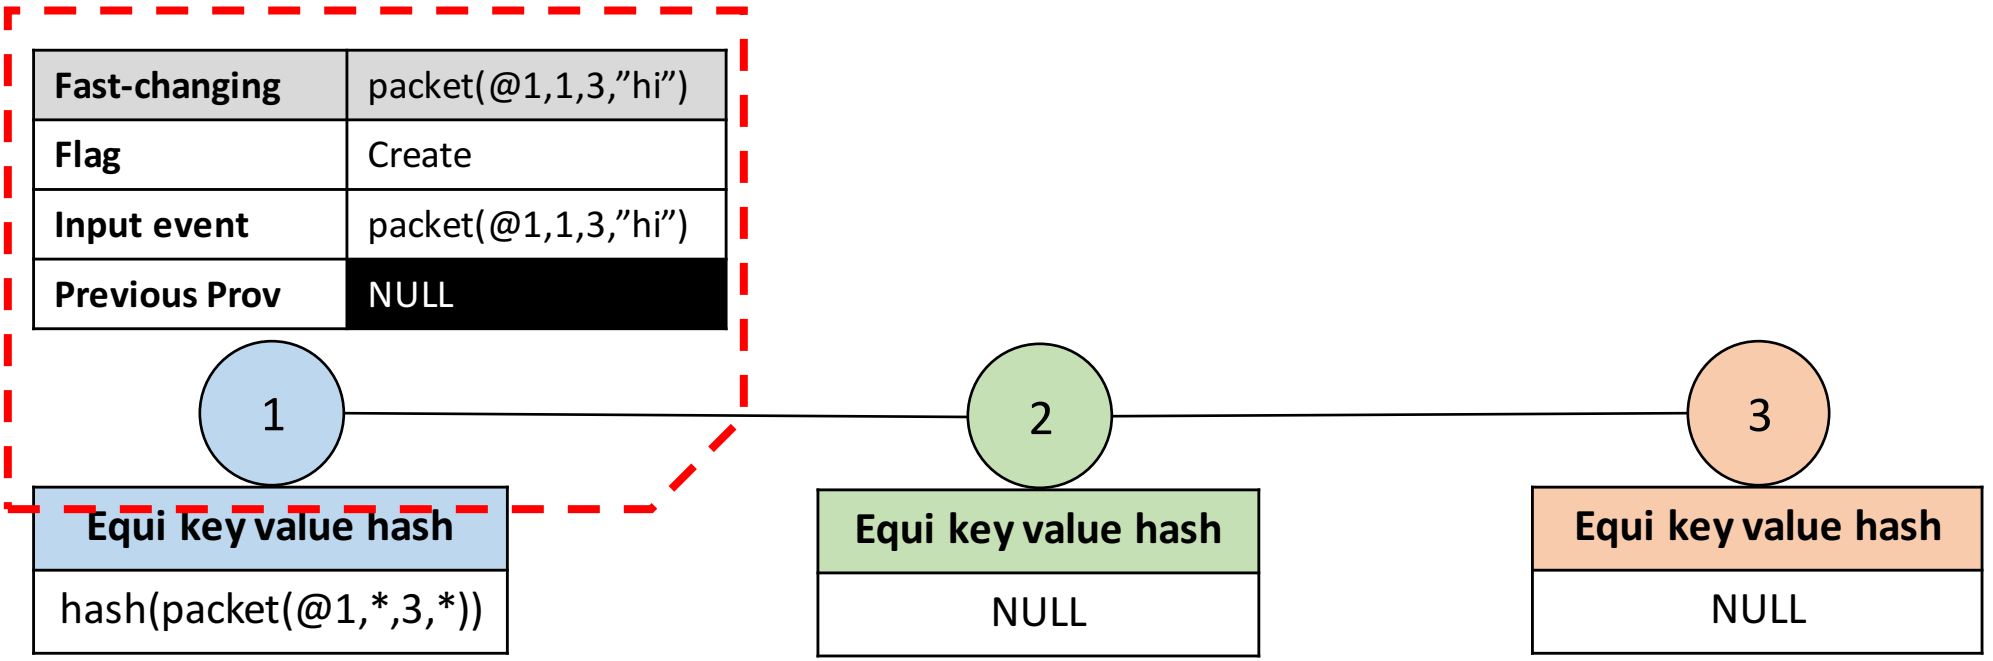

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

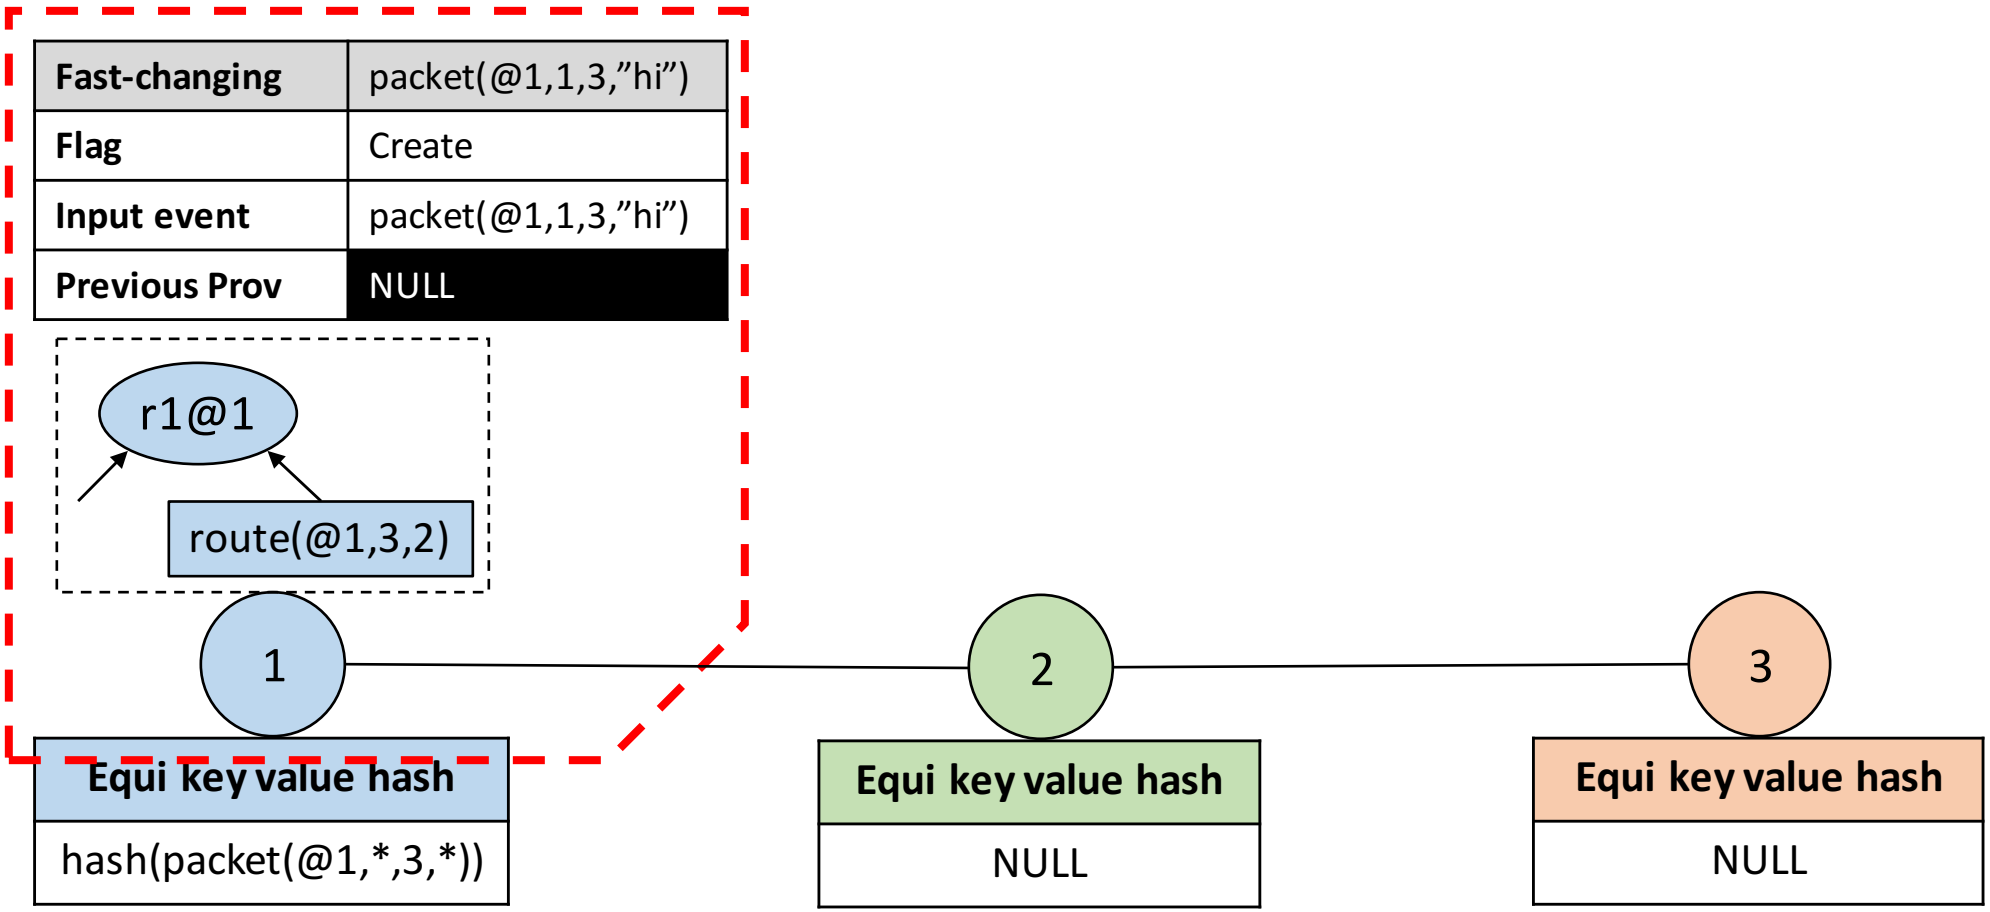

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

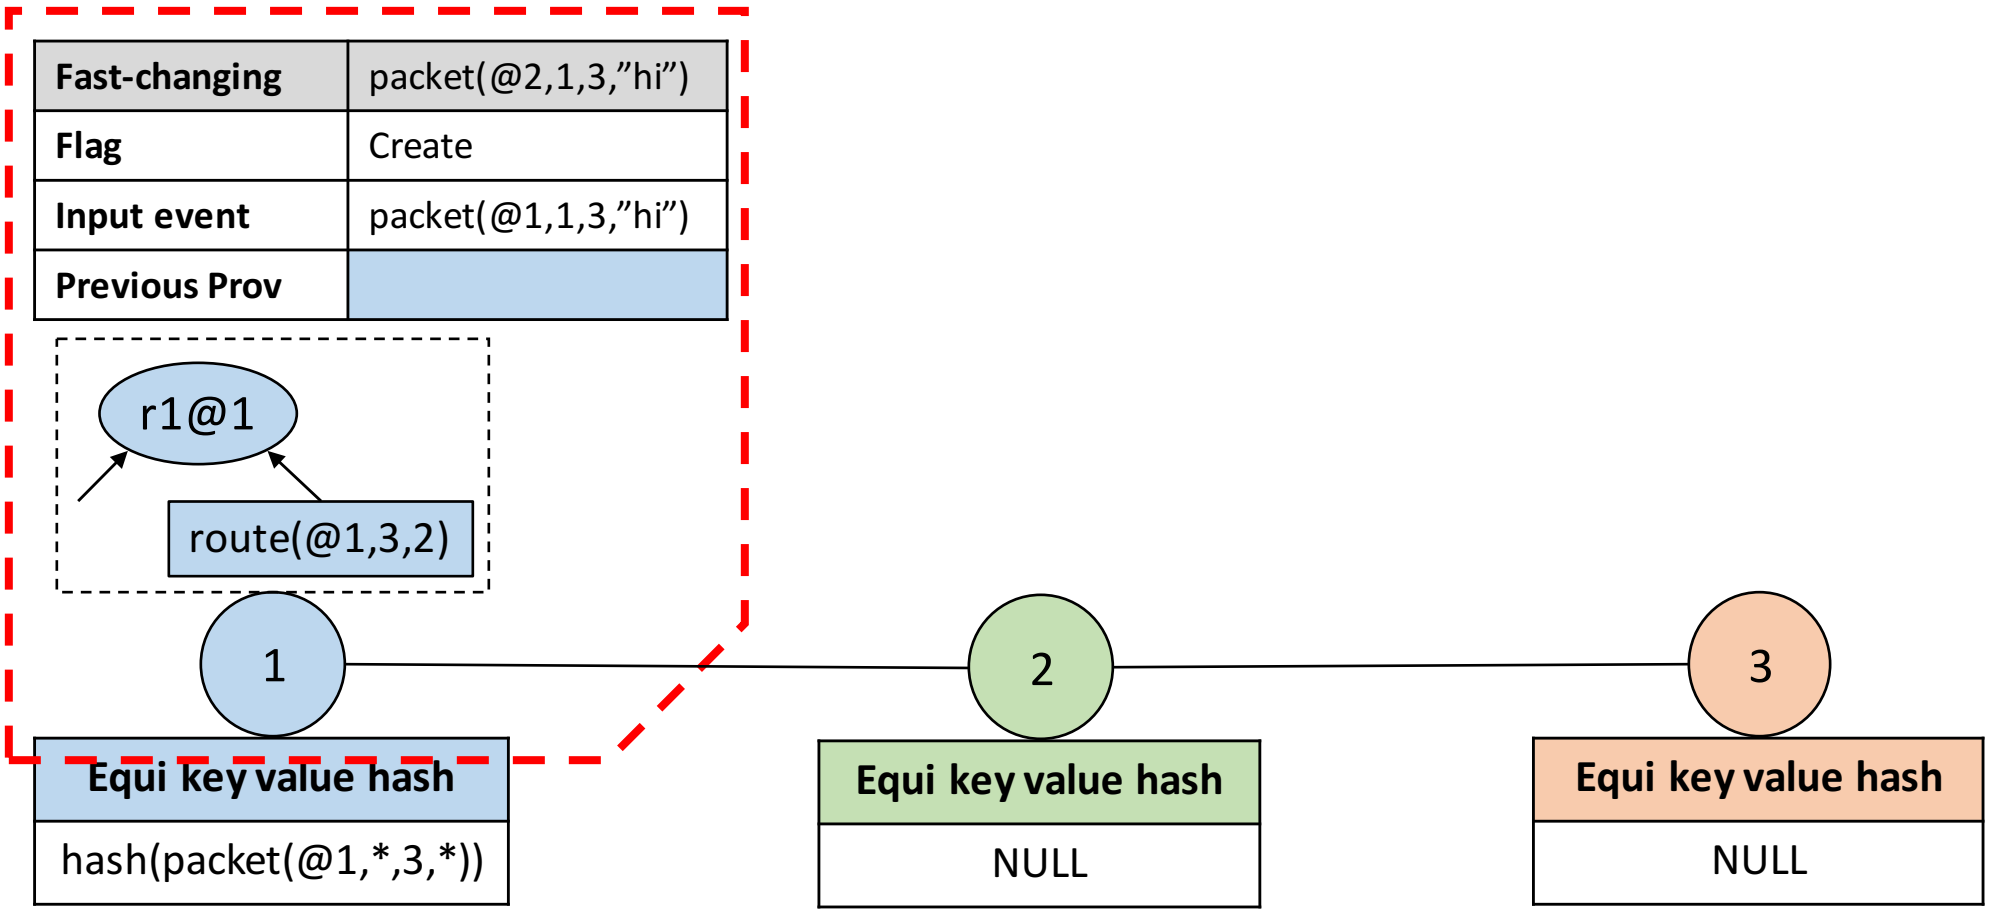

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

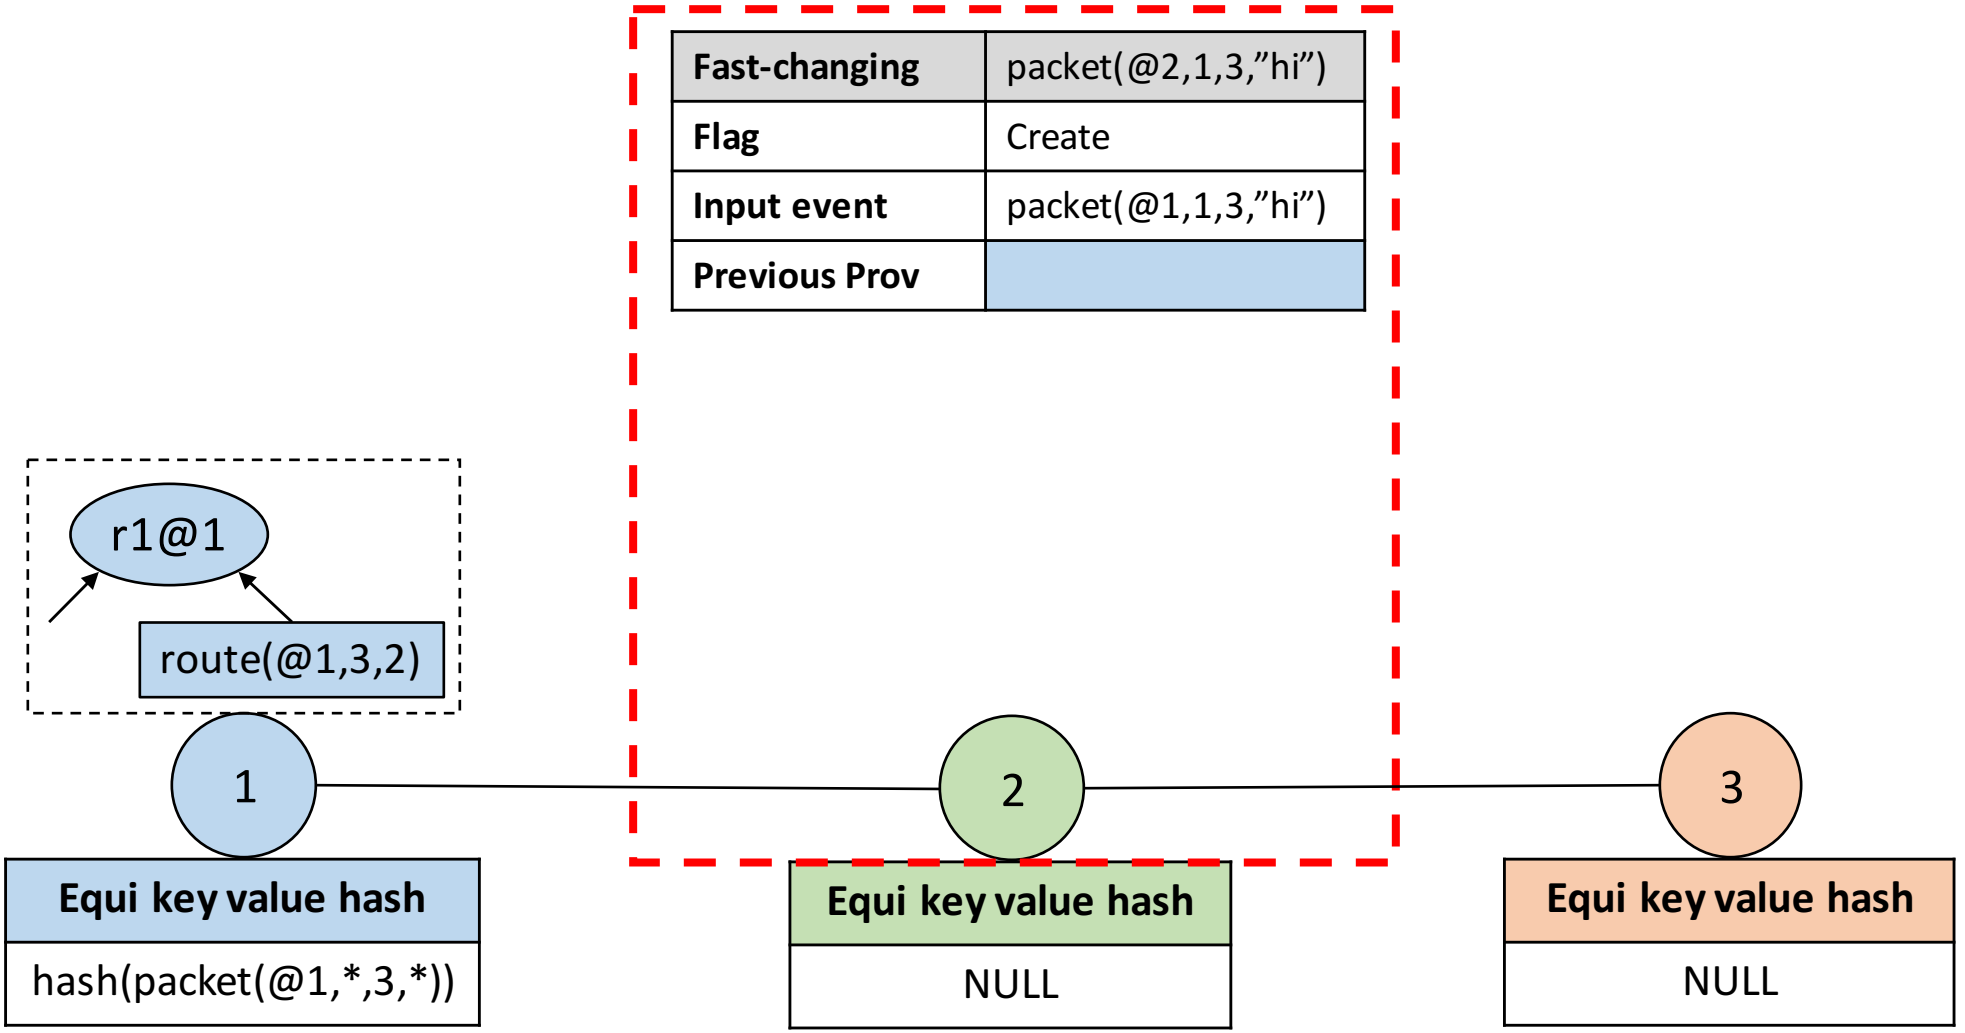

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

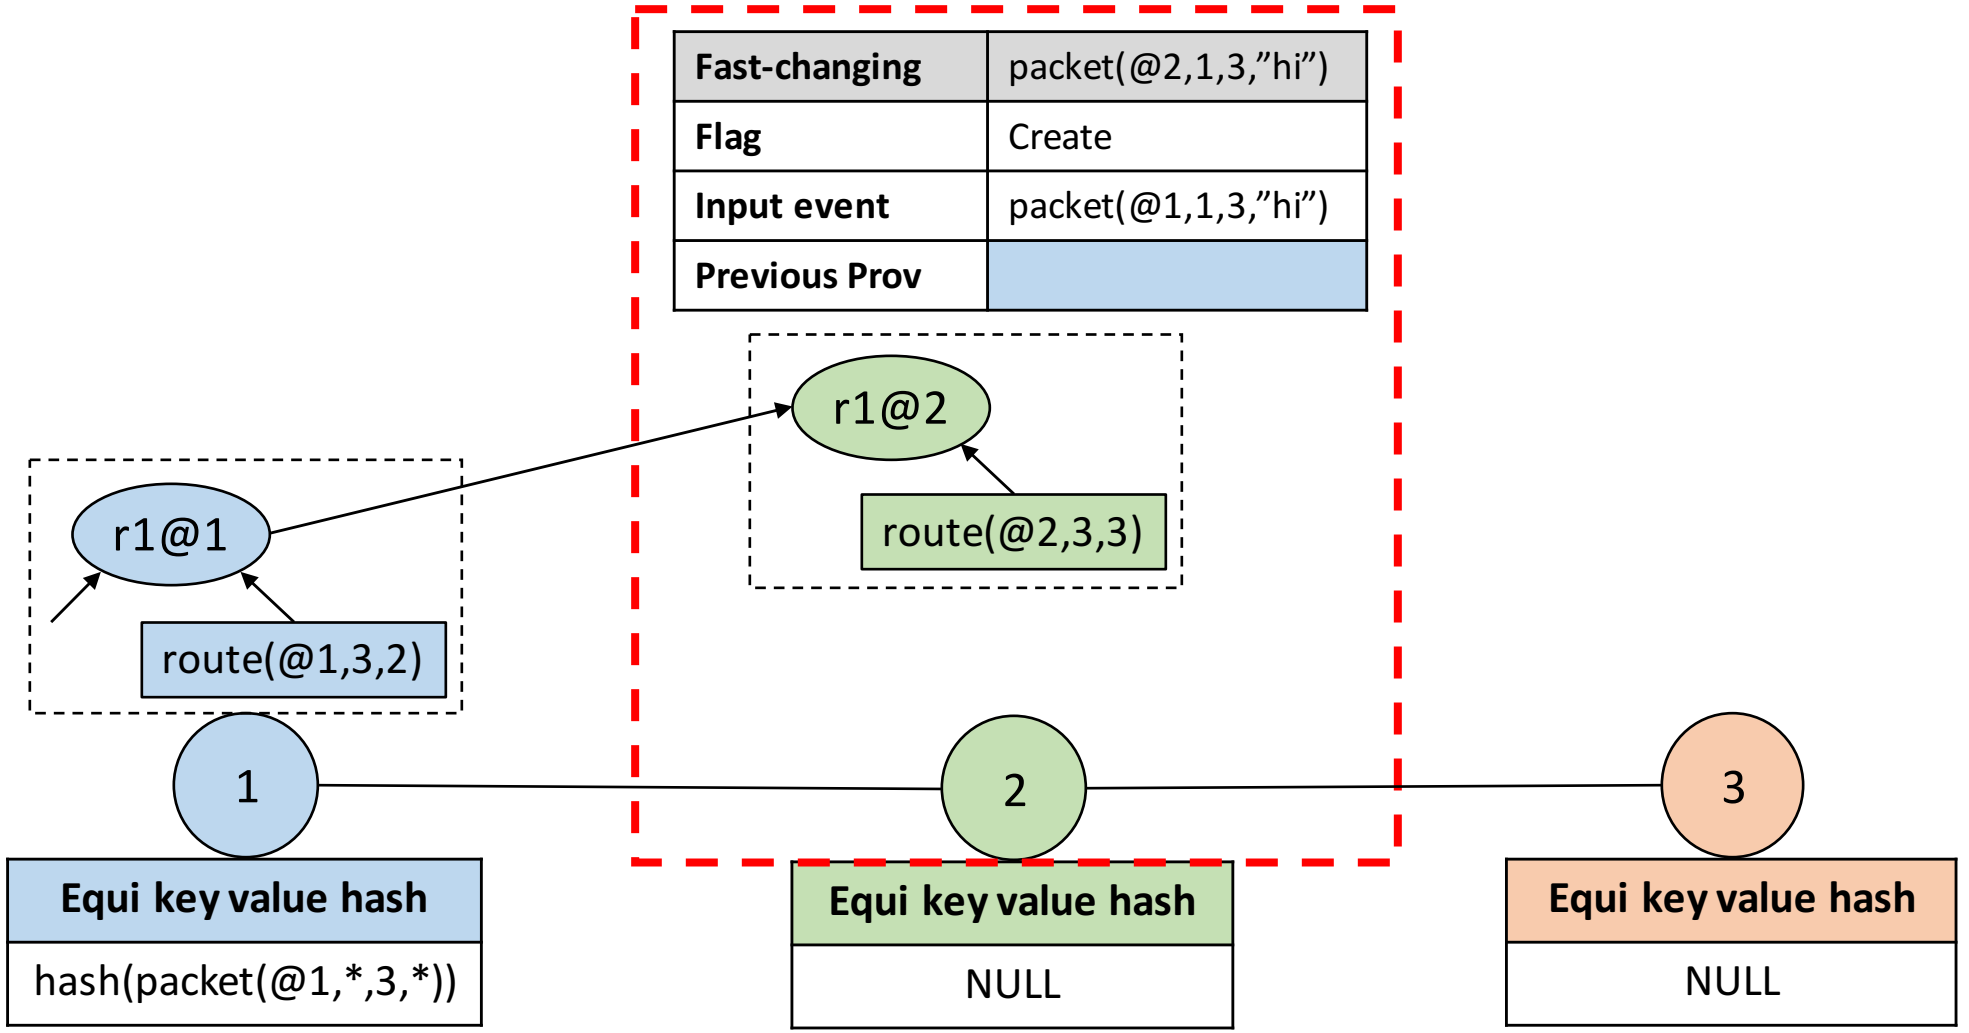

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

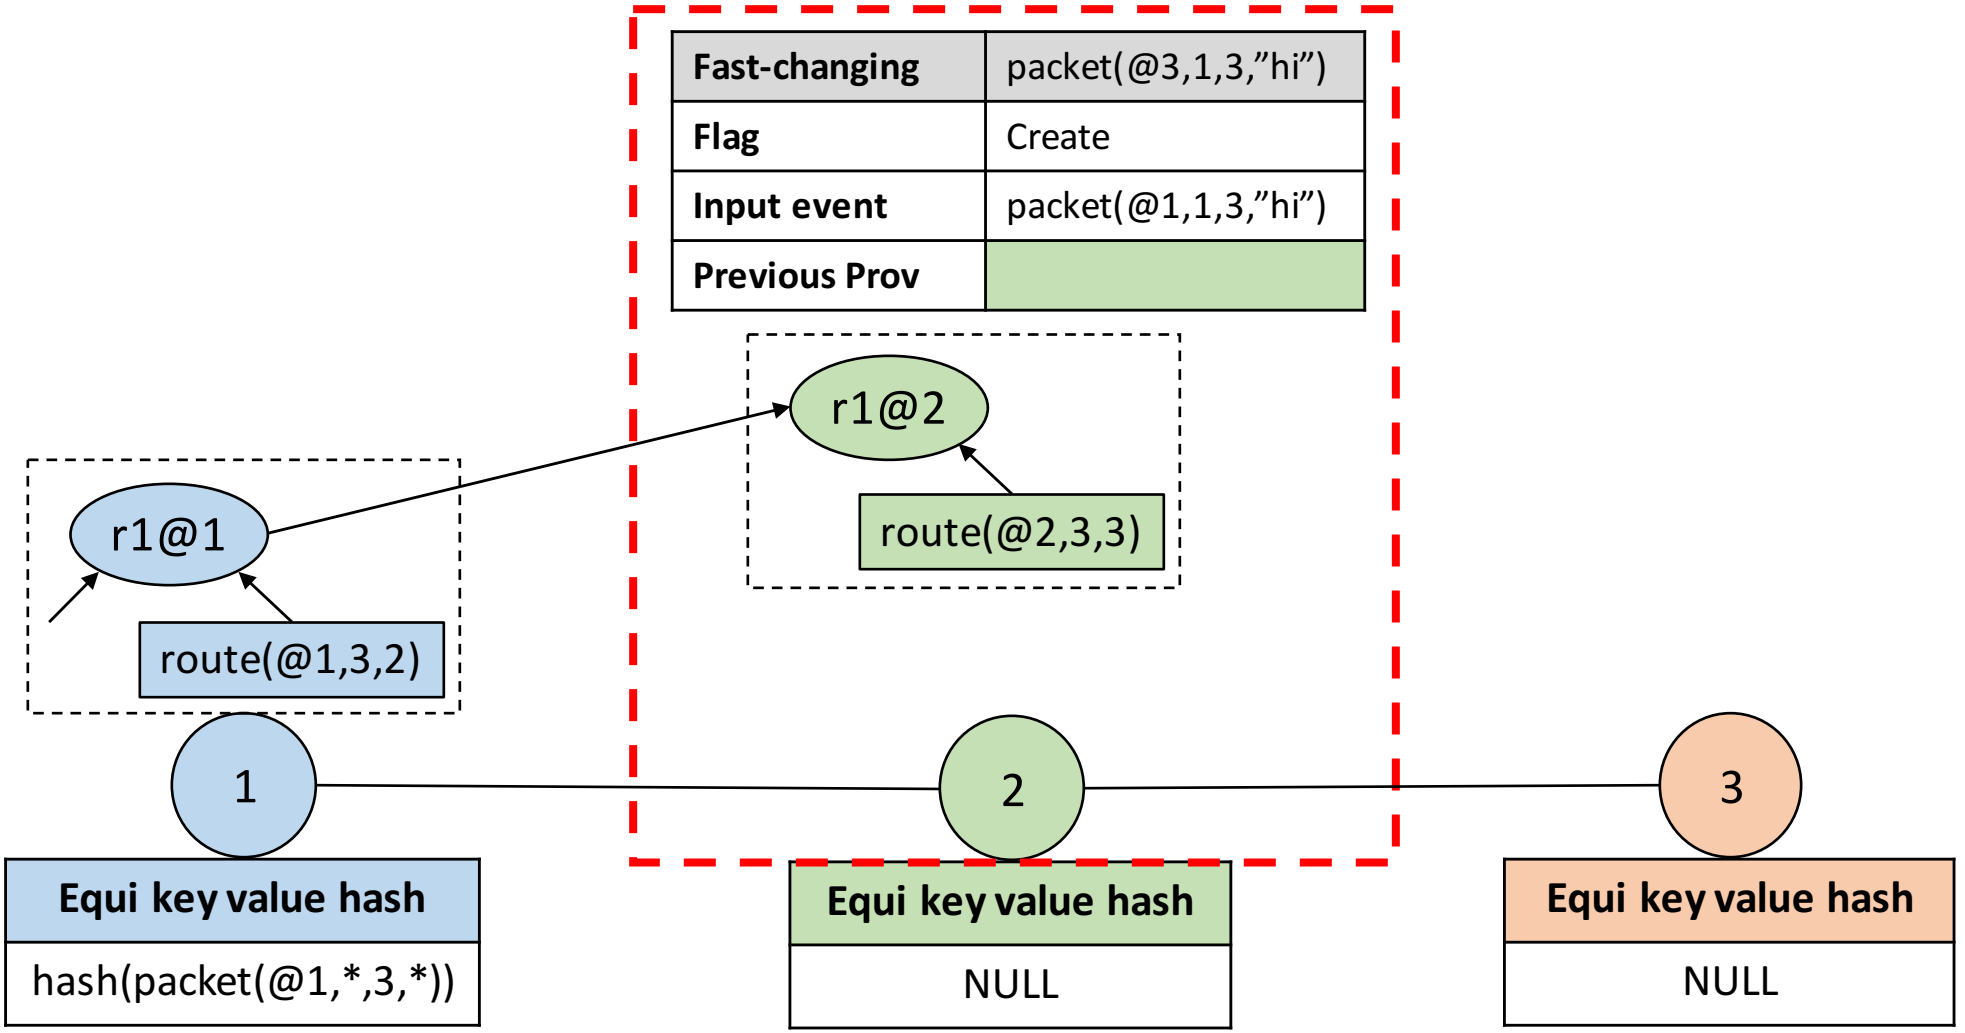

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

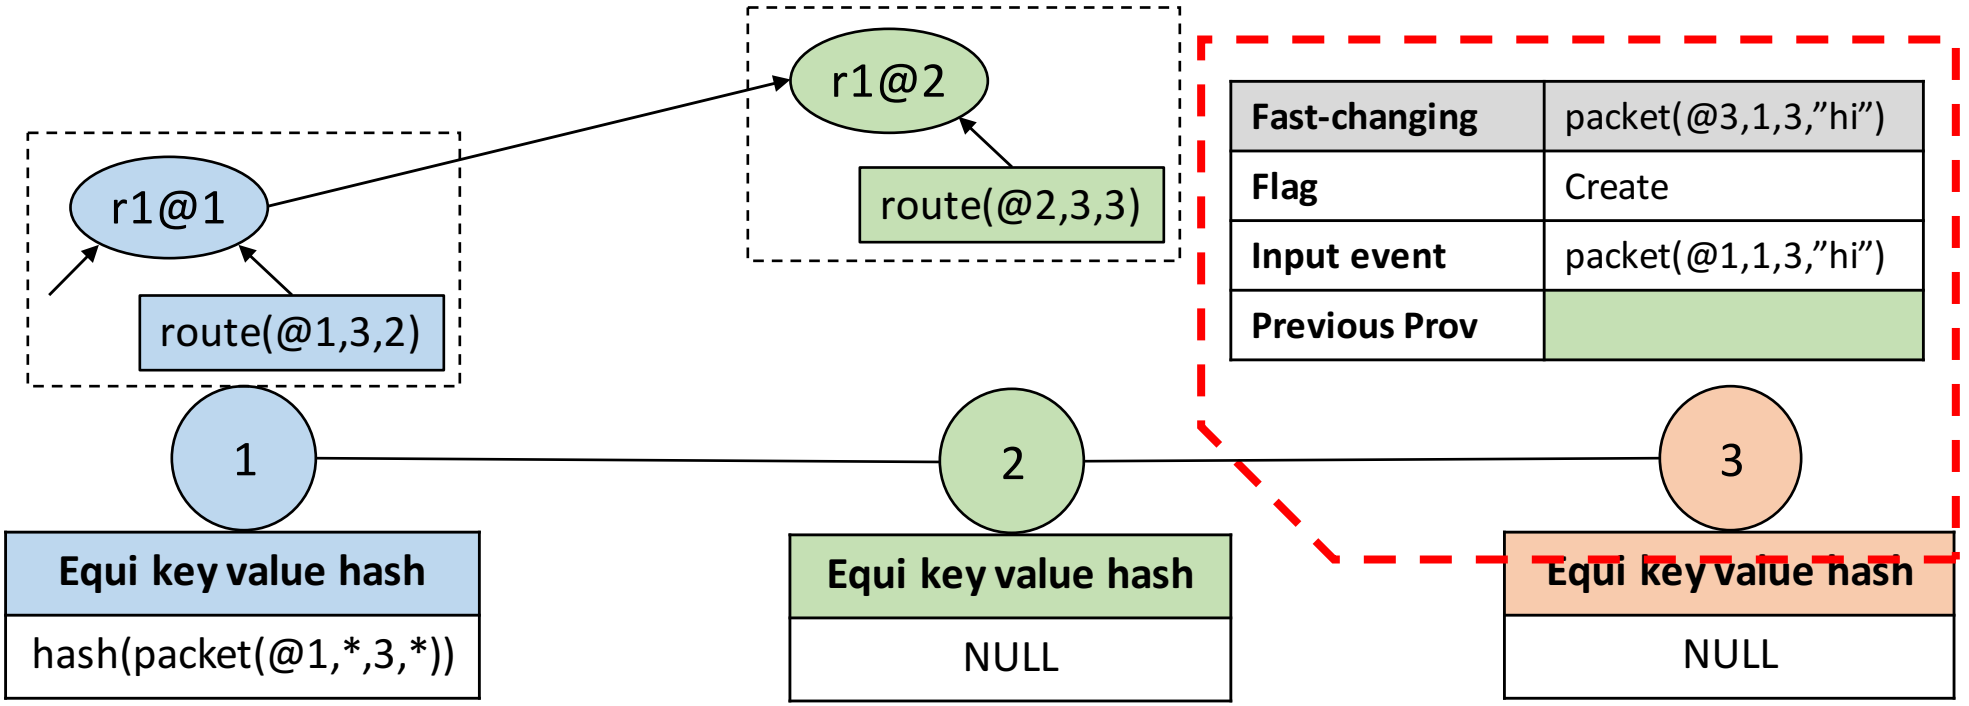

Stage 2: Provenance maintenance

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

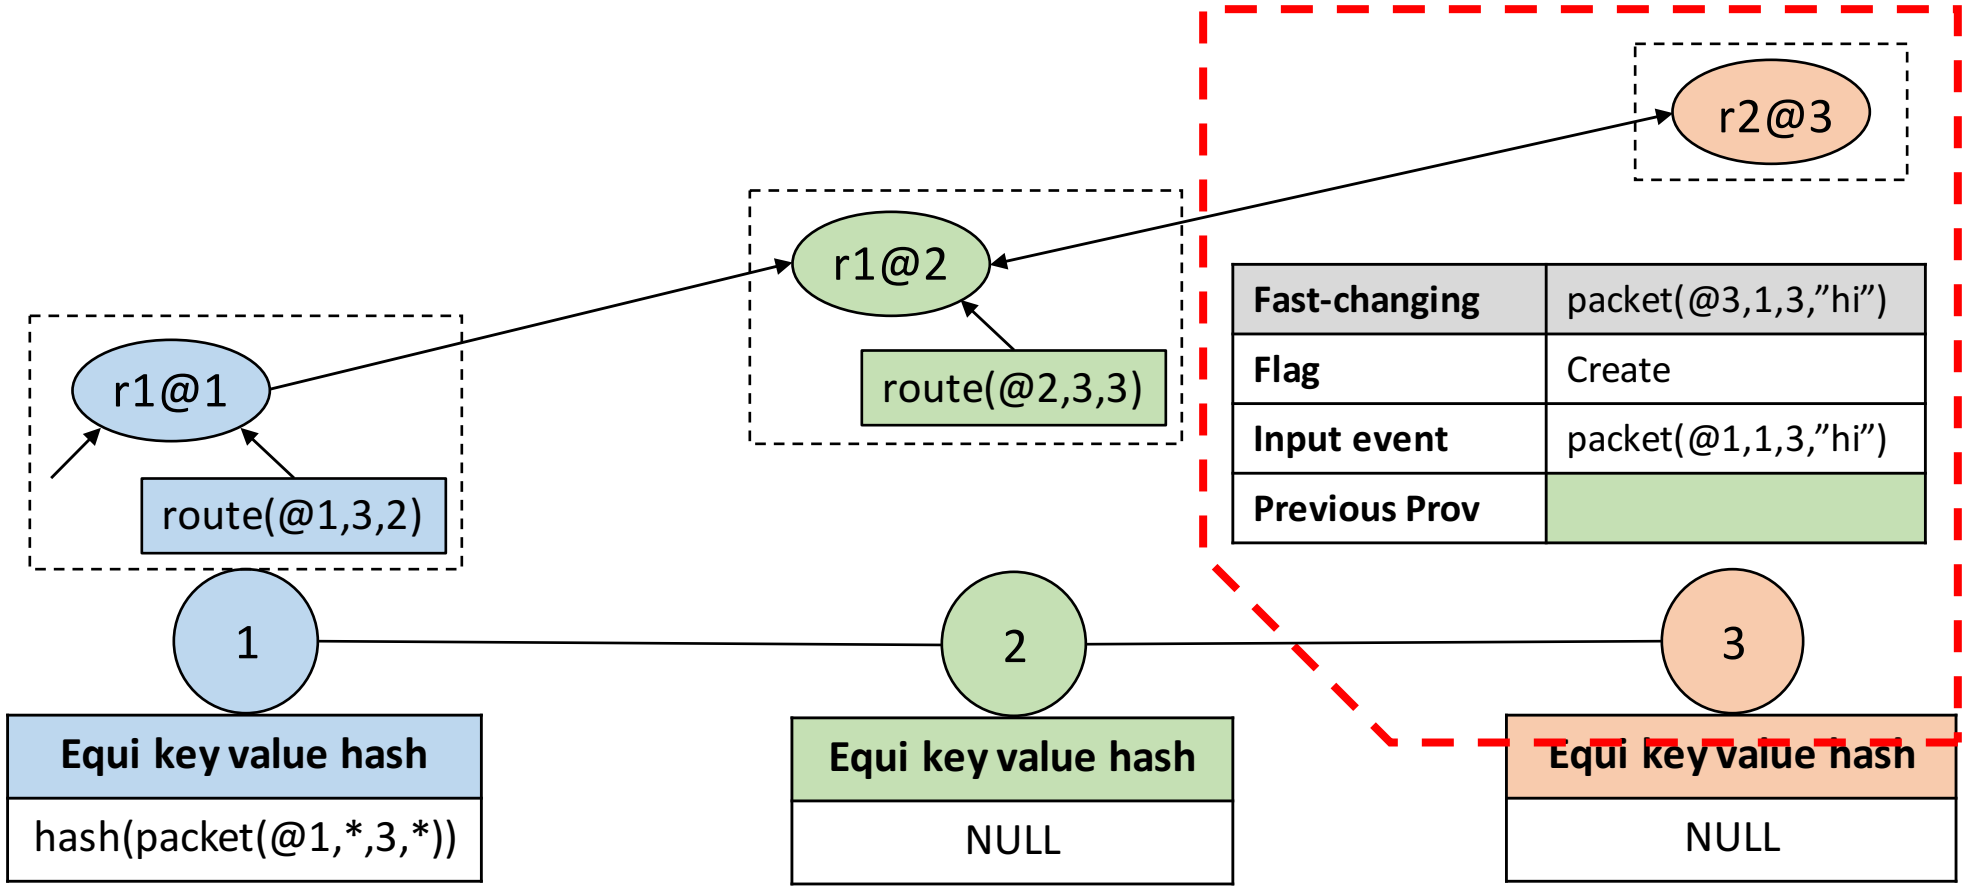

Stage 3: Provenance Association

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

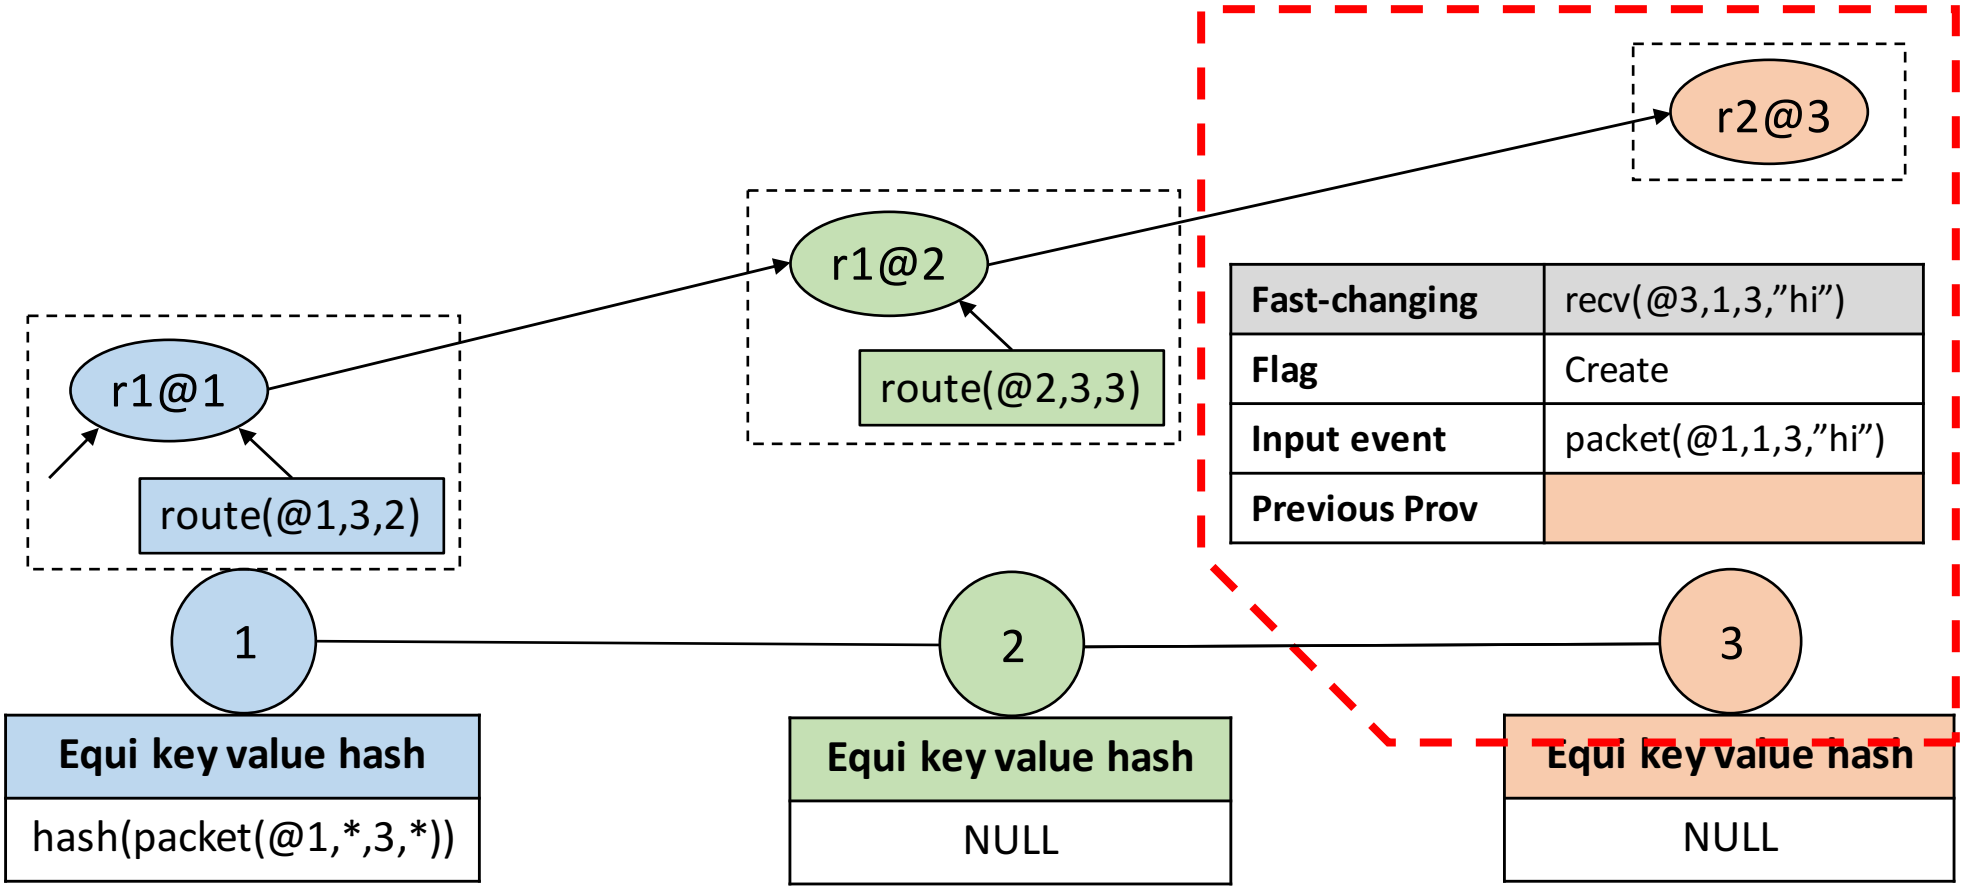

Stage 3: Associate output tuple to its stored provenance tree

Packet Forwarding
r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).
r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

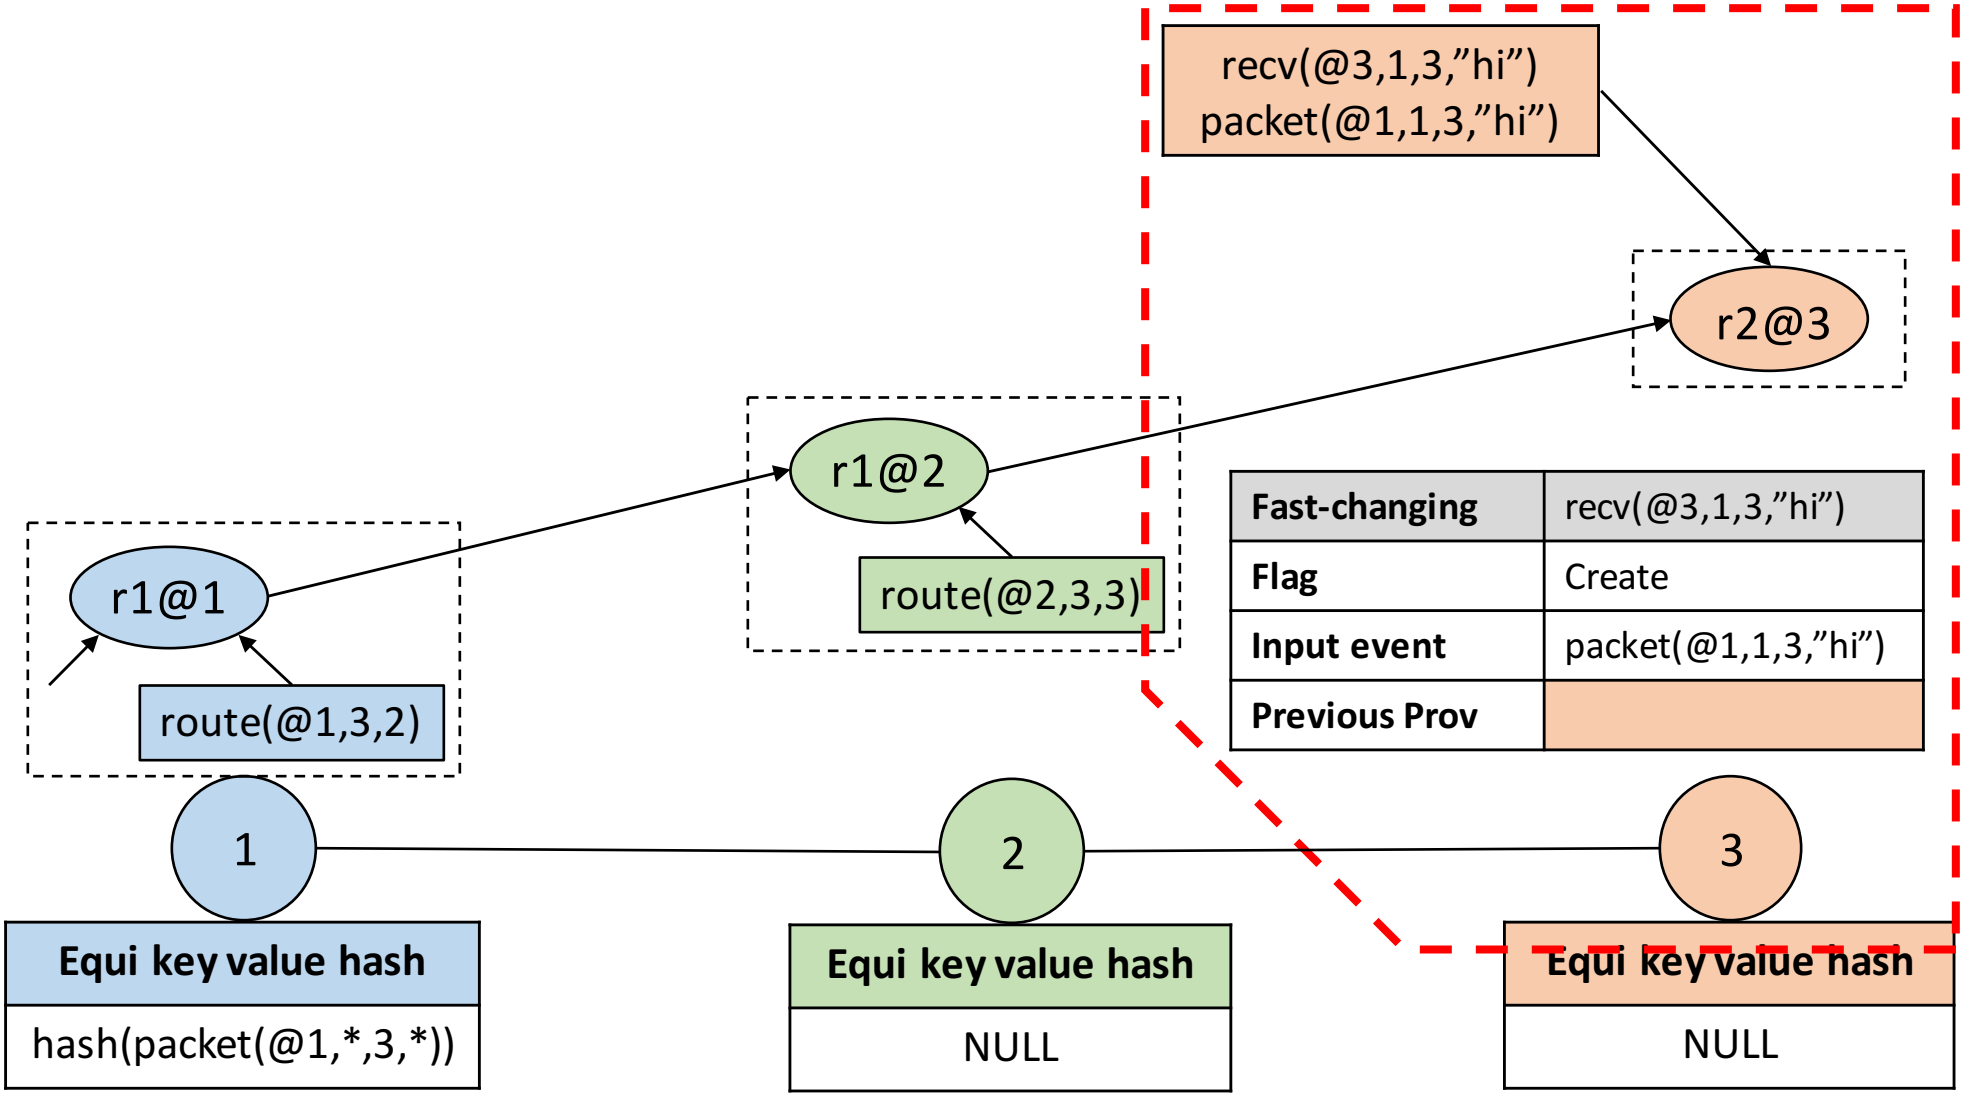

No provenance stored

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

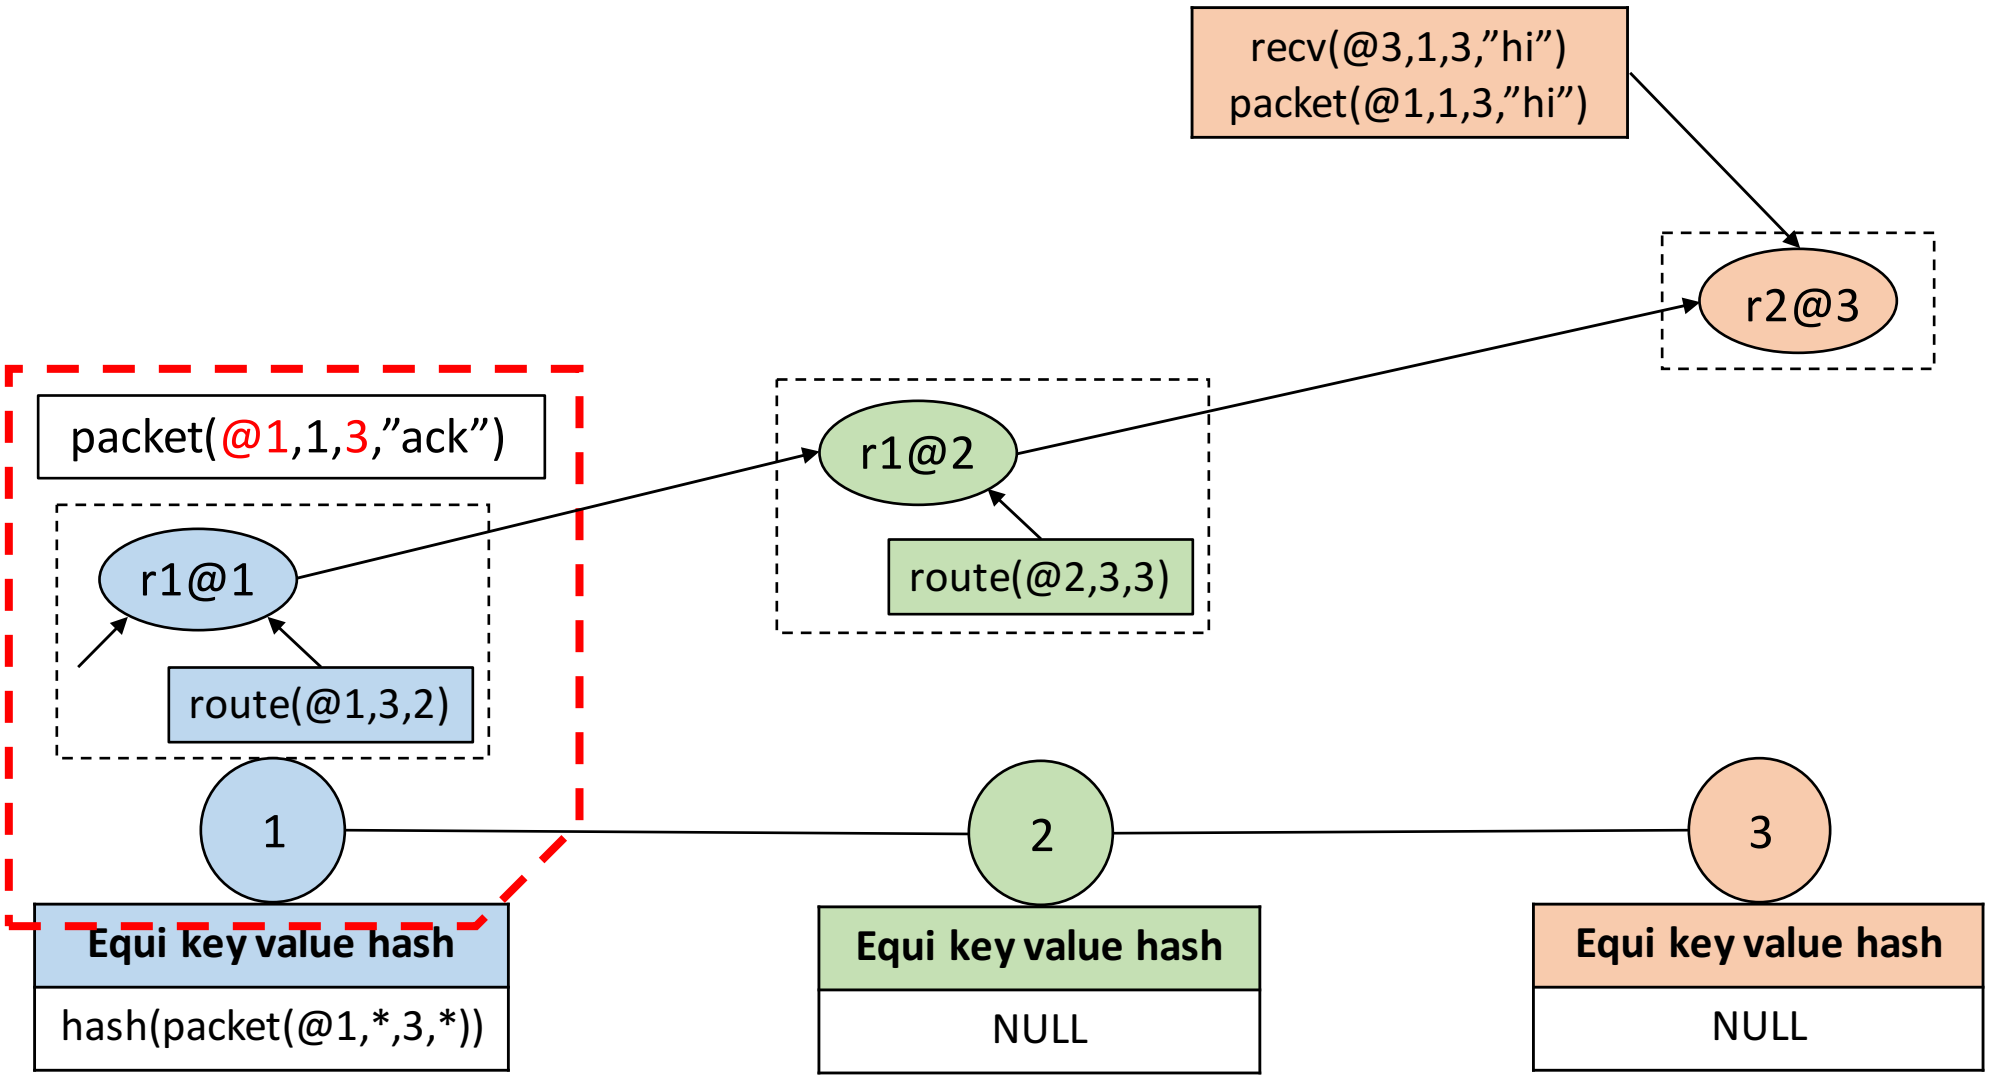

No provenance stored

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

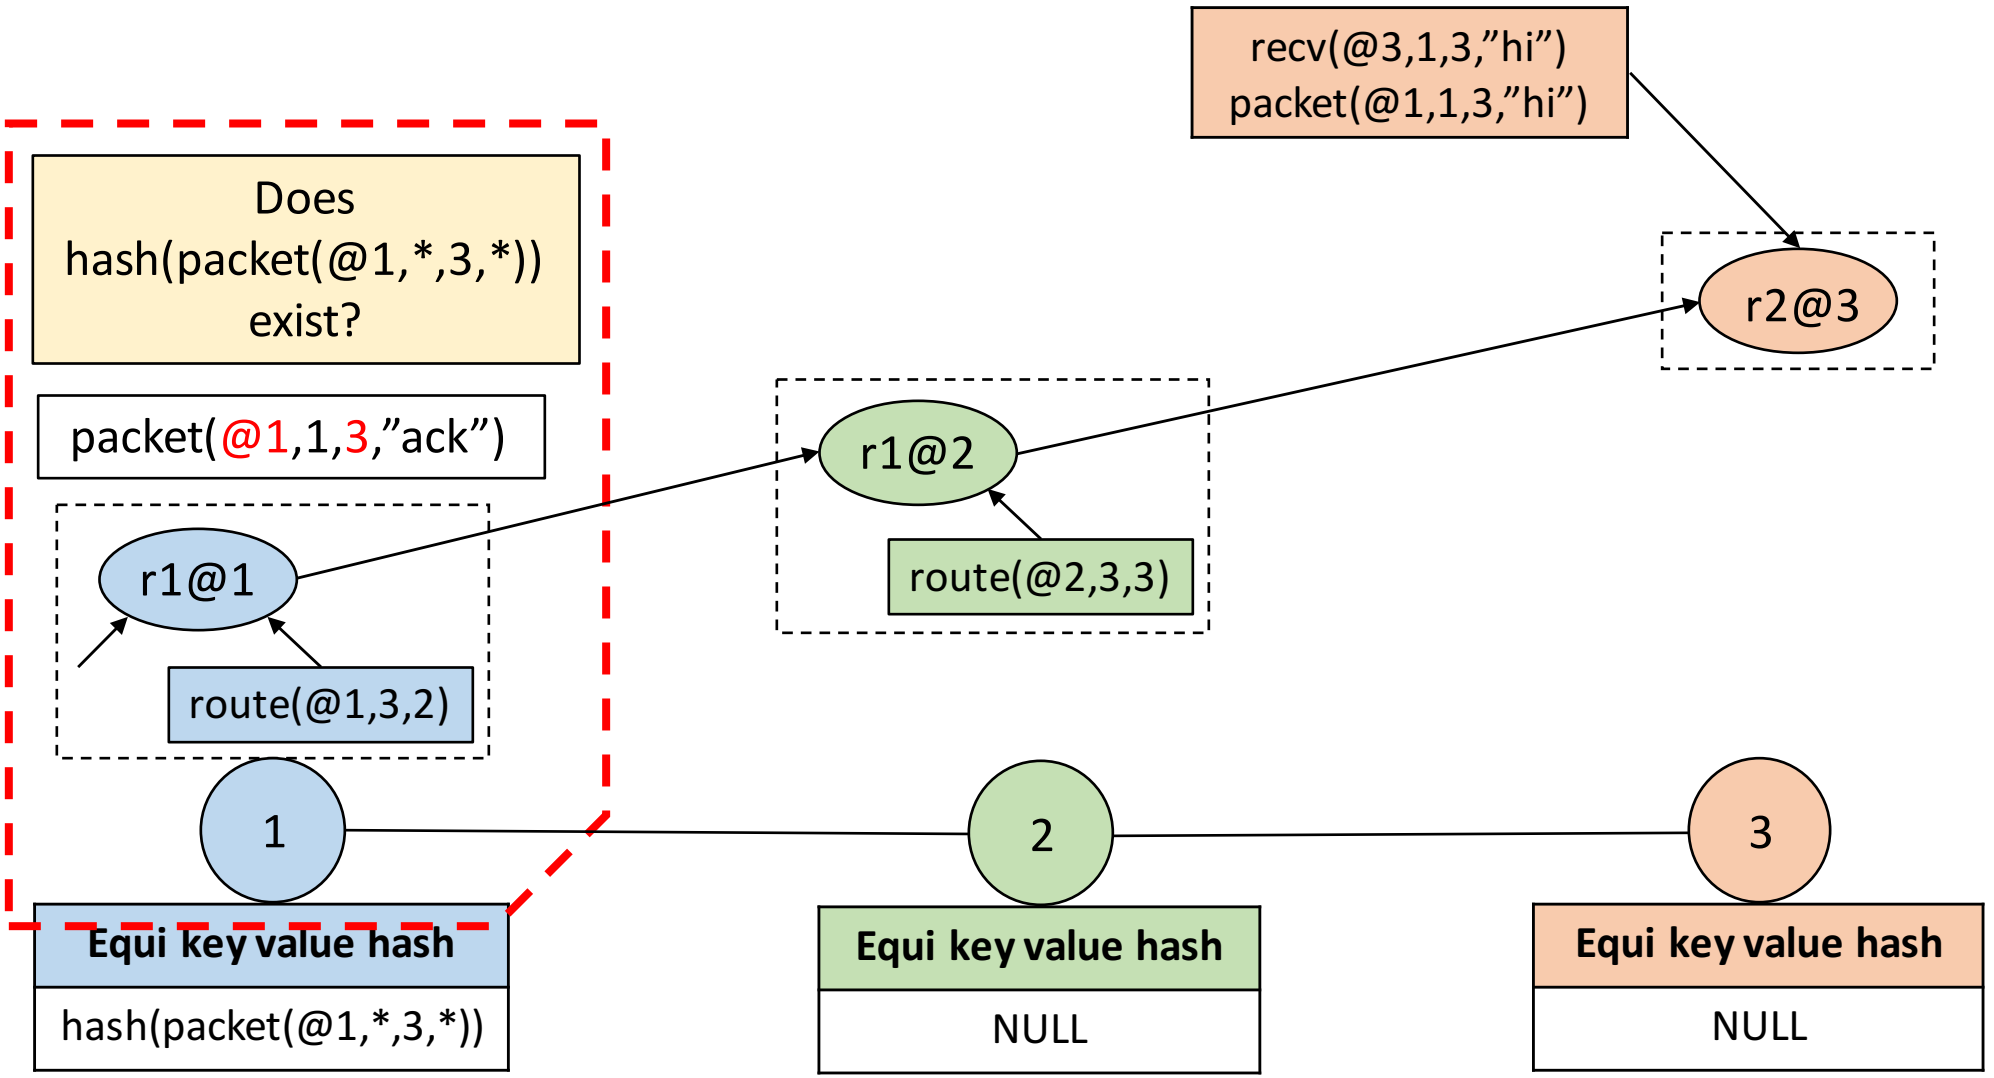

No provenance stored

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

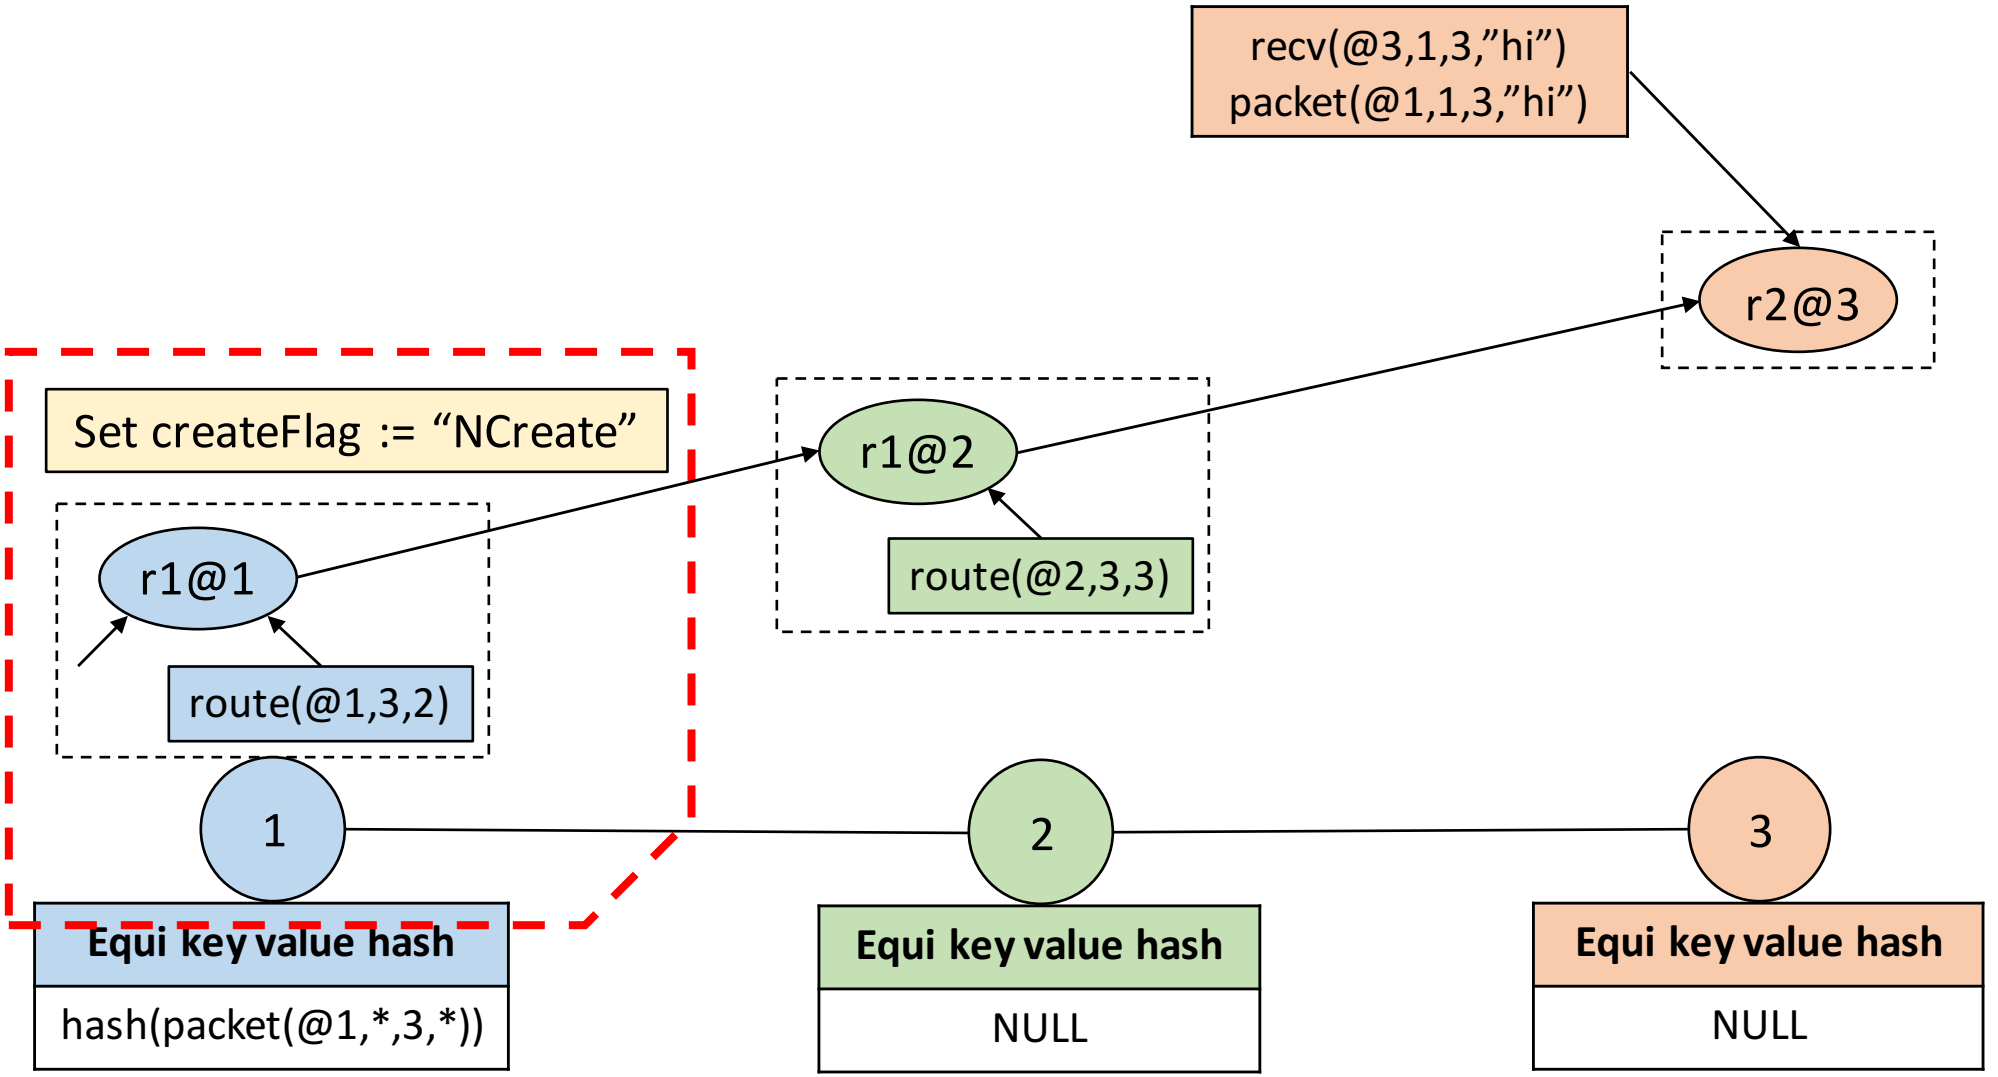

No provenance stored

Stage 1: Equivalence Key Checking

Packet Forwarding

r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).

r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

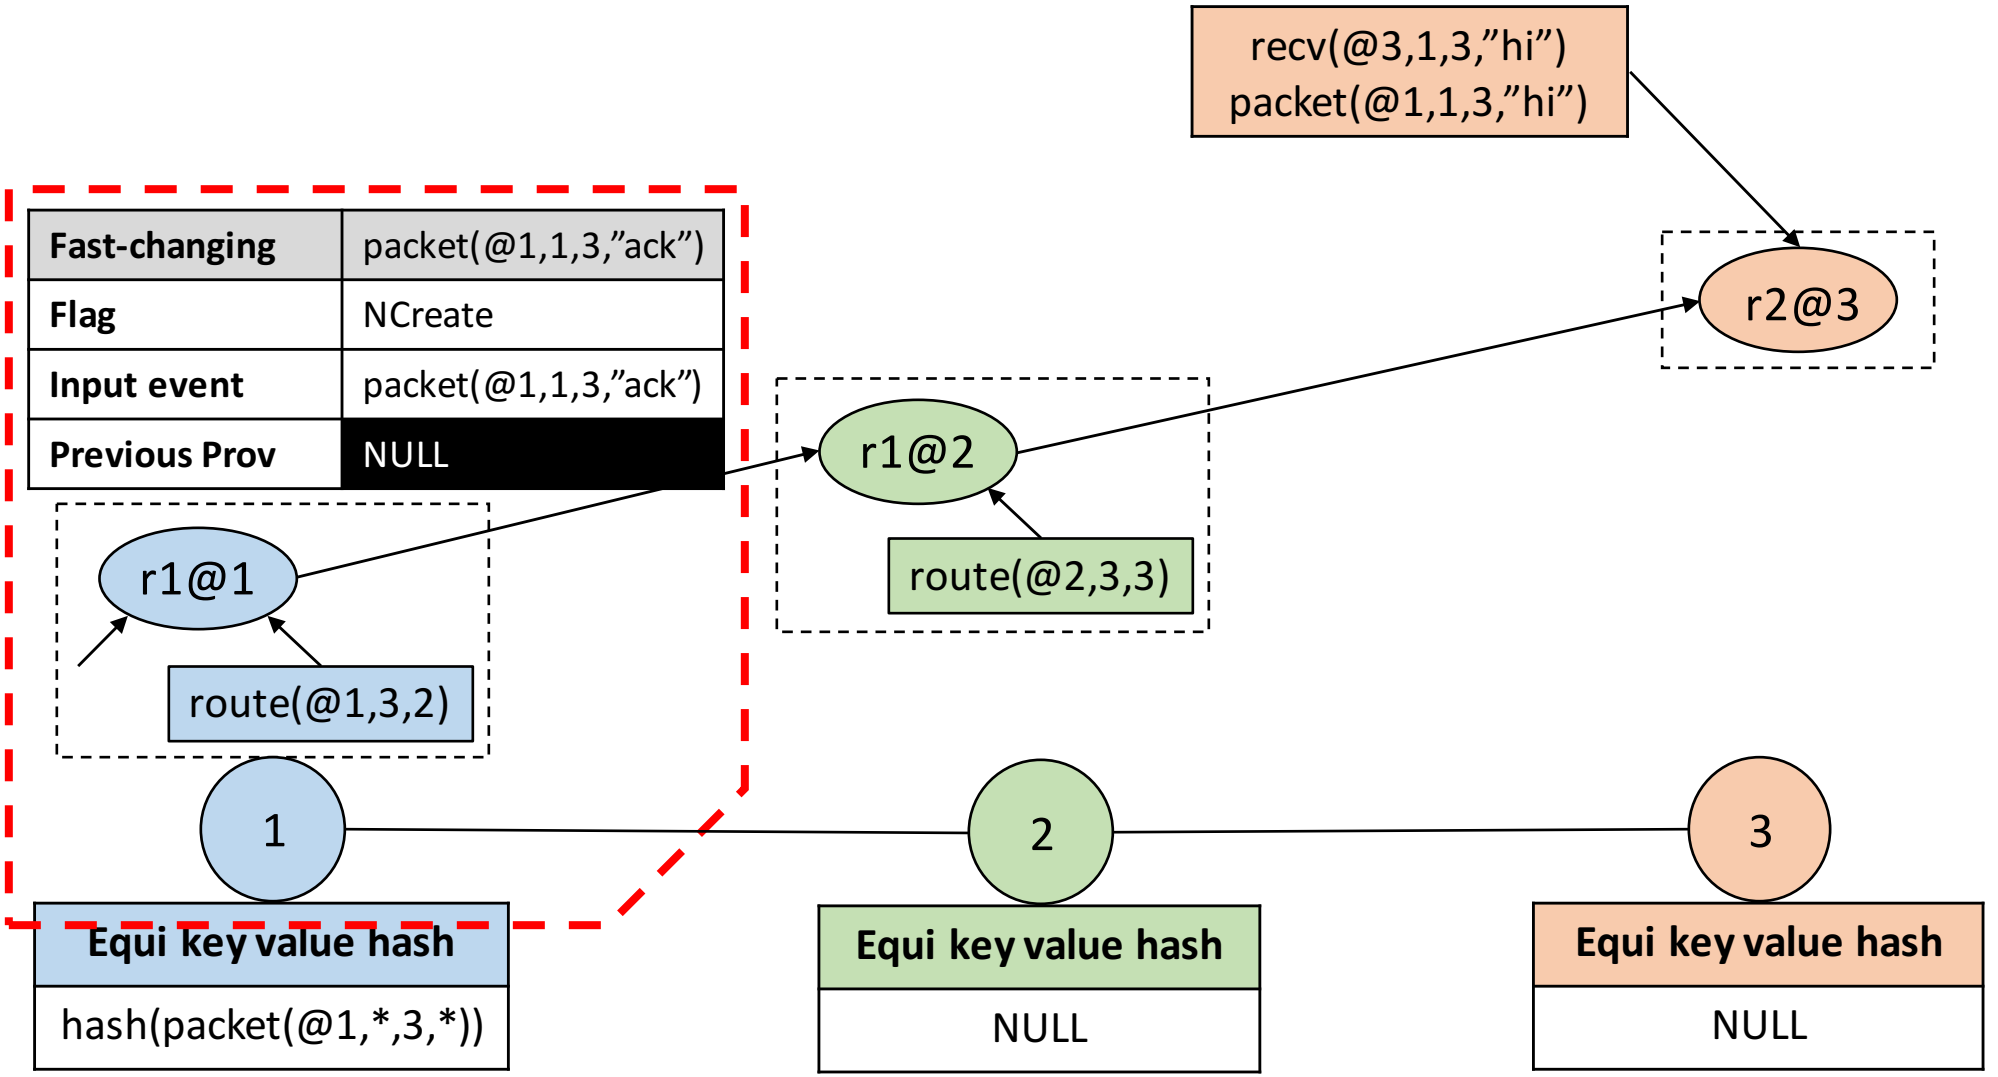

No provenance stored

Stage 3: Associate output tuple to its stored provenance tree

Packet Forwarding
r1 packet(@N,S,D,T) :- packet(@L,S,D,T), route(@L,D,N).
r2 recv(@L,S,D,T) :- packet(@L,S,D,T), D==L.

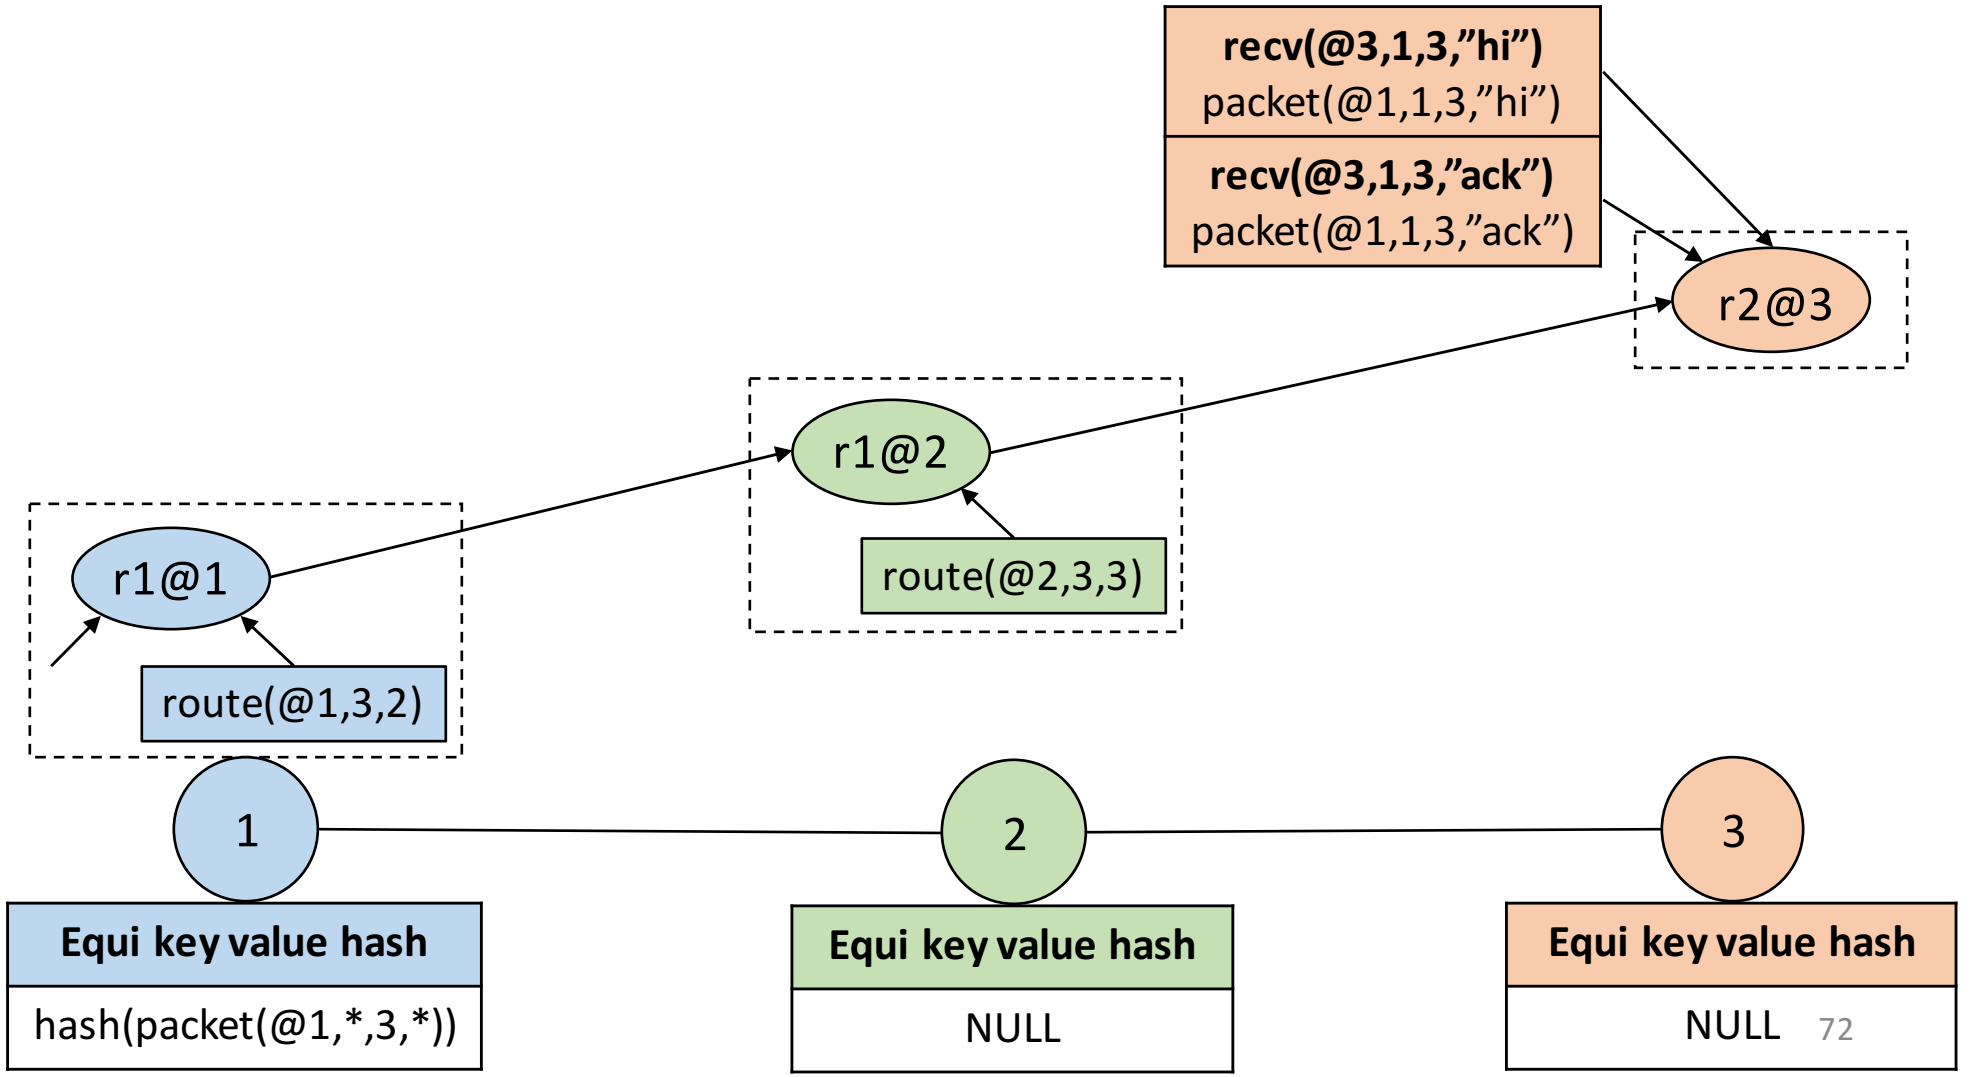

Initial network states

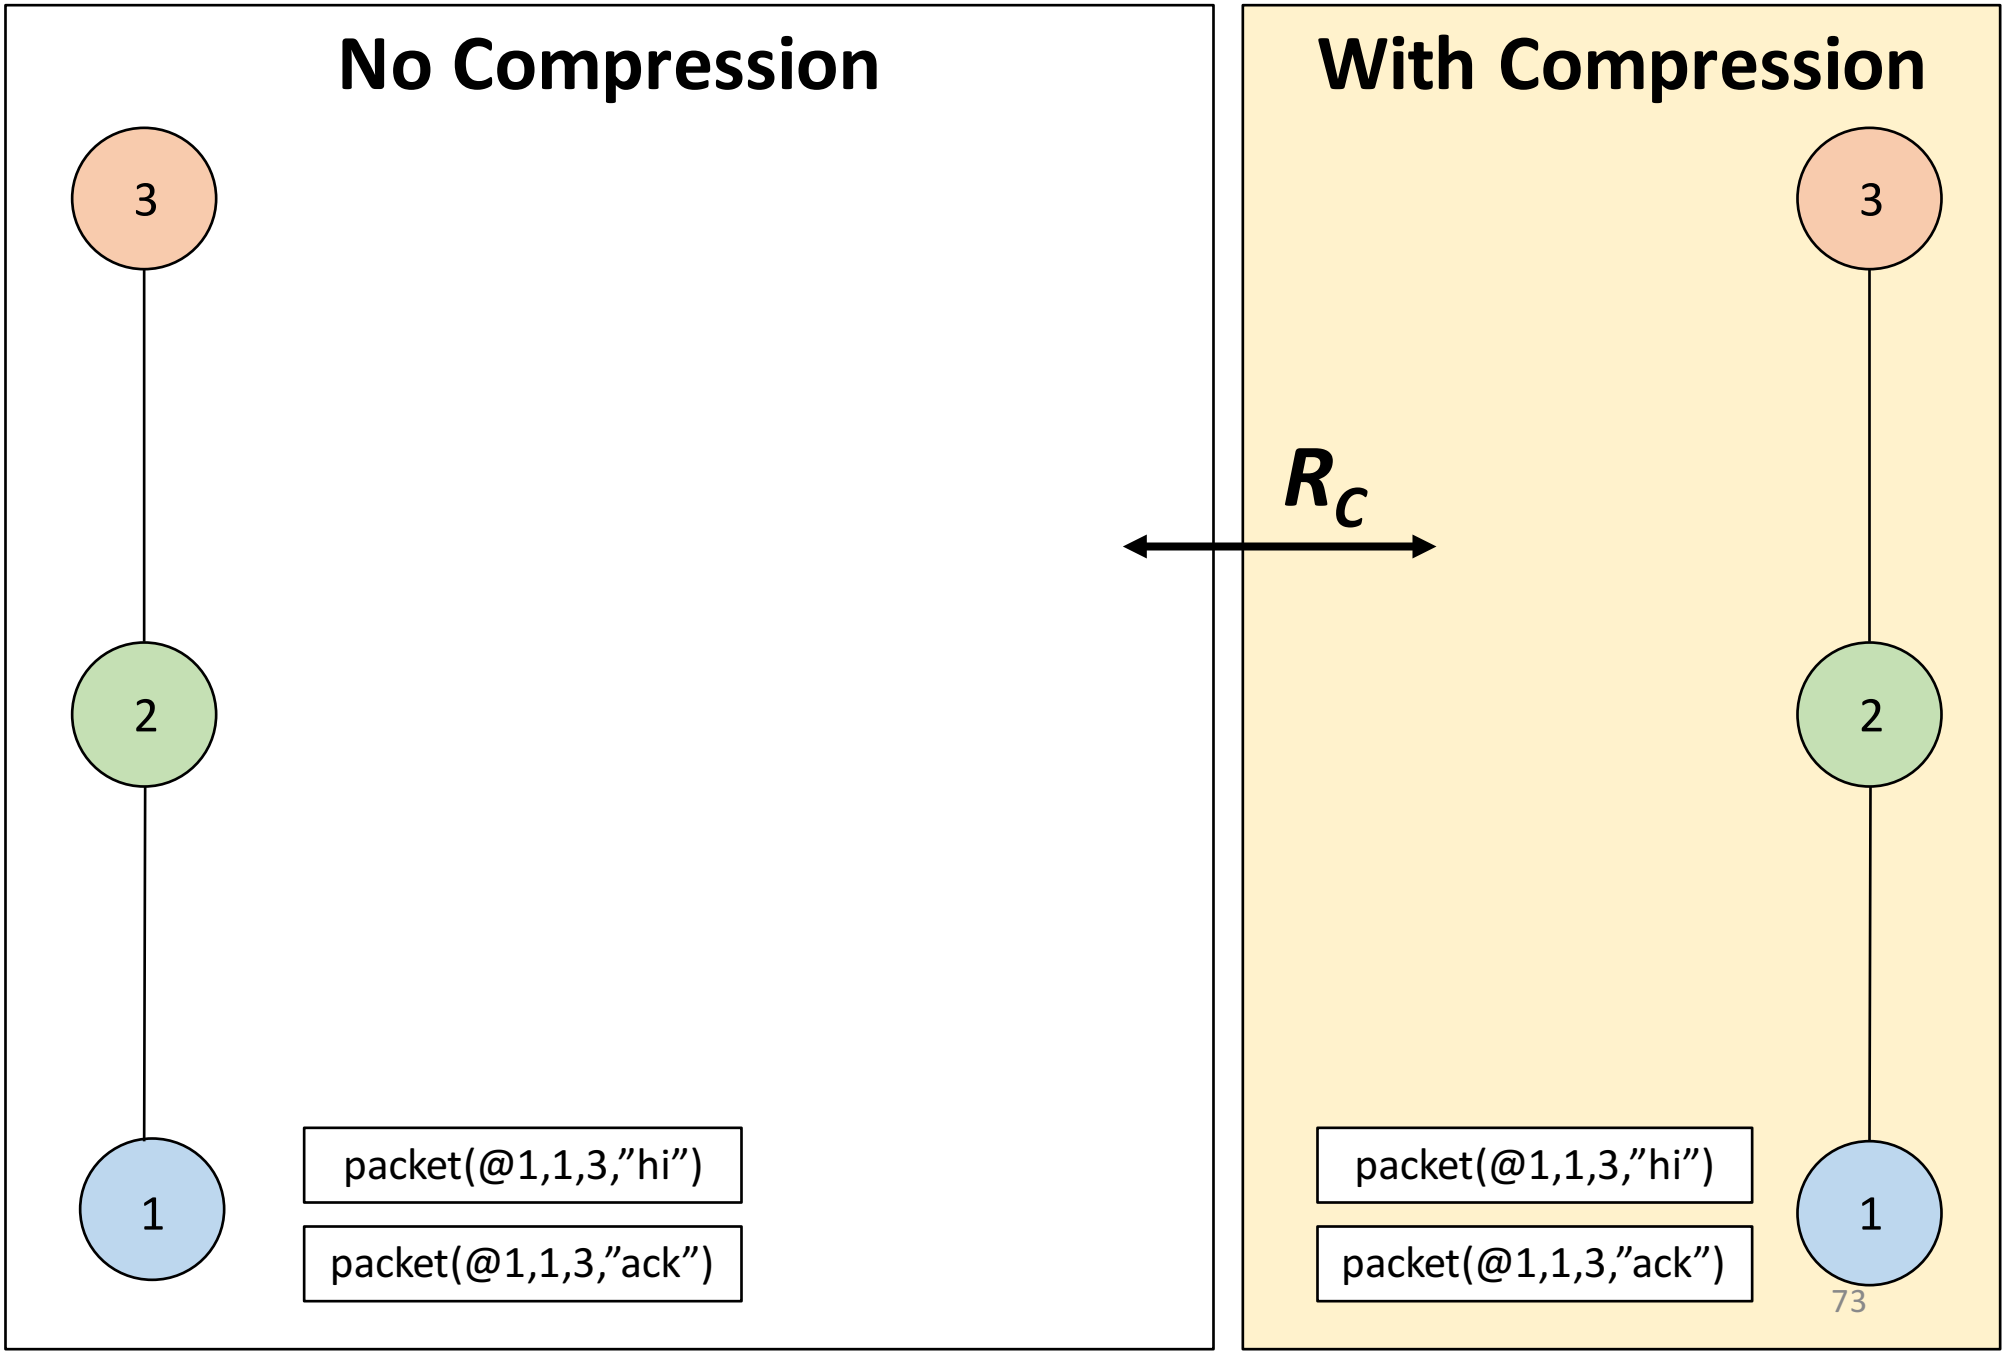

Final network states

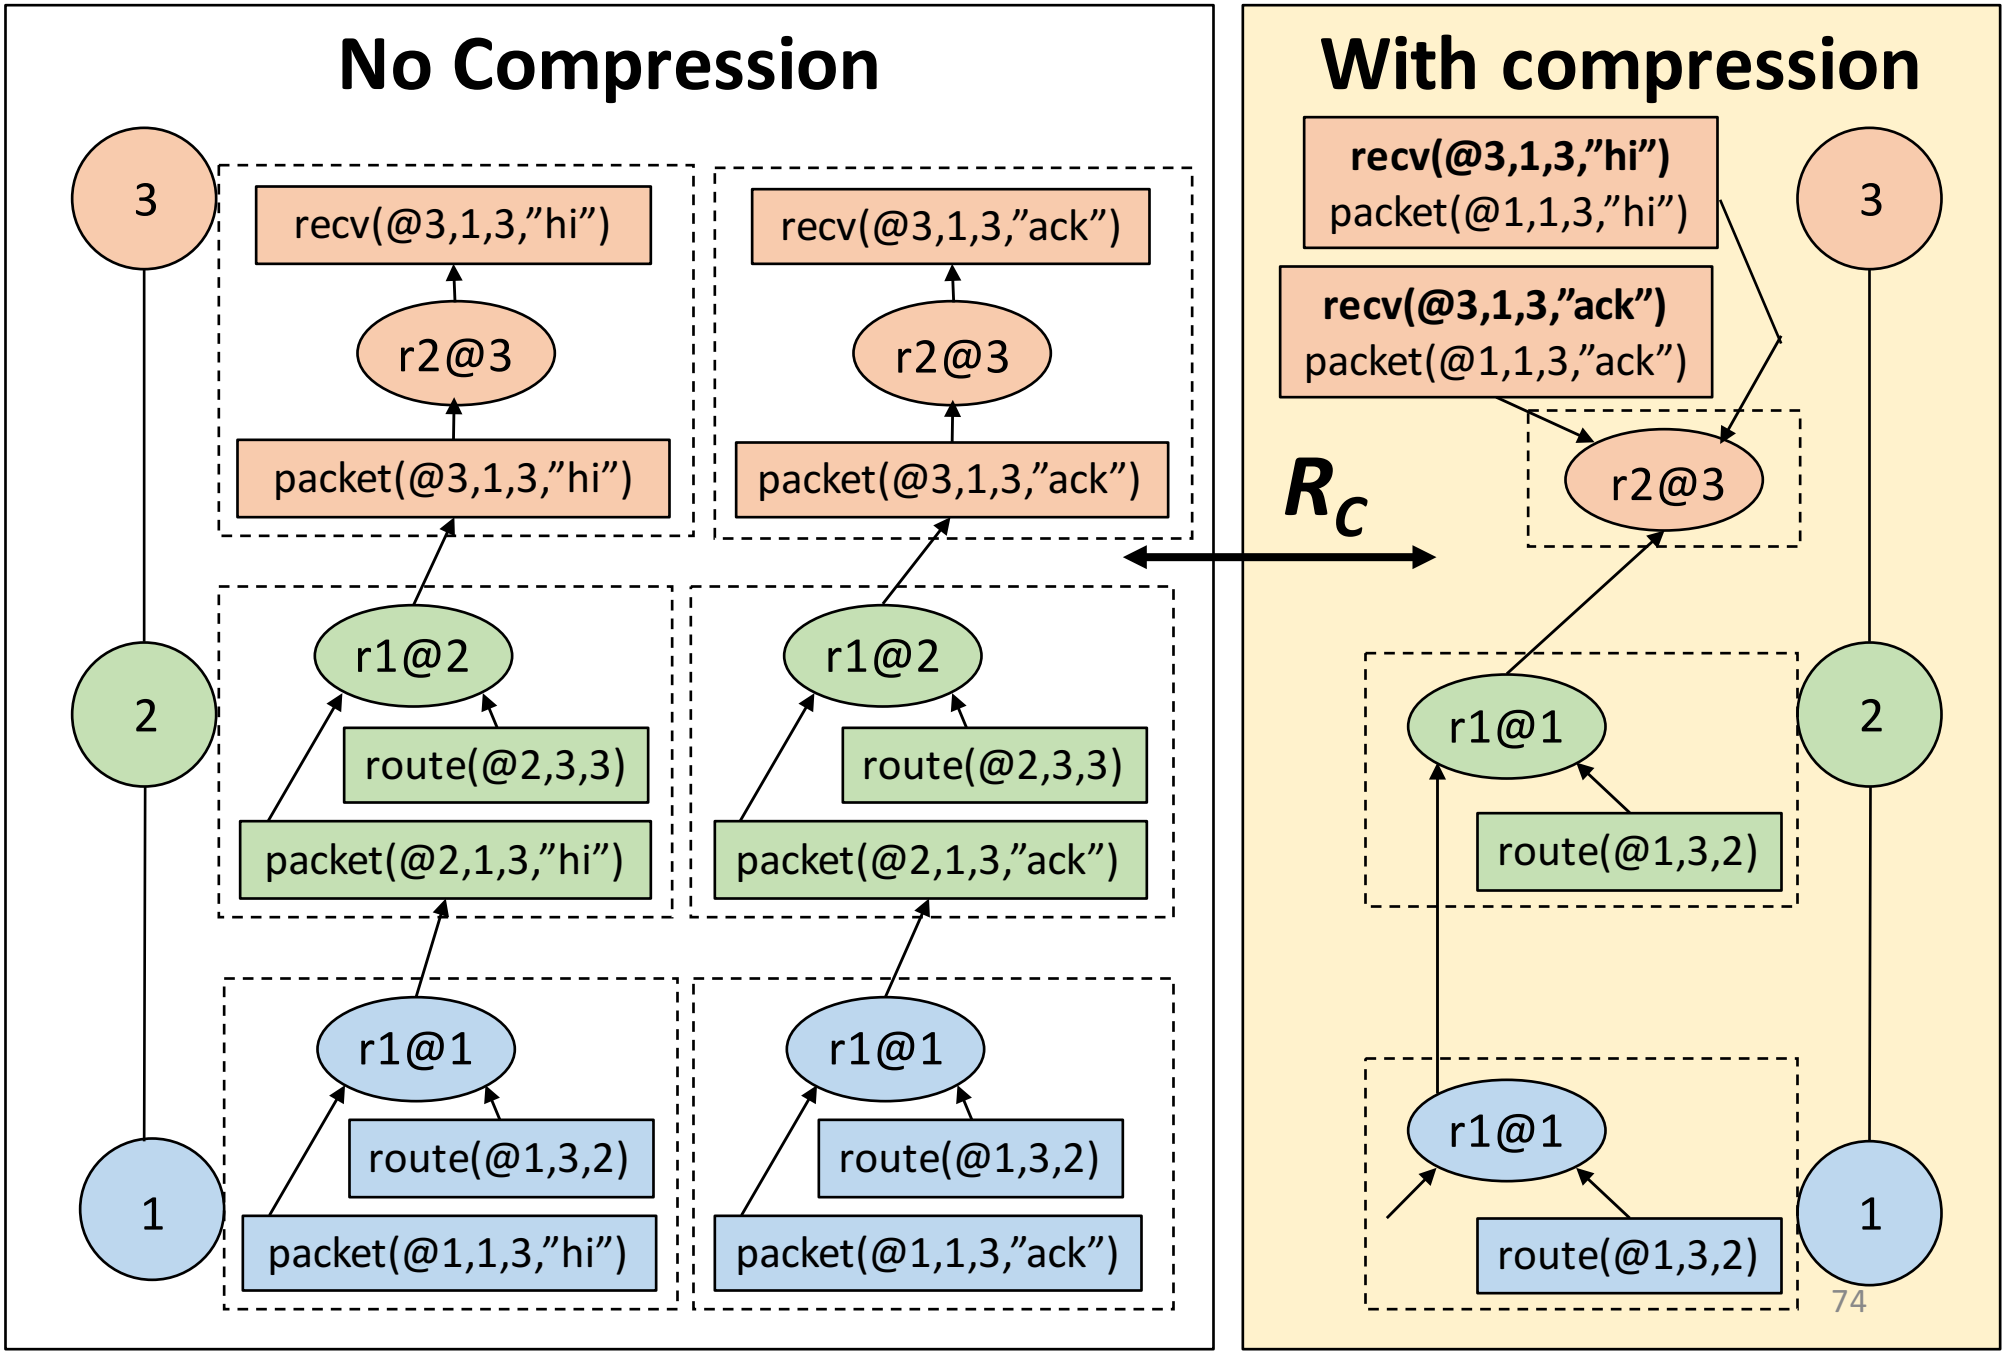

Relating network states

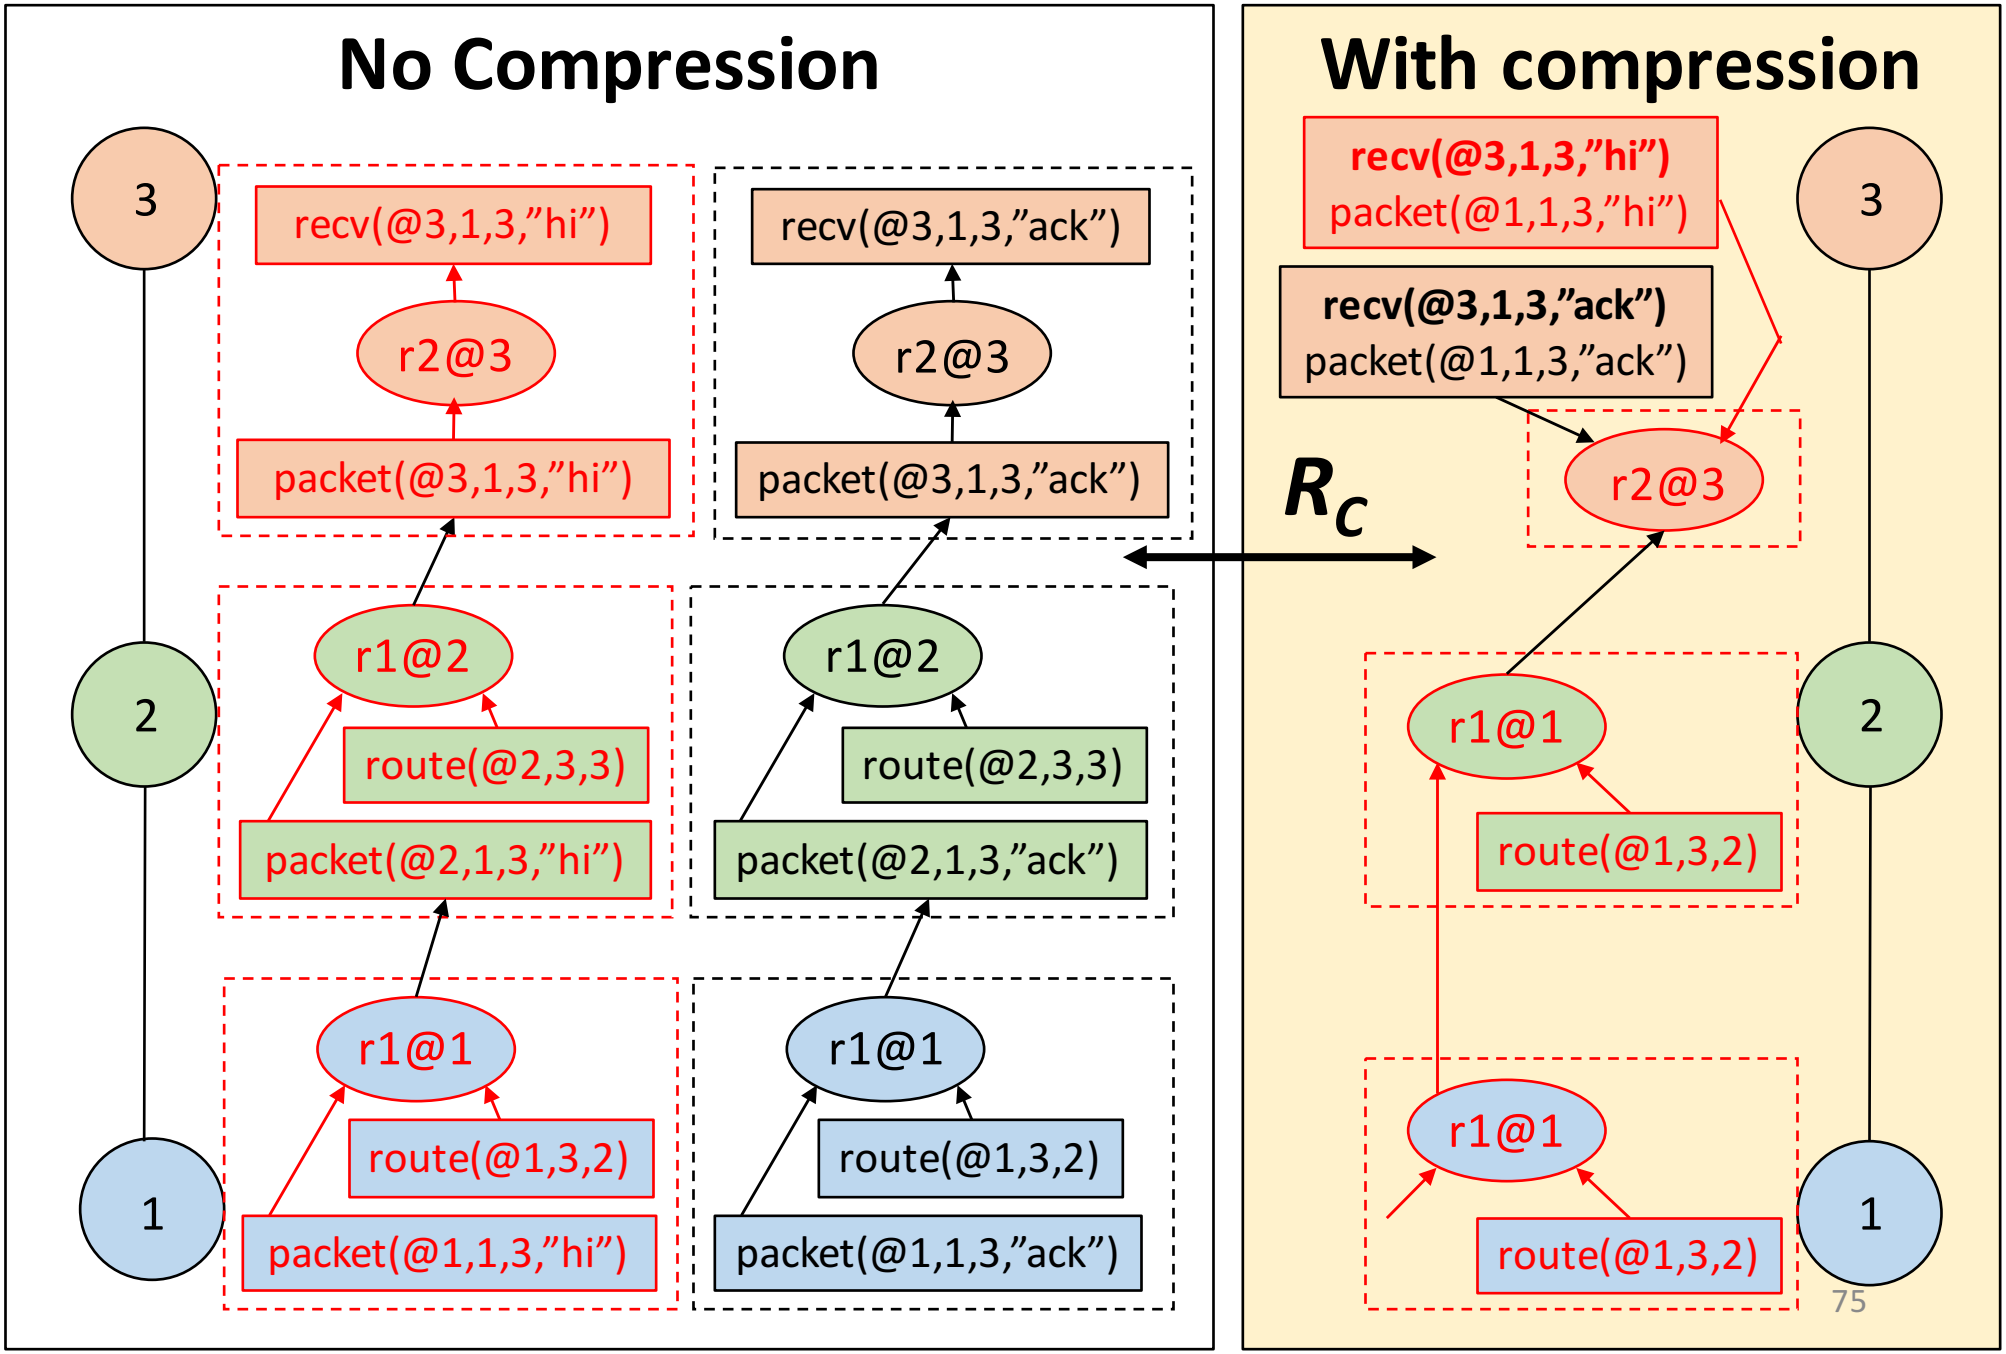

Relating network states

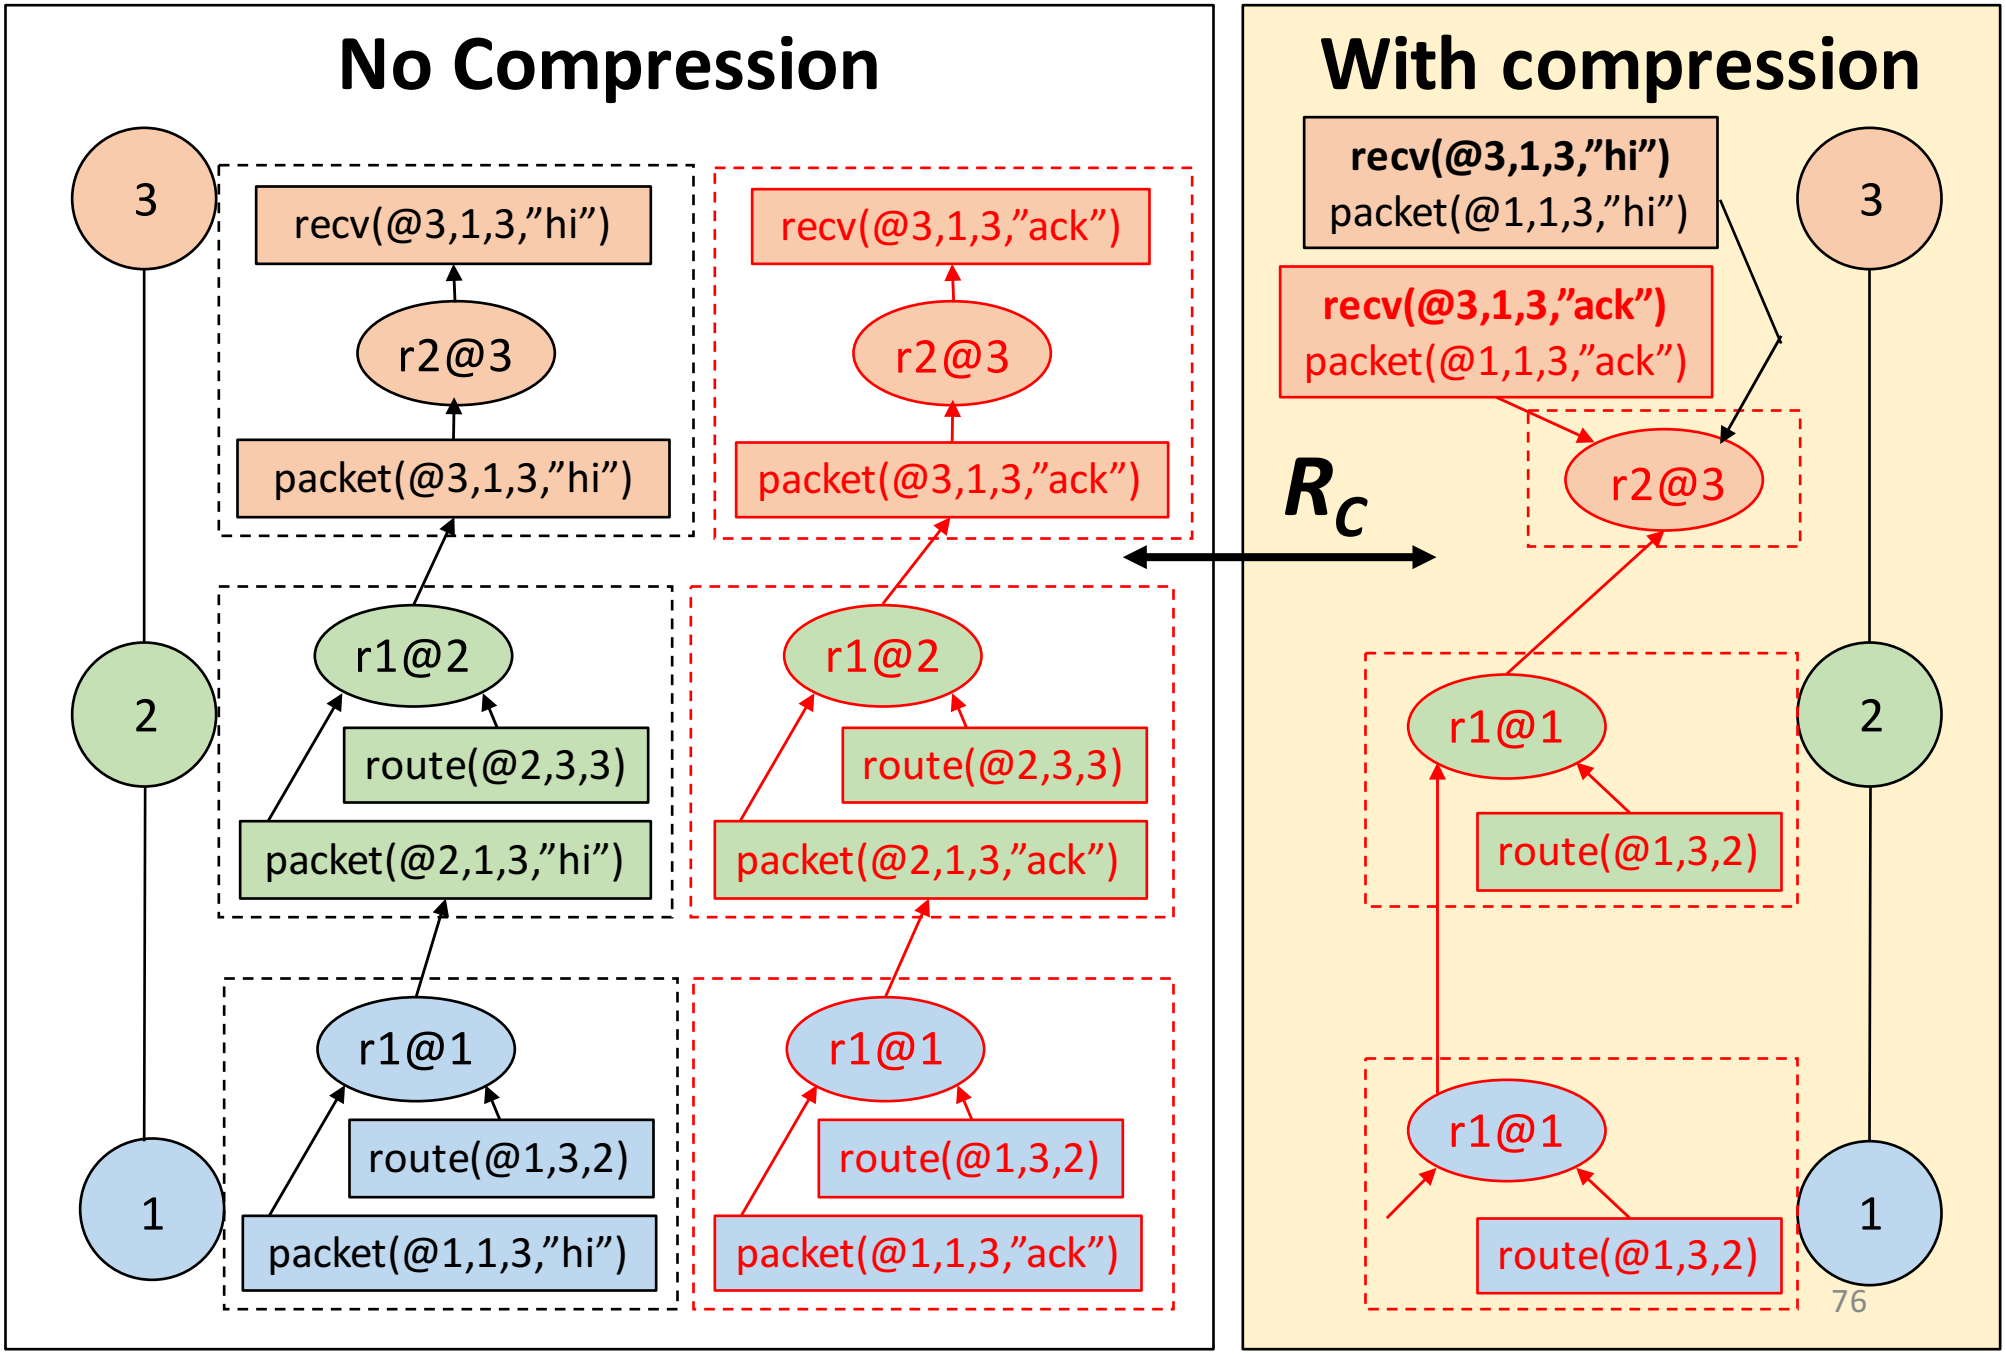

Roadmap

- Background
- Key insights
- Our compression scheme
- ***Conclusion***

Challenge

Large amount of storage needed to maintain network provenance in a distributed setting

Our Solution

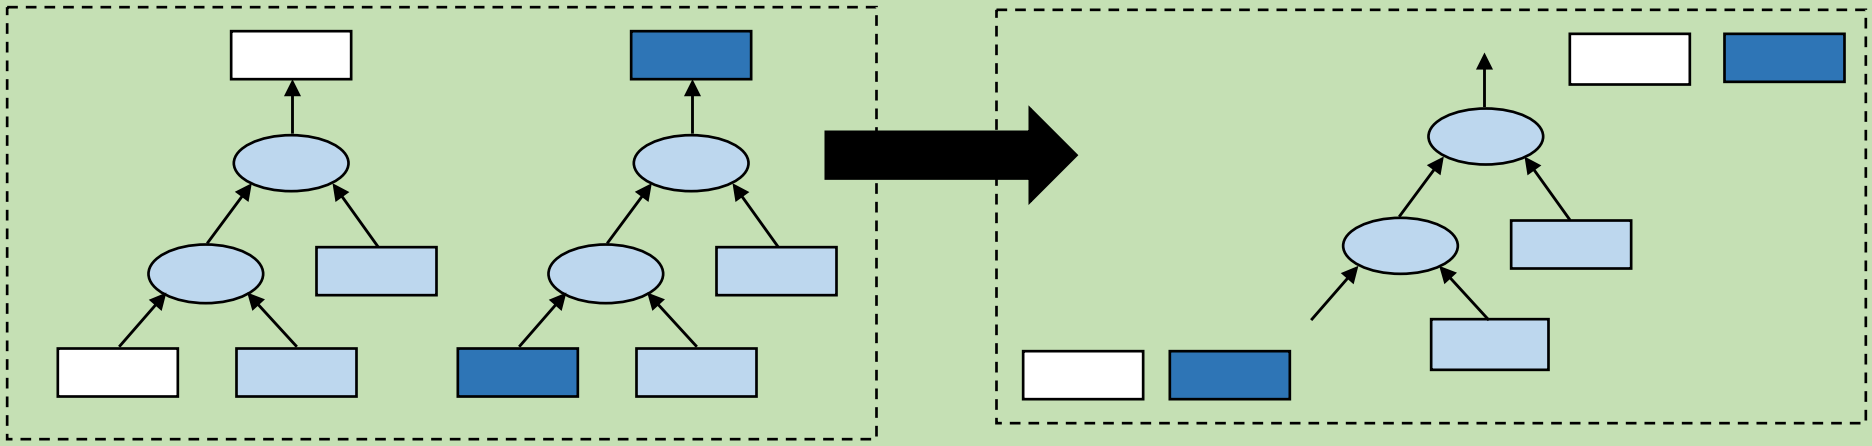

Summary

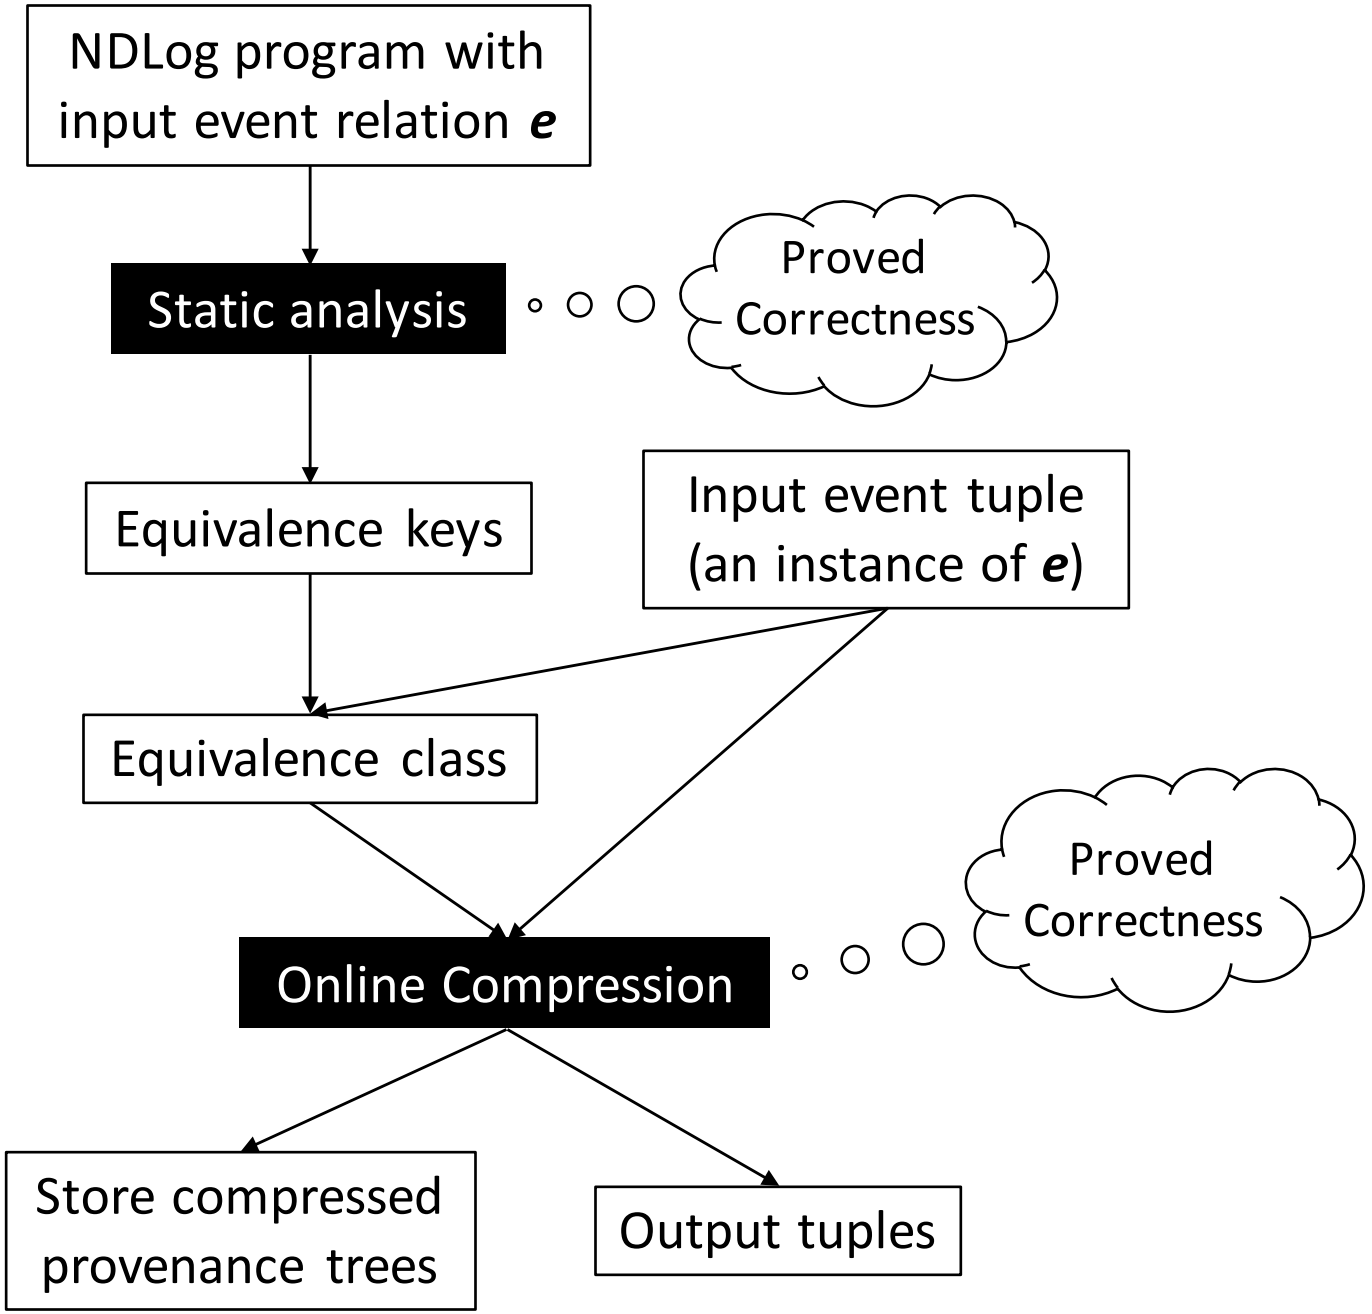
NDLog program with input event relation *e*
Static analysis
Proved Correctness
Equivalence keys
Input event tuple (an instance of *e*)
Equivalence class
Online Compression
Proved Correctness
Store compressed provenance trees
Output tuples